

Neural networks and Deep learning

1 Theory

1.1 Introduction

1.1.1 Introduction to Machine Learning

1.1.1.1 Supervised learning

- Where an algorithm learns from **labeled data**.
- More formally:
 - Given a training set $\mathcal{X} = \{(\mathbf{x}_i, y_i)\}_{i \in [N]}$.
 - Learn a function $g(\mathbf{x}; \theta)$ that approximates the *true* function $f(\mathbf{x})$ that generated the data.
 - $g(\circ)$ is the model, θ are its parameters.
 - g is assumed to **generalize** to unseen data.

1.1.1.2 Tasks

- Tasks in ML specify the nature of the problem to be solved.
 - Classification, regression, clustering, association rules extraction, etc.
- **Classification:**
 - A **supervised learning** task.
 - *Goal:* build a model to **classify** instances of an underlying concept using given **labels**.
 - eg Spam detection, sentiment analysis, object detection, LLM, etc.
- **Regression:**
 - $y \in \mathbb{R}$ while in classification y is a categorical value.
 - eg stock price prediction, weather forecasting, image generation, etc.

1.1.1.3 Induction

Induction: inference of a generalized conclusion from particular instances.

- ML algorithms **optimizes** the parameters, θ , in such a way a **loss function** L is minimized.
 - $\theta^* = \arg \min_{\theta} L(g(\mathbf{X}; \theta), \mathbf{y})$.
 - This assumes that the **induction** process hold in the applicative settings and generalizes to **unseen data**.
- **Induction** is precisely what every algorithms aims to achieve when learning a classification function.
- **No Free Lunch Theorem:** «all algorithms are equivalent, on average, by any of the following measures of risk».
 - Or, «all models are wrong, but some models are useful».
 - A learning algorithms is only as good as its **inductive bias**.
 - How well the inductive bias of the algorithms fits the problem at hand must be considered.
- eg Examples of inductive bias:
 - A **linear classifier** assumes that:
 - Examples lie in a metric space.
 - The underlying concept is linear in nature.
 - A **convolutional neural network** assumes that:
 - The underlying concept is translation invariant.
 - The input is an image (a 2D grid of pixels where adjacent pixels are correlated).
 - A **recurrent neural network** assumes that:
 - The underlying concept is sequential in nature.

1.1.1.4 Training a neural network

- NN are a incredible flexible family of models that can be used to approximate any function.
 - This flexibility comes at a cost: they depend on a vast number of **hyperparameters**.
 - Like the number of layers, neurons in each layer, learning rate, optimizer, etc.
 - These hyperparameters need to be set correctly in order to achieve good performance.
 - It is then important to understand how to **set these hyperparameters correctly**.
 - And how to validate the quality of the inferred model **without overestimating its performance**.
- Problem statement:
 - Define the correct procedure to:
 - Find the best possible way to **set the hyperparameters**.
 - Get an **accurate estimation of the generalization error** at the end of the process.

1.1.1.5 Errors

Generalization error

- **def Generalization errors:** $R = \mathbb{E}_{(x,y) \sim p^*} [L(y, g(x, \theta))]$.
- The error that the acquired classifier g commit, on average, on examples drawn from the same distribution used for sampling the examples in the training set.
- p^* is the true distribution of the data.
- L is a loss function used to measure how bad the error is when y is predicted as $g(x, \theta)$.
- Problems:
 - No actual access to p^* .
 - Even with access to p^* , it's not possible to compute an average on an infinite number of samples.

Loss functions

- **def 0 – 1 loss:** $L(y, y') = \mathcal{I}_{y \neq y'} = 1$ if $y \neq y'$, 0 otherwise.
- **def Quadratic loss:** $L(y, y') = (y - y')^2$.
- But many others are possible (eg hinge, exponential, logistic, etc).

Empirical error

- **def Empirical error:** $\hat{R}_T = \frac{1}{|T|} \sum_{(x,y) \in T} L(y, g(x; \theta))$.
- It should be evident that (in general) R cannot be computed.
 - To overcome this problem in most cases in R is approximated with the empirical error.
- T is a finite sample drawn from p^* .

Training and test errors

- **Training error** (\hat{R}_{Tr}): when T is the **training set** ($T = Tr$).
 - Since it's implicitly (sometimes explicitly) optimized by the learning algorithm, it has an **optimistic bias**.
- **Test error** (\hat{R}_{Te}): when T is the **test set** ($T = Te$).
 - It serves as an unbiased estimator of the generalization error R .
 - It may exhibit a **pesimistic bias** because retraining the model on the entire dataset is likely to produce a model with a lower error.

1.1.1.6 Overfitting

- When $\hat{R}_{Te} - \hat{R}_{Tr} > 0$, the algorithm is said to be **overfitting**.
- Overfitting is common problem for learning algorithms.
 - And it is usually necessary to counter it by some method.
- But minimizing the generalization error on the test set is a sort of tuning.
 - This makes the test set error an **optimistic estimator** of the generalization error.
 - And therefore **overfitting the test set**.
- To overcome this problem, an extra set, the **validation set**, is needed.
 - Used to solely evaluate model performance under a given choice of hyperparameters.
 - Once the hyperparameters are set, Te is used to asses the final model's quality.
 - When data is scarce, choices have been made and the final quality is fine, retrain the classifier on $Tr + Va$.

- When the test set is no longer needed, it can be merged to the training set.
 - Selecting best hyperparameters for a model (?):
 - Both validation and test sets are **finite samples**, measuring the error on them will introduce some noise in the estimation of the generalization error.
 - $\hat{R}_{Te} = R + \epsilon_{Te}$ and $\hat{R}_{Val} = R + \epsilon_{Val}$.
 - The model is chose according to the hyperparameters that minimize the error on the validation set.
 - Determine if at the end of the process \hat{R}_{Te} and \hat{R}_{Val} are **unbiased estimators of the GE**.
 - Which is, $\mathbb{E}[\hat{R}_{Te}] = \mathbb{E}[\hat{R}_{Val}] = R$.
 - The distribution of Test errors is uniform and has mean μ_0 .
 - The distribution of Validation errors is skewed to the left, resulting in an under-estimation.
 - Training, validation and test sets:
 - **Training set**: used during learning.
 - **Validation set**: used to assess the quality of the current choice of hyperparameters.
 - **Test set**: used to asses the quality of the final classifier.
-

1.2 Mathematical foundations

1.2.1 Matrices

- In ANN, vectors and matrices are everywhere.
 - Inputs are vectors, weights are matrices.
 - Vectors and matrices must be multiplied to calculate network output.
- A **scalar** is just a single number x .
- A **vector** is an array of numbers.
- A **matrix** is a $2D$ array of numbers.
 - Indices: i for the row and j for the column.
- A **tensor** is an array with more than two dimensions.
 - Indices: i for the batch (slice), j for the row and k for the column.
- The **transpose** \mathbf{A}^T of a matrix \mathbf{A} is a mirror image, $(\mathbf{A}^T)_{i,j} = \mathbf{A}_{j,i}$.
- If matrices have the same shape, they can be **added**, $C_{i,j} = A_{i,j} + B_{i,j}$.
 - A scalar can be added or multiplied to, $D_{i,j} = a \cdot B_{i,j} + c$.
- **Matrix multiplication**: if A has shape $m \times n$ and has shape B $n \times p$, $\mathbf{A} \cdot \mathbf{B}$.
 - $\mathbf{A} \cdot \mathbf{B} \neq \mathbf{B} \cdot \mathbf{A}$.
- **Vector multiplication**: $\mathbf{x} \cdot \mathbf{y} = \mathbf{x}^T \mathbf{y}$.

1.2.1.1 Geometric interpretation

- Vector can be represented geometrically.
 - $\mathbf{x} = [x_1, x_2]$ will be a vector originating at $(0, 0)$ and arriving at (x_1, x_2) .
- def **Euclidean norm** (or L^2 norm): $\|\mathbf{x}\| = \sqrt{x_1^2 + x_2^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}}$.
 - The euclidean distance from the origin to the point identified by \mathbf{x} .
 - def **Euclidean norm**: $\|\mathbf{x}\| = \sqrt{\sum_i x_i^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}}$.
 - The squared euclidean norm is also used.
- def **L^1 norm**: $\|\mathbf{x}\|_1 = \sum_i |x_i|$.
- def **Dot product**: $\mathbf{x} \cdot \mathbf{y} = \|\mathbf{x}\| \|\mathbf{y}\| \cos(\theta)$.

1.2.1.2 Linear transformation

- Matrices are a tool to **linearly transform** vectors (and vector space).
- **Determinant** of a squared matrix, $\det(\mathbf{A})$, is a function mapping matrices to real scalar.
 - $\det(\mathbf{A})$ gives a measure of how multiplication by the matrix expands or contract space.
 - If $|\det(\mathbf{A})| = 1$ the volume remains unchanged.
 - If $\det(\mathbf{A}) = 0$ the space is contracted along at least one dimension.
 - For \mathbf{A} with dimension 2, $\det(\mathbf{A}) = a_1 \cdot a_4 - a_2 \cdot a_3$.

1.2.1.3 Special matrices

- **Identity matrix**: it does not change any vector when multiplied the matrix by the vector.
- **Inverse matrix** (\mathbf{A}^{-1}): the matrix such as $\mathbf{A}^{-1} \mathbf{A} = I_n$.
 - Matrices for which \mathbf{A}^{-1} exists are **invertible** (only squared matrices).

- prop \mathbf{A} is invertible iff $\det(\mathbf{A}) \neq 0$.
- When \mathbf{A}^{-1} exists, there are several algorithms for finding it.
- **Symmetric matrix:** any matrix equal to its transpose ($\mathbf{A} = \mathbf{A}^T$).
- **Unit vector:** a vector with unit norm ($\|\mathbf{x}\|_2 = 1$).
- **Orthogonal vectors:** \mathbf{x} and \mathbf{y} to each other if $\mathbf{x}^T \mathbf{y} = 0$.
 - Vectors orthogonal to each other and with unit norm are **orthonormal**.
- **Orthogonal matrix:** a square matrix whose rows are mutually orthonormal and whose columns are mutually orthonormal.
 - prop For orthogonal matrices $\mathbf{A}^T \mathbf{A} = \mathbf{A} \mathbf{A}^T = \mathbf{I}$.
 - Hence $\mathbf{A}^{-1} = \mathbf{A}^T$.

1.2.2 Calculus

1.2.2.1 Derivatives

- The slope of the tangent line to f at point x (denoted as $f'(x)$ or $\frac{df}{dx}(x)$).
- It specifies how to scale a small change in input in order to obtain the corresponding change in output.
 - $f(x + \epsilon) \approx f(x) + \epsilon f'(x)$.
- The secant line becomes the tangent line when $\epsilon \rightarrow 0$, $f'(x) = \lim_{h \rightarrow 0} \frac{f(x+\epsilon)-f(x)}{\epsilon}$.
- prop Properties of derivatives:
 - Linearity: $(\alpha f(x) + \beta g(x))' \equiv \alpha f'(x) + \beta g'(x)$.
 - **Chain rule:** $(f(g(x)))' \equiv f'(g(x))g'(x)$.
 - Product rule: $(g(x)h(x))' \equiv g'(x)h(x) + g(x)h'(x)$.
 - Quotient rule: $(\frac{f(x)}{g(x)})' \equiv \frac{f(x)'g(x) - f(x)g'(x)}{(g(x))^2}$.
 - Power rule: $(x^r)' \equiv rx^{r-1}$.

1.2.2.2 Integrals

- The integral value of f between a and b is the area under f in the given region.
 - When the function is below 0, the area contributes negatively.
- def **Fundamental Theorem of Calculus:**
 - If f admits an antiderivative F (if it exists F such that $F'(x) = f(x)$) then:
 - $\int f(x) dx = F(x) + C$.
 - $\int_a^b f(x) dx = F(x)|_a^b = F(b) - F(a)$.
- Approximation integral area with rectangles: $\int_a^b f(x) dx \approx \sum_i f(x) dx$.
- prop Properties of integrals:
 - Linearity: $\int \alpha f(x) + \beta g(x) dx \equiv \alpha \int f(x) dx + \beta \int g(x) dx$.
 - Constant rule: $\int k dx \equiv kx + C$.
 - Power rule: $\int x^n dx \equiv \frac{x^{n+1}}{n+1} + C$ ($n \neq -1$).
 - Log rule: $\int \frac{1}{x} dx \equiv \ln(|x|) + C$.
 - Exponential rule: $\int a^{kx} dx \equiv \frac{a^{kx}}{k \ln a}$ ($a > 0, a \neq 1$).
 - Sine rule: $\int \sin(x) dx \equiv -\cos(x) + C$.
 - Cosine rule: $\int \cos(x) dx \equiv \sin(x) + C$.
 - Derivatives can be easily computed automatically, but for antiderivatives this is not true.
 - But those properties can be applied.

1.2.2.3 Partial derivatives

- A function $y = f(x_1, \dots, x_n) = f(\mathbf{x})$ is given, where $y \in \mathbb{R}$ and $\mathbf{x} \in \mathbb{R}^n$.
- $\frac{\partial}{\partial x_j} f(\mathbf{x})$ measures how f changes as only the x_j variable increases at point \mathbf{x} .
 - $\frac{\partial}{\partial x_j} f(\mathbf{x}) = \lim_{h \rightarrow 0} \frac{f(\mathbf{x} + h\hat{i}_j) - f(\mathbf{x})}{h} = \lim_{h \rightarrow 0} \frac{f(x_1, \dots, x_j + h, \dots, x_n) - f(x_1, \dots, x_n)}{h}$.
 - h multiplied by the versor \hat{i}_j is a *nudge* in the direction given by the current versor.
- def **Gradient** (of f): $\nabla_{\mathbf{x}} f$ (or ∇f) = $[\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n}]^T$.
 - The vector collecting all partial derivatives.
 - When $\nabla f(\mathbf{x}_0) = \mathbf{0}$, the tangent plane is horizontal.
 - The function doesn't grow in any direction along that point.

- prop **Chain rule for multivariate calculus**: $\frac{dz}{dt} = \frac{\partial z}{\partial x} \frac{dx}{dt} + \frac{\partial z}{\partial y} \frac{dy}{dt}$.
 - $z = f(x, y)$ is assumed and let x, y depend on an additional variable t (z can be seen as $f(x(t), y(t))$).
 - prop **Chain rule for multivariate calculus**: $\frac{df}{dt} = \sum_{i=1}^n \frac{\partial f}{\partial x_i} \frac{dx_i}{dt}$.
 - More in general for $f: \mathbb{R}^n \rightarrow \mathbb{R}$ when x_1, \dots, x_n depend on a variable t .

1.2.2.4 Directional derivatives

- def **Directional derivative**: $D_{\mathbf{u}}f(\mathbf{x}) = \lim_{h \rightarrow 0} \frac{f(\mathbf{x}+h\mathbf{u}) - f(\mathbf{x})}{h}$.
- DD of f at \mathbf{x} in \mathbf{u} direction: the rate of change in the direction given by the unit vector \mathbf{u} .
- The derivative of $f(\mathbf{x} + \alpha\mathbf{u})$ w.r.t. α evaluated at $\alpha = 0$.
- Using the chain rule the expression for $D_{\mathbf{u}}f(\mathbf{x})$ can be easily computed.
 - $D_{\mathbf{u}}f(\mathbf{x}) = \frac{d}{d\alpha}f(\mathbf{x} + \alpha\mathbf{u})|_{\alpha=0} = \sum_{i=1}^n \frac{\partial f(\mathbf{x} + \alpha\mathbf{u})}{\partial x_i}|_{\alpha=0} \frac{dx_i}{d\alpha} = \nabla f(\mathbf{x}) \cdot \mathbf{u} = \mathbf{u}^T \nabla f(\mathbf{x})$.

1.2.2.5 Gradient and optimization

- Goal: to find the direction in which the function increases the most.
 - To find \mathbf{u} such that $\nabla_{\mathbf{u}}f$ is largest.
- $\max_{\mathbf{u}, \mathbf{u}^T \mathbf{u}=1} D_{\mathbf{u}}f(\mathbf{x}) = \max_{\mathbf{u}, \mathbf{u}^T \mathbf{u}=1} \mathbf{u}^T \nabla f(\mathbf{x}) = \max_{\mathbf{u}, \mathbf{u}^T \mathbf{u}=1} |\mathbf{u}| |\nabla f(\mathbf{x})| \cos(\theta)$.
 - $\mathbf{u}^T \mathbf{u} = 1$ is used to ensure that the \mathbf{u} is still a unit vector.
 - $|\mathbf{u}|_2 = \sqrt{\sum_i u_i^2} = \sqrt{\mathbf{u} \cdot \mathbf{u}} = \mathbf{u}^T \mathbf{u}$.
 - $|\mathbf{u}| = 1$ and since $\nabla f(\mathbf{x})$ doesn't depend on \mathbf{u} .
 - Therefore, only to find \mathbf{u} that maximises $\cos \theta$ is needed.
 - Therefore the maximum is attained when \mathbf{u} is in the same direction as $\nabla f(\mathbf{x})$.
- ! The gradient points in the direction in which f increases the most.

1.2.2.6 Jacobian Matrix

- def **Jacobian matrix**: $\mathbf{J}_{i,j} = \frac{\partial}{\partial x_j} f(\mathbf{x})_i$ or $\mathbf{J} = [\nabla [f(\mathbf{x})_i]^T]_{i=1}^m$.
- A multi-valued, multi-variable function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m, f(\mathbf{x}) = [f(\mathbf{x}_1), \dots, f(\mathbf{x}_n)]$ is considered.
- The Jacobian matrix $\mathbf{J} \in \mathbb{R}^{m \times n}$ contains the partial derivative of all $f(\mathbf{x})_i$ for all variables x_j .
 - With $1 \leq i \leq m$ and $1 \leq j \leq n$.

1.2.2.7 Hessian Matrix

- A **second derivatives** is a derivative of a derivative.
 - Specify how the first derivative will change as the input is varied (it measures **curvature**).
 - With $f: \mathbb{R}^n \rightarrow \mathbb{R}$, n^2 second derivatives can be computed as $\frac{\partial^2}{\partial x_i \partial x_j} f(\mathbf{x})$.
 - There is a second derivative for each partial derivative.
- def **Hessian matrix** ($H(f)$): $H(f) = \mathbf{J}(\nabla f)$.
 - It contains all these partial derivatives.
- prop Properties of Hessian matrices:
 - The Hessian Matrix is **symmetric** where the second partial derivatives are continuous (Schwarz's Theorem).
 - Anywhere that the second partial derivatives are continuous, differential operators are commutative.
 - Therefore their order can be swapped: $\frac{\partial^2}{\partial x_i \partial x_j} f(\mathbf{x}) = \frac{\partial^2}{\partial x_j \partial x_i} f(\mathbf{x})$.
 - When $\nabla f(\mathbf{x}_0) = 0$, the Hessian helps detect **minimum**, **maximum**, or **saddle points**.
 - **Maximum**: if the Hessian is negative definite (i.e. all eigenvalues are < 0).
 - **Minimum**: if the Hessian is positive definite (i.e. all eigenvalues are > 0).
 - If the Hessian is neither positive nor negative definite (at least one zero eigenvalue exists):
 - **Saddle point**: if there is at least one positive eigenvalue and one negative.
 - **Inconclusive test**: otherwise.

1.2.3 Probability theory

- Can be seen as the extension of logic to deal with uncertainty.
- Logic provides a set of formal rules for determining what proposition are implied to be true.
 - Given the assumption that some other set of propositions is true or false.
- PT provides a set of formal rules for determining the likelihood of a proposition being true.

- Given the likelihood of other propositions.

1.2.3.1 Probability distributions

- Random variable:** a variable that can take on different values randomly.
 - Random variables may be **discrete** or **continuous**.
 - The set of all possible values taken by a random variable x is denoted Ω_x .

Discrete probability distributions

- Described using a **probability mass function** (PMF, $P(x)$):
 - It maps from a state of a random variable to the probability of that random variable taking on that state.
 - $x \sim P$: the random variable x follows the $P(x)$ distribution.
- prop Properties of a PMF (not to be confused with axioms of probabilities):
 - To be on a PMF on a random variable x , a function P must:
 - The domain of P must be the set of all possible states of x .
 - $\forall x \in x: 0 \leq P(x) \leq 1$.
 - $\sum_{x \in x} P(x) = 1$.
 - With S , S_1 and S_2 being sets of possible outcomes:
 - $P(S) = \sum_{x \in S} P(x)$ ($P(S)$ is a shorthand for $P(x \in S)$).
 - $P(S_1 \cup S_2) = P(S_1) + P(S_2) - P(S_1 \cap S_2)$.
 - $P(\Omega - S) = 1 - P(S)$.

Continuous probability distributions

- When working with continuous random variables, a **probability density function** (PDF, $p(x)$) is used.
- prop Properties of a PDF:
 - To be on a PDF on a random variable x , a function p must:
 - The domain of p must be the set of all possible states of x .
 - $\forall x \in x: p(x) \geq 0$ ($p(x) \leq 1$ is not required).
 - $\int p(x) dx = 1$.

Marginal probability

- The probability distribution over a subset of the set of variables.
- Marginal probabilities are computed by summing over all values of other variables.
- eg With discrete variables x and y with a joint distribution $P(x, y)$:
 - The marginal distribution $P(x)$ is $\forall x \in x: P(x) = \sum_y P(x = x, y = y)$.
- eg With continuous variables x and y with a joint distribution $p(x, y)$:
 - The marginal distribution $p(x)$ is $\int p(x, y) dy$.

1.2.3.2 Conditional probability

- Conditional probability:** $P(y = y | x = x) = \frac{P(y = y, x = x)}{P(x = x)}$.
- In many cases, it's useful to have the probability of some event given that some other has happened.
- def **Chain rule of Conditional probabilities:** $P(x^{(1)}, \dots, x^{(n)}) = P(x^{(1)}) \prod_{i=2}^n P(x^{(i)} | x^{(1)}, \dots, x^{(i-1)})$.
 - Any joint probability distribution over many random variables may be decomposed.
 - Into products of conditional distribution over only one variable.
- def **Bayes rule:** $P(x | y) = \frac{P(y | x)}{P(y)}$.
 - Used to find where $P(x | y)$ where $P(y | x)$ is known.

Independence

- Independent variables** ($x \perp y$): x and y if their probability distribution can be expressed as a product of two factors, one involving only x and one involving only y .
 - $\forall x \in x, y \in y: p(x = x, y = y) = p(x = x)p(y = y)$.
 - $x \perp y \iff \forall x \in x, y \in y: p(x | y) = p(x) \wedge p(y | x) = p(y)$.
- Conditional independent variables** ($x \perp y | z$, given a random variable z): x and y if the conditional probability over x and y factorizes in this way for every value of z .
 - $\forall x \in x, y \in y, z \in z: p(x = x, y = y | z = z) = p(x = x | z = z)p(y = y | z = z)$.

- x and y may not be independent per se, but they may be come it when another value is known.

1.2.3.3 Core statistical measures

- **Expectation (μ):**
 - The expected value of $f(x)$ with respect to $P(x)$ is the mean value that f takes on where x is drawn from P .
 - Discrete variables: $\mathbb{E}_{x \sim P}[f(x)] = \sum_x P(x)f(x)$.
 - Continuous variables: $\mathbb{E}_{x \sim P}[f(x)] = \int p(x)f(x) dx$.
 - prop **Linearity of the expectation operator**: $\mathbb{E}[\alpha f(x) + \beta g(x)] = \alpha \mathbb{E}[f(x)] + \beta \mathbb{E}[g(x)]$.
- def **Variance (σ^2)**: $Var[f(x)] = \mathbb{E}[(f(x) - \mathbb{E}[f(x)])^2] = \mathbb{E}[f(x)^2] - \mathbb{E}[f(x)]^2$.
 - It gives a measure of how much the values of a function of x vary as a different values of x from its probability distribution is sampled.
- **Standard deviation (σ)**: the square root of the variance.
 - For a normally distributed variable, about 95% of the points fall into the range $\mu \pm 2\sigma$.

Covariance

- def **Covariance**: $Cov(x, y) = \mathbb{E}_{x,y \sim P(x,y)}[(x - \mathbb{E}[x])(y - \mathbb{E}[y])] = \mathbb{E}_{x,y \sim P(x,y)}[(x - \mu_x)(y - \mu_y)]$.
- It gives a sense of how much two random variables are related to each other.
- It is affected by the scale of variables.
- def **Correlation**: $Corr(x, y) = \frac{Cov(x,y)}{\sigma_x \sigma_y}$.
 - Adjust the scale of each variable, ensuring that the relation between the variables is measured without being influenced by their individual magnitude.
 - A correlation close to ± 1 indicates a strong relationship between the variables.
 - A correlation close to 0 suggests that the variables may be independent.
- **Covariance** and **dependence** are related, but are in fact **distinct** concepts.
 - Two variables that are independent have zero covariance.
 - Two variables that have non-zero covariance are dependent.
 - Two variables can have zero covariance and be dependent nonetheless.

1.2.3.4 Common probability distribution

Bernoulli distribution

- **Bernoulli distribution**: a distribution over a single **binary** random variable.
 - Controlled by a single parameter $\phi \in [0, 1]$, which gives the probability of $x = 1$.
 - prop Properties of Bernoulli distributions:
 - $P(x = 1) = \phi$ and $P(x = 0) = 1 - \phi$.
 - $P(x = x) = \phi^x(1 - \phi)^{1-x}$.
 - $\mathbb{E}[x] = \phi$ and $Var(x) = \phi(1 - \phi)$.
- **Multinulli distribution** (or categorical): a distribution over a single discrete variable with k (finite) different states.
 - It is parametrized by a vector $\mathbf{p} \in [0, 1]^k$, with $\mathbf{1}^T \mathbf{p} = 1$.
 - Where p_i gives the probability of the i -th state.
 - Often used to refer to distributions over **categories** of objects.
 - The state 1 having numerical value 1, etc, is not assumed.
 - usually no expectation or variance is computed.

Binomial distribution

- It gives the probability of observing a given number of success in a **repeated Bernoulli experiment**.
- It is parametrized by:
 - p : the probability of success of the Bernoulli experiment.
 - N : the number of total repetitions of the Bernoulli experiment.
- If $x \sim Bi(p, N)$ then:
 - $P(x = k) = \binom{N}{k} p^k (1 - p)^{N-k}$ ($\binom{N}{k} = \frac{N!}{k!(N-k)!}$).
 - $\mathbb{E}[x] = Np$ and $Var[x] = Np(1 - p)$.

Gaussian (or normal) distribution

- Most commonly used distribution over real numbers.
- It is parametrized by mean μ and variance σ^2 .
- If $x \sim \mathcal{N}(\mu, \sigma^2)$ then:
 - $p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$.
 - $\frac{1}{\sqrt{2\pi\sigma^2}}$ to ensure that its integral will be 1.
 - It goes to 0 exponentially fast as x goes far from μ .
- Many distribution in the real world are truly close to being normal distribution.
 - **th Central limit theorem of PT:**
 - Let X_1, \dots, X_n be a sequence of iid random variable with finite mean μ and variance σ^2 .
 - Then the distribution of $S = \sum_{i=1}^n X_i$ approaches $\mathcal{N}(n\mu, n\sigma^2)$ as n approaches infinity.
 - The average $\frac{1}{n}S$ also approaches $\mathcal{N}(n\mu, n\sigma^2)$ as n approaches infinity.
- Out of all possible probability distribution with the same mean and variance, \mathcal{N} encodes the maximum amount of uncertainty.
 - It is the distribution having the **maximum entropy**.
 - With an unknown phenomenon, to assume that it behaves as \mathcal{N} takes the least amount of assumptions.
- \mathcal{N} generalizes to \mathbb{R}^n : $\mathcal{N}(\mathbf{x}, \mu, \Sigma) = \frac{1}{\sqrt{(2\pi)^n \det(\Sigma)}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu)\right)$.
 - μ : a vector denoting the mean of the distribution.
 - Σ : the covariance matrix of the distribution.
 - Covariance matrices are **symmetric** and **positive semi-definite**.
 - Their main diagonal contains variances.
 - If $\Sigma = \sigma^2 \mathbf{I}$ the variance is the same in every direction: the distribution is said to be **isotropic**.

Exponential and Laplace distributions

- def **Exponential distribution**: $p(x; \lambda) = \lambda \exp(-\lambda x)$, $x \geq 0$.
 - In DL, a distribution with a sharp point at $x = 0$ is needed.
- def **Laplace distribution**: $\text{Laplace}(x; \mu, \gamma) = \frac{1}{2\gamma} \exp\left(-\frac{|x-\mu|}{\gamma}\right)$.
 - Closely related to the exponential one.
 - It allows to place a sharp peak of probability mass at an arbitrary point μ .

Mixtures of distributions

- It's common to define probability distributions by combining other simpler distributions.
- One common way is to construct a **mixture distribution**: $P(\mathbf{x}) = \sum_i P(c=i)P(\mathbf{x} | c=i)$.
 - Where $p(c)$ is a multinomial distribution over the components of the mixture.
- **Latent variables**:
 - Random variables that **cannot be observed directly**.
 - The component identity variable c of the mixture model provides an example.
 - Latent variables may be related to \mathbf{x} through the joint distribution.
 - Then $P(\mathbf{x}, c) = P(\mathbf{x} | c)P(c)$.

1.2.3.5 Useful properties of common functions

- def **Sigmoid**: $\sigma(x) = \frac{1}{1+\exp(-x)}$.
 - Often used to produce the ϕ parameter of a Bernoulli distribution.
 - It **saturates** for large (negative/positive) values of x .
- def **Softplus**: $\zeta(x) = \log(1 + \exp(x))$.
 - A smoothed version of $x^+ = \max(0, x)$.

1.2.3.6 Continuous random variables

Measure theory

- A proper formal understanding of Continuous random variables and PDF requires **measure theory**.
 - Without it, one might encounter paradoxical situations.
 - eg To construct two sets S_1 and S_2 where $S_1 \cap S_2 = \emptyset$ such that $p(x \in S_1) + p(x \in S_2) > 1$.

- These paradoxes usually involve constructing very exotic sets, but the possibility exists.
- One of MT key contributions is that it provides a framework to characterize sets where probabilities can be consistently computed, thus avoiding paradoxes.
- **Negligibly small** set of points:
 - It has **measure zero** (eg a line in \mathbb{R}^2 has measure zero).
 - Any **union of countably many sets** having measure zero has measure zero.
 - So the set of all rational numbers has measure zero.
 - A property that **holds almost anywhere**:
 - When a property holds throughout all of space except for points in a set of measure zero.

Continuous variables that are deterministic functions of one another

- Two random variables x and y , such that $y = g(x)$, are given.
 - Where g is an invertible, continuous differentiable transformation.
 - Since g is invertible, $x = g^{-1}(y)$.
 - By inverting the transformation, the probability in the initial space should be computable.
 - Unfortunately: $p_y(y) \neq p_x(g^{-1}(y))$.
 - Since the **transformation changes the space**.
- This approach fails to account for the **distortion of space** introduced by g .
 - The probability of x lying in an infinitesimally small region with volume δx is given by $p(x)\delta x$.
 - Since g can expand or contract space, the infinitesimal volume surround x in x space may have different volume in y space.
- To correct the problem the following property must be preserved: $|p_y(g(x))dy| = |p_x(x)dx|$.
 - Which yields $p_y(g(x))|dy| = p_x(x)|dx| \implies p_x(x) = p_y(g(x))|\frac{d}{dx}g(x)|$.
 - Or equivalently, $p_y(y)|dy| = p_x(g^{-1}(y))|dx| \implies p_y(y) = p_x(g^{-1}(y))|\frac{d}{dy}g^{-1}(y)|$.
- In **higher dimension** $g: \mathbb{R}^m \rightarrow \mathbb{R}^n$ the derivative generalizes to the Jacobian matrix.
 - And the absolute value to the absolute value of the determinant.
 - $P_x(x) = p_y(g(x))|\det(J)|$.

1.2.3.7 Graphical models

- Often probability distribution can be split into many factors.
 - Like when a random variable influences another, and the latter influences another one, etc.
 - eg $P(a, b, c) = p(a)p(b | a)p(c | b)$ (the second expression is way simpler).
 - With 10 values for each variable, $10 + 10^2 + 10^2 = 210$ instead of $10^3 = 1000$.
 - These factorizations can greatly reduce the number of parameters needed to describe the distribution.
- Factorization over distributions can be visually described using **graphs** (directed or undirected).

Directed models

- **Directed models:** use graphs with directed edges.
- They represent factorizations into conditional probabilities distributions.
- One factor for every random variable x_i .
- An edge from a to b represent the dependency $P(b | a)$.
- **def Directed model factorization:** $p(\mathbf{x}) = \prod_i p(x_i | P_{a_g}(x_i))$.
 - Where $P_{a_g}(x_i)$ is the set of parents of x_i .
 - It consists of the conditional distribution of x_i given its parameters.

Undirected models

- **Undirected models:** use graphs with undirected edges.
- Variables influence each others.
- They represent factorizations using a set of functions.
 - Not necessary, nor common that they are probabilistic distributions.
- One factor $\phi^{(i)}$ per clique $C^{(i)}$ in the graph.
- **def Undirected models factorization:** $p(\mathbf{x}) = \frac{1}{Z} \prod_i \phi^{(i)}(C^{(i)})$.
 - Where $Z = \sum_{\mathbf{x} \in \mathcal{X}} \prod_i \phi^{(i)}(C^{(i)})$ (very problematic to compute).
 - i ranges between all the cliques (fully-connected graph subsets) in the graph.
 - The normalized product of all factors.

1.2.4 Information Theory

- Revolves around quantifying how much information is present in a signal.
- Useful applications of IT for DL:
 - To characterize probability distributions.
 - To quantify similarity between probability distributions.
 - eg Minimize the distance to a goal distribution (the one that generate Tr data).
- **Quantifying information** (main intuition behind IT):
 - The quantity of information carried by a message depends on **how likely it is**.
 - Learning that an **unlikely event** has occurred is **more informative**.
 - The **less probable** an even is, the more surprising it is.
 - The more surprising it is the **more information** it yields.
 - An event with probability 100% is **perfectly unsurprising and yields no information**.
 - **Independent events** should have **additive information**.

1.2.4.1 Shannon's Entropy Measure

- def **Shannon's Entropy Measure**: $H(x) = \mathbb{E}_{x \sim P}[I(x)] = -\sum_x P(x) \log P(x)$.
- It captures the average amount of *information* across all possible outcomes of a random variable.
- \sum if the probability is discrete, otherwise \int .
- def **Self-information**: $I(x) = -\log P(x)$.
 - Used to quantify the uncertainty of an event.
 - $-\log$ used to guarantee that the more surprising an event is the **more information** it yields.
- th **Shannon's Source Coding Theorem**: $H(x)$ provides a lower bound for the average length of code-words in an optimal encoding of the possible values of x .
 - An optimal encoding for messages (with different probabilities) is needed.
 - Ideal: if a message is very frequent, a smaller code is assign to it.
 - Since it's very frequent, a smaller code is better to minimize communication.
- The entropy of a random variable $x \sim \text{Bernoulli}(\phi)$ as ϕ varies from 0 to 1.
 - 0 nats (**natural unit of information**) at $\phi = 0$ and $\phi = 1$, 0.7 nats at $\phi = 0.5$ (upward parabola).
 - The quantity of information send with a message with $\phi = 0$ or $\phi = 1$ is 0 (no information gained).
 - The entropy of a distribution is maximized for a uniform distribution.

1.2.4.2 Kullback-Leibler divergence

- def **Kullback-Leibler divergence**: $D_{KL}(P||Q) = \mathbb{E}_{x \sim P}[\log \frac{P(x)}{Q(x)}] = \mathbb{E}_{x \sim P}[\log P(x) - \log Q(x)]$.
- $P(x)$ and $Q(x)$ are two separate distribution over the same random variable x .
- The KL is used to measure how much different are these distributions.
 - If two distributions are the same, there is no information loss $\Rightarrow D_{KL} = 0$.
 - If two distributions are very dissimilar, there will be a substantial information loss.
- In the case of discrete variables:
 - It is the **extra amount of information** needed to send a message containing symbols drawn from P .
 - When it is used a code that was designed to minimize the length of message drawn from Q .
- prop Properties of the KL divergence:
 - KL divergence is always **non-negative**.
 - KL divergence is 0 $\iff P$ and Q are the same distribution.
 - Or are equal *almost anywhere* in the case of continuous variables.
 - **KL is not a distance**.
 - A distance should also be symmetric and satisfy the triangle inequality.
 - The fact that KL is not symmetric has important consequences when the distance between P and Q has to be minimized.
 - Most of the times, minimizing $D_{KL}(p||q)$ yields different output than minimizing $D_{KL}(q||p)$.
- eg An optimization problem is used to approximate a bimodal distribution with a unimodal gaussian distribution.
 - By optimizing $\min_q [D_{KL}(q||p)] = \min_q [\mathbb{E}_{x \sim q}[\log \frac{q(x)}{p(x)}]]$:
 - p is given and can't be modified, instead q is *tuned* to approximate the target distribution.
 - The result is a narrow distribution that approximate only one of the two modes.
 - By optimizing $\min_q [D_{KL}(p||q)] = \min_q [\mathbb{E}_{x \sim p}[\log \frac{p(x)}{q(x)}]]$:
 - The result is a shallow distribution that approximate both of the two modes, but in a more inaccurate way.
 - It is possible to approximate by composing the best parts of both results.

- By just chaning the order of KL argument, there are two very different solution to the optimization problem.

1.2.4.3 Cross-Entropy

- def **Cross-Entropy**: $H(P, Q) = H(P) + D_{KL}(P\|Q) = -\mathbb{E}_{x \sim P}[\log Q(x)]$.
- The average number of bits needed to encode message for code Q with a code designed for P .
- Similar to $D_{KL}(P\|Q) = \mathbb{E}_{x \sim P}[\log P(x) - \log Q(x)]$.
 - Similar expression, but it lacks the term $\log P(x)$.
 - D_{KL} measures the expected **extra** number of bits.
 - While H measures the total number of bits.
- Minimizing $H(P, Q)$ wrt Q is the same as minimizing $D_{KL}(P\|Q)$.
 - $\min_q D_{KL}(P\|Q) = \min_q \mathbb{E}_{x \sim P}[\log P(x)] - \mathbb{E}_{x \sim P}[\log Q(x)]$.
 - Since minimization is wrt Q , the first element term doesn't contribute to anything.
 - The resulting term is exactly the cross-entropy.

1.3 Introduction to Neural networks

1.3.1 Perceptrons

- Introduced by Frank Rosenblatt in 1958.
 - Presented as a pattern recognition device (although is used for more general problems).
 - It mimics a part of what is known of the mammalian visual system.
- Simple form of neural network used to classify **linearly separable** examples.
- Basic structure of a perceptron:
 - *Input:* a vector \mathbf{x} (eg pixel values, phonemes, outputs from previous neurons, etc).
 - *Output:* a vector \mathbf{y} (single value if only a neuron).
 - *Weight vector:* \mathbf{w} .
 - Training set: set of tuples (input, desired output).
 - A **pattern** is an element of the training set.
- **Biological neural networks:**
 - In the brain the basic computational unit is the **neuron**.
 - A neuron collect, elaborate and propagate electric signals.
 - The output electric signal is proportional to the received input.
 - Neurons collect and transmit informations **in parallel** (as in ANN).
 - There are 10 billion neurons and 60 millions of millions of millions of synapses.
 - In the same way, the basic computational unit of NN is the neuron.
 - Several neurons are connected to each other by synapses.
 - Very simple operations: deciding if the **total input is higher than a threshold**.

1.3.1.1 McCulloch & Pitts' model of a neuron [1943]

- def **Network input** (*netinput*) to j : $\mathbf{x} \cdot \mathbf{w} = \sum_{i=1}^p x_i w_i$.
 - For each input value x_i what reaches the neuron j is $w_{ji}x_i$.
 - w_{ji} is the weight of the synapse from i to j .
 - A synapse can be **excitatory** or **inhibitory**, therefore the weight can be $<, >, = 0$.
- The output of the neuron will depend on the netinput and the bias.
 - The **bias** represents the natural predisposition of the neuron to activate.
 - If the bias < 0 , for the neuron to be activated, incoming input values must be high.
 - If the bias > 0 , the neuron is often activated, also with low inputs.
 - To simplify the computation, the bias is represented as an extra input element $x_0 = 1$ with $w_0 = b$.
 - In this case, $v + j = \sum_{j=0}^p w_{ji}x_i$.
 - def **Perceptron output**: $y_j = \varphi(u_j + b)$.
 - φ is the **activation function** (eg sigmoid).
 - **Perceptron activation function:**
 - 1 for all input vectors \mathbf{x} s.t. $\mathbf{x} \cdot \mathbf{w} > 0$.
 - -1 for all input vectors \mathbf{x} s.t. $\mathbf{x} \cdot \mathbf{w} \leq 0$.
 - $u_j + b$ is the **activation potential** or local field of j .

1.3.1.2 Perceptron Learning

- The **decision boundary** is made by all \mathbf{x} s.t. $\mathbf{x} \cdot \mathbf{w} = 0$.
 - In 3-D with $w_0 = 0, w_1x_1 + w_2x_2 = 0$ is the equation of the DB line.
- The perceptron **learn** to classify examples linearly separable by **adjusting the weights**.

```

n=1;
initialize w(n) randomly;
while (there are misclassified training examples)
    Select a misclassified augmented example (x(n),d(n))
    if(d(n) = 1), w(n+1) = w(n) + ηx(n);
    if(d(n) = -1), w(n+1) = w(n) - ηx(n); → w(n+1) = w(n) + ηd(n)x(n);
    n = n+1;
end-while;
  
```

Figure 1: Perceptron Learning Algorithm.

Algorithm: Perceptron Learning

- Adjust the weights to obtain the desired classification.
 - The activation function cannot be modified, only the weights.
 - Learning based on **correction of the error** (for each misclassified pattern).
- **Weights update:**
 - Updates are performed **pattern by pattern** (and not by **epoch**, as in other neural networks).
 - An **epoch** is a iteration over all elements of the training set.
 - Instead, here weights are updated for each pattern misclassified.
 - Correct output: $w(n+1) = w(n)$.
 - Incorrect output:
 - $w(n+1) = w(n) - \eta(n) \times (s)$ if $x(n) \cdot w(n) > 0$ and $x(n) \in C_2$ (output too high).
 - $w(n+1) = w(n) + \eta(n) \times (s)$ if $x(n) \cdot w(n) \leq 0$ and $x(n) \in C_1$ (output too low).
 - η is the **learning rate** (with high η updates will be abrupt).
- **th Convergence theorem:** if the problem is **linearly separable** the learning algorithm will find a solution.
 - At each iteration, the weight vector is modified and, as a consequence, the decision boundary.
 - This theorem guarantees that **weight adjustment terminates**.
 - The proof is based on a geometric interpretation about how the weights vector is updated.
 - A: $\|w(k+1)\|^2 \geq \frac{k^2\alpha^2}{\|w^*\|^2}$ (lower bound).
 - B: $\|w(k+1)\|^2 \leq k\beta$ (upper bound).
 - A and B are compatible only if $\frac{k^2\alpha^2}{\|w^*\|^2} \leq k\beta \rightarrow k \leq \frac{\beta\|w^*\|^2}{\alpha^2}$.
 - k can't be larger than this **finite** quantity, the algorithm **terminates** (QED).
- It converges for **linearly separable problems**.
 - But simple (eg XOR) and non-linearly solvable problems cannot be solved [Minsky & Papert, 1969].
 - **To test if a dataset is linearly separable:**
 - The model should be trained on the whole dataset (train and test set).
 - Its training accuracy should be 100% (if not, then the dataset is not linearly separable).

1.3.1.3 Multilayer networks

- To solve more complex problems **hidden units** are introduced.
 - Each unit solves part of the problem, then solutions are combined.
 - But a different learning algorithm (**backpropagation**) is introduced.
- With enough hidden units, the network can **discriminate a convex zone**.

1.3.2 Multilayer neural networks

- In ML, a task is approached by trying to create a system that can learn from examples.
- By combining several perceptrons into a **complex networks**, a very large quantity of function can be modelled.
 - A perceptron can easily represent a NAND gate (eg with $x_1, x_2 \in \{0, 1\}, (-2, -2) \cdot x + 3$).
 - Any other logic operator can be built in terms of NAND operators only.
 - Any NAND-only complex circuit (eg binary adder with carry) can be translated into a perceptron NN.
- Networks as computational device:
 - Perceptron networks can perform all sorts of computation.
 - Algorithms that **learns those computations** can be made.

1.3.2.1 Universal Approximation Theorem

th Universal Approximation Theorem: a **feedforward network** with a **linear output layer** and at least one hidden layer with any **squashing activation function** (eg logistic sigmoid) can **approximate any Borel measurable function** from one finite-dimensional space to another with **any desired non-zero amount of error**, provided that the network is given **enough hidden units**. The derivatives of the feedforward network can also approximate the derivatives of the function arbitrarily well.

- A **linear output layer** computes simply $w \cdot x + b$ (no > 0).
- Any continuous function on a closed and bounded subset of \mathbb{R}^n is **Borel measurable**.
 - And therefore may be approximated by a neural network.
- With a big enough NN, any arbitrarily-complex practical problem can be approximated.
 - Empirical observations proved that with enough data, any problem of overfitting can be overcome.

- Unfortunately, in the worst case, an **exponential number of hidden units** may be required.
 - Possibly with one hidden unit for each input configuration that needs to be distinguished.
 - This is easiest to see in the binary case:
 - The number of possible binary functions on vectors $\mathbf{v} \in \{0, 1\}$ is 2^{2^n} .
 - Selecting one such function requires 2^n bits, which will in general require $O(2^n)$ degrees of freedom.

1.3.2.2 Sigmoid neurons

- It is hard to learn anything useful using the perceptrons.
- *Desiderata:* small changes in any weight (or bias) causes a small changes in the output.
 - $\mathbf{w} + \Delta\mathbf{w} \rightarrow \text{output} + \Delta\text{output}$.
- def **Sigmoid neuron:** a neuron where the output is computed as $\sigma(\mathbf{w} \cdot \mathbf{x} + b)$, where $\sigma(z) = \frac{1}{1+e^{-z}}$.
 - The output approximate the Step function when $\mathbf{w} \cdot \mathbf{x} + b$ is very large or small.
 - But it does not change abruptly in proximity of 0.
 - Since the output is in $[0, 1]$, each internal node input ranges in $[0, 1]$ as well.
- The information relating to changes can be used to **updates the weights**.
 - A Step function doesn't provide any meaningful information about those changes.
 - When the activation function has a larger slope, it's easier to **guide the learning**.
- It is important that the **activation function** is **smooth** in order to control how varies Δoutput .
 - Calculus implies that $\Delta\text{output} \approx \sum_j \frac{\partial\text{output}}{\partial w_j} \Delta w_j + \frac{\partial\text{output}}{\partial b} \Delta b$.
 - The *output* must be **differentiable w.r.t. the weights and the biases**.
 - The Step function is not differentiable since it's not continuous (it also has derivatives = 0, bad).
 - It is not the exact shape of σ to be crucial, but rather its smoothness.
 - In other situations, the sigmoid activation function is not the best choice.
 - Everything in NN is **situational**, and to develop a proper model is *more an art than a science*.

1.3.2.3 Architecture of Neural networks

- A typical neural network is composed by: an **input layer**, **1-n hidden layers**, an **output layer**.
 - This architecture is called **multilayer perceptrons** (despite not being built out of perceptrons).
 - The input layer is not a proper *layer*, it only represents the inputs.
 - While all other nodes in other layers are sigmoid neurons.
 - In the output vector, an output neuron can be used to represent each of the possible labels.
 - The final output will be the output unit with the **largest activation**.
 - In a **feedforward network**, a node is connected only to *successive* nodes in the following layer.
 - In a **fully connected FN**, a node is connected to **every** nodes in the following layer.
- The design of input and output layers is application-dependent.
 - eg An output neuron can be used to represent each one of the possible labels.
 - The final output is pick up as the output unit with the largest activation.
 - eg Each element of the output vector can be modeled as an output neuron.
- Hidden layers design is not as straight forward as the one for input or output layers.
 - NN researchers have developed over time many design heuristics for it.
 - eg Sharing weights between neurons allow to learn a single feature detector and apply it in many different locations.
 - eg Stacking layers may allow to learn complex features from simpler ones.
 - eg Using a large number of hidden units is powerful but it's prone to overfitting or require more data.

1.3.3 Gradient descent

- A learning algorithm to find the weights for a NN.
 - A **cost function** to measure how well a network is doing is defined.
 - $C(\mathbf{w}, \mathbf{b})$ measures how well the network performs on *Tr* data (\mathbf{X}, \mathbf{y}) .
 - A way to **minimize this cost function** needs to be found.
 - To solve the optimization problem minimize _{\mathbf{w}, \mathbf{b}} $C(\mathbf{w}, \mathbf{b})$.
 - def **Cost function:** $C(\mathbf{w}, \mathbf{b}) = \frac{1}{2n} \sum_{\mathbf{x}} \|y(\mathbf{x}) - \mathbf{a}(\mathbf{x})\|^2$.
 - \mathbf{a} is the vector of outputs with \mathbf{x} has an input.
 - It depends on \mathbf{w}, \mathbf{b} and \mathbf{a} , more properly denoted as $\mathbf{a}_{\mathbf{w}, \mathbf{b}}(\mathbf{x})$.
 - Accuracy is a natural evaluation measure, but it is **not smooth**.

1.3.3.1 Gradient descent derivation

! GD is an **iterative optimization algorithm** that at each step move the current position in the **opposite direction** w.r.t. the gradient, $\mathbf{v}' \leftarrow \mathbf{v} - \eta \nabla C$. The error function C is descended following the **direction of the steepest descent**.

- Consider a function $C(\mathbf{v})$ defined as $C : \mathbb{R}^n \rightarrow \mathbb{R}$.
- In general it might be very difficult to pinpoint the global minimum.
 - Using Calculus to find **closed form solution** is not feasible.
- The variables are moved by a little $\Delta \mathbf{v} = (\Delta v_1, \Delta v_2)^T$:
 - The vector C changes as $\Delta C \approx \frac{\partial C}{\partial v_1} \Delta v_1 + \frac{\partial C}{\partial v_2} \Delta v_2 = \nabla C \cdot \Delta v$.
 - $-\nabla C$ is the direction where C decrease the most.
 - $-\nabla C \cdot \nabla C = -\|\nabla C\|_2^2$ (the norm is a positive number).
 - The Taylor expansion is **truncated at the first derivative to use a plane** to approximate the problem.
 - Then the variables are moved along this plane.
- By choosing $\Delta \mathbf{v} = -\eta \nabla C$, it is guaranteed to make the **cost function to decrease**.
 - The step η , called **learning rate**, should be not big nor too small (typically of the order 0.1-0.001).
- The GD is applied until the found derivatives are very small.
- **Second order approximation**:
 - The Taylor expansion can be truncated not at the first but at the second derivative.
 - But with n parameters, there are n^2 derivatives (not feasible with large n).
 - Methods to mitigate this have been developed, but GD still remains the most common implementation.
- **Applying Gradient descent in a NN**:
 - Applying GD to weights and biases allows for training a NN.
 - Start from a fixed point and iterate through the data updating at each step the weights as:
 - $w'_k \leftarrow w_k - \eta \frac{\partial C}{\partial w_k}$.
 - $b'_l \leftarrow b_l - \eta \frac{\partial C}{\partial b_l}$.

1.3.3.2 Stochastic gradient descent (SGD)

- Consider the cost function $C = \frac{1}{n} \sum_{\mathbf{x}} C_{\mathbf{x}}$ (generic averaged CF).
 - The gradient of C is then $\nabla C = \frac{1}{n} \sum_{\mathbf{x}} \nabla C_{\mathbf{x}}$ (∇ is a **linear operator**).
 - To compute the gradient ∇C :
 - The gradients $\nabla C_{\mathbf{x}}$ has to be computed separately for each training input \mathbf{x} .
 - And then average them.
 - When the dataset is very big, this becomes too slow.
- In **mini-batch SGD**, a **small subsample** of the dataset is used to approximate the gradient over the entire dataset.
 - Let x_1, \dots, x_m be m randomly chosen training inputs.
 - If m is large enough, it should be apparent that $\sum_{j=1}^m \frac{\nabla C_{x_j}}{m} \approx \frac{\sum_{\mathbf{x}} \nabla C_{\mathbf{x}}}{n} = \nabla C$.
- Procedure:
 - A mini-batch is extracted and the weights are updated using the following rules:
 - $w'_k \leftarrow w_k - \frac{\eta}{m} \frac{\partial C_{x_j}}{\partial w_k}$.
 - $b'_l \leftarrow b_l - \frac{\eta}{m} \frac{\partial C_{x_j}}{\partial b_l}$.
 - Then a **new mini-batch** is randomly chosen and the model is trained with those.
 - When all examples have been used an **epoch of training** is completed.
 - At that point a new training epoch is started over.
- The estimate won't be perfect (there will be statistical fluctuations), but it isn't needed.
 - The point is to **move in a general direction** that will help decrease C .

1.3.4 Backpropagation

- Gradient descent is a general procedure for optimizing differentiable functions.
 - Backpropagation is its instantiation in the context of NN.

1.3.4.1 Mathematical preamble to backpropagation

- Notation:

- w_{jk}^l : the weight from the k -th neuron in level $l - 1$ to the j -th neuron in the level l .
- b_j^l : the bias of the j -th neuron in the l -th layer.
- z_j^l : the activation value of the j -th neuron in the l -th layer.
- **Activation:** $z_j^l = \sigma(\sum_k w_{jk}^l z_k^{l-1} + b_j^l)$.
 - The sum is on all neurons k in the $(l - 1)$ -th layer.
- \mathbf{W}^l : the matrix having element w_{jk}^l at row j and column k .
 - The x -th row of \mathbf{W}^l corresponds to the weights entering the x -th neuron in the l layer.
- \mathbf{b}^l : the column vector having b_j^l as its j -th element.
- \mathbf{z}^l : the column vector having z_j^l as its j -th element.
- $f(\mathbf{v})$ is the element-wise application of f to a matrix (output: a matrix).
- **def Activation:** $\mathbf{z}^l = \sigma(\mathbf{W}^l \mathbf{z}^{l-1} + \mathbf{b}^l)$.
 - Faster to compute thanks to matrix computation libraries.
 - The **weighting input** to a neuron as level l is: $\mathbf{a}^l \equiv \mathbf{W}^l \mathbf{z}^{l-1} + \mathbf{b}^l$.
 - The component j of a vector \mathbf{a}^l is $a_j^l = \sum_k w_{jk}^l z_k^{l-1} + b_j^l$.
- Working assumptions about the **cost function**:
 - Can be written as an average $C = \frac{1}{n} \sum_x C_x$ over C_x for individual Tr examples.
 - It is a function of the outputs of the neural network (**self-contained**).
 - eg Those assumptions are satisfied by the quadratic cost function.
- **Hadamard Product** (\odot): the element-wise product of two vectors.

- **(BP1):** $\delta^L = \nabla_{\mathbf{z}} C \odot \sigma'(\mathbf{a}^L)$
- **(BP2):** $\delta^l = ((\mathbf{W}^{l+1})^T \delta^{l+1}) \odot \sigma'(\mathbf{a}^l)$
- **(BP3):** $\frac{\partial C}{\partial b_j^l} = \delta_j^l$
- **(BP4):** $\frac{\partial C}{\partial w_{jk}^l} = z_k^{l-1} \delta_j^l$

Figure 2: Backpropagation four fundamental equations.

1.3.4.2 Backpropagation four fundamental equations

- The **gradient of the cost function** w.r.t. weights and biases is $\nabla_{\mathbf{w}, \mathbf{b}} C = [\frac{\partial C}{\partial w_{jk}^l}, \dots, \frac{\partial C}{\partial b_j^l}, \dots]$.
- **def Neuron error:** $\delta_j^l = \frac{\partial C}{\partial a_j^l}$.
 - The error at level l and neuron j .
 - BP provides a procedure to compute it for each neuron in each layer.
 - And then will relate δ_j^l to the quantities of real interest.
 - Measures how much the CF varies when the inputs of the neurons at that level are slightly perturbed.
 - The term *error* derives from the idea of introducing a *small* error in the inputs of the layer.
 - And then observe how it propagates to the cost function.

First equation of BP

- **def First equation of BP (BP1):** $\delta_j^L = \frac{\partial C}{\partial z_j^L} \sigma'(a_j^L)$.
- Specifies how to compute the *error* δ^L at the output layer (L).
- How much the cost function C changes when there is a small change in the input of σ .
- **Easily computable:**
 - a_j^L is computed by letting the inputs to flow through the network up to neuron j the level L .
 - The exact form of $\frac{\partial C}{\partial z_j^L}$ will depend on the form of the cost function.
 - eg For the quadratic cost: $C = \frac{1}{2} \sum_j (y_j - z_j^L)^2 \implies \frac{\partial C}{\partial z_j^L} = (z_j^L - y_j)$.
- **First equation of BP (BP1):** $\delta^L = \nabla_{z^L} C \odot \sigma'(\mathbf{a}^L)$.

- The gradient is by definition the collection of the partial derivatives.

Second equation of BP

- def **Second equation of BP** (BP2): $\delta^l = ((\mathbf{W}^{l+1})^T \delta^{l+1}) \odot \sigma'(\mathbf{a}^l)$.
- Allows to compute the error for non-output layer.
- Multiplication between:
 - $\sigma'(\mathbf{a}^l)$: measures how fast the output at this level varies in response to a variation of its input.
 - $(\mathbf{W}^{l+1})^T$: measures how that changes propagate through the NN via the connecting weights.
 - δ^{l+1} : how fast that level will change in response to changes in its input.
- Ignore anything before l and after $l + 1$ (already accounted by δ^{l+1}).
- **Easily computable:**
 - The only non-trivial term, δ^{l+1} can be computed via BP1 applied to the last layer.
 - And then apply BP2 to the other layers up to the current layer.
 - Information is computed from the back (output) and then **propagated backwards**.

Third and fourth equations of BP

- Those equations are the ones directly related to **applying gradient descent to a NN**.
 - How to compute the cost function given the biases and weights terms.
- def **Third equation of BP** (BP3): $\frac{\partial C}{\partial b_j^l} = \delta_j^l$.
 - That is, the error δ_j^l is *exactly equal* to the rate of change $\frac{\partial C}{\partial b_j^l}$.
- def **Fourth equation of BP** (BP4): $\frac{\partial C}{\partial w_{jk}^l} = z_k^{l-1} \delta_j^l = z_{in} \delta_{out}$.
 - How to compute the rate of change of the cost w.r.t. any of the weights in the network.
 - It is intended that the weight w.r.t. of which the derivative is taken determines the levels and the indices of z_{in} and δ_{out} .

Consequences of the four equations of BP

- Considering BP4:
 - Whenever z_{in} is small, the gradient will also tend to be small and the learning rate will be slow.
 - Weights originating from **low activation** units will **evolve slowly**.
- Considering BP1:
 - The sigmoid derivative is almost zero when its argument is either very large or very small.
 - The learning will proceed slowly for an output neuron if it's either *low activation* (≈ 0) or *high activation* (≈ 1).
 - When the derivative of the activation function is almost zero, a neuron is **saturated**.
- Same considerations can be made about the implications of BP2.
 - If the neuron saturates the learning will be slow.
- These observations do not rely on the activation function being the sigmoid.
 - They do **apply to any activation function**.
- These observations can be applied to design novel activation functions.
 - eg An activation function σ so that σ' is always positive (to avoid saturation).

1.3.4.3 Algorithm: Gradient descent + Backpropagation

- Weights and biases are initialized randomly.
 - And then (hopefully) they will converge to useful values.
- Backpropagation:
 - *Input*: x will be used to set the corresponding activation \mathbf{z}^1 for the input layer.
 - *Output*: the gradient of the cost function given by $\frac{\partial C}{\partial w_{jk}^l} = z_k^{l-1} \delta_j^l$ and $\frac{\partial C}{\partial b_j^l} = \delta_j^l$.
- Backpropagation provides only the gradient of the error committed on a fixed input.
 - To get (S)GD, iterate over all examples and average the results.
- eg $\mathbf{a}^{x,l}$: a computed at the level l on the example x .
- To fully implement SGD, one needs:
 - An additional loop that iterates over the **minibatches**;
 - An additional loop to iterate through the **epochs**.

Input a minibatch of training examples of length m .

1. **For each training example \mathbf{x} :** Set the corresponding input activation $\mathbf{z}^{\mathbf{x},1}$, and perform the following steps:
 - **Feedforward:** For each $l=2,3,\dots,L$ compute $\mathbf{a}^{\mathbf{x},l} = \mathbf{W}^l \mathbf{z}^{\mathbf{x},l-1} + \mathbf{b}^l$ and $\mathbf{z}^{\mathbf{x},l} = \sigma(\mathbf{a}^{\mathbf{x},l})$.
 - **Output error $\delta^{\mathbf{x},L}$:** Compute the vector $\delta^{\mathbf{x},L} = \nabla_{\mathbf{z}} C_{\mathbf{x}} \odot \sigma'(\mathbf{a}^{\mathbf{x},L})$.
 - **Backpropagate the error:** For each $l=L-1, L-2, \dots, 2$ compute $\delta^{\mathbf{x},l} = ((\mathbf{W}^{l+1})^T \delta^{\mathbf{x},l+1}) \odot \sigma'(\mathbf{a}^{\mathbf{x},l})$.
2. **Gradient Descent:** For each $l=L, L-1, \dots, 2$ update the weights according to the rule $\mathbf{W}^l \leftarrow \mathbf{W}^l - \frac{\eta}{m} \sum_{\mathbf{x}} \delta^{\mathbf{x},l} (\mathbf{z}^{\mathbf{x},l-1})^T$, and the biases according to the rule $\mathbf{b}^l \leftarrow \mathbf{b}^l - \frac{\eta}{m} \sum_{\mathbf{x}} \delta^{\mathbf{x},l}$.

Figure 3: Gradient descent + Backpropagation algorithm.

Algorithm efficiency

- To understand how efficient this algorithm it has to be considered:
 - A numerical approximation of the derivative based on computing $\frac{\partial C}{\partial w_j} \approx \frac{C(w+\epsilon e_j) - C(w)}{\epsilon}$:
 - Would require to compute the approximation:
 - For each possible weight (they can be millions).
 - Each time performing a forward pass to compute C with the new weights.
 - The computation of the exact derivative performed by BP requires a **single forward pass + a single backward pass**.
- It is trivial to exploit **dynamic programming** to avoid exponential computation.
 - The influence of nodes on later layers onto nodes of the l -th layers is totally captured by δ^{l+1} vector.
 - If this was not the case, to compute the derivative of the cost w.r.t. a weight w_{jk}^l one would need:
 - To integrate the effects over the exponential # of paths that connect that node k at level $l-1$ to the output node.

1.3.5 Best practices in neural networks

1.3.5.1 Slow learning

- Networks based on sigmoid units with a quadratic loss function often learn **very slowly**.
 - When a sigmoid is combined with a squared loss, GD requires to compute the CF derivative wrt the weights.
 - The obtained expression is the sigmoid derivative (flat in most of its domain).
 - Therefore the derivative is almost 0 in most of its domain (**slow learning**).

Cross-entropy

- The **cross-entropy** CF solves this problem without changing the activation function.
- σ will be still defined to be the sigmoid function.
- **def Cross-entropy:** $C = -\frac{1}{n} \sum_{\mathbf{x}} [y \ln z + (1-y) \ln(1-z)]$.
 - Where y is the desired (binary) output and $z \in [0, 1]$ is the activation of the neuron.
 - A single neuron with multiple inputs is considered.
 - It still behaves as a cost function:
 - If $y = 1$ and $z \approx 1$, then $-(y \ln z + (1-y) \ln(1-z)) \approx 0$.
 - If $y = 1$ and $z \approx 0$, then $-(y \ln z + (1-y) \ln(1-z)) \approx \infty$.
 - The same goes for $y = 0$ and $z \approx 0$ and $y = 0$ and $z \approx 1$.
 - Always positive and tend to 0 as the computed outputs tend to the desired outputs.
- The sigmoid derivative is $\sigma'(a) = \frac{d}{da} \left(\frac{1}{1+e^{-a}} \right) = \frac{e^{-a}}{(1+e^{-a})^2} = \sigma(a)(1-\sigma(a))$.
 - Which yields $\frac{\partial C}{\partial w_j} = \frac{1}{n} \sum_{\mathbf{x}} x_j (\sigma(a) - y)$.
 - The rate of learning depends only on how well the output unit is approximating the desired output.
- **Learning fast when in wrong and learning slow when correct** (desirable).
 - And by just changing the CF and not the activation function.
- **def Multi-layer Cross-entropy:** $C = -\frac{1}{n} \sum_{\mathbf{x}} \sum_j [y_j \ln z_j^L + (1-y_j) \ln(1-z_j^L)]$.
 - Generalization of CE to many-neurons, multi-layers NN.

- $y = y_1, y_t, \dots$ are the desired values at the output neurons.
- Each neuron is treated as a binary problem where the usual CE is applied.
 - And then all these cross-entropies are summed up.
- For numerical reasons, ML libraries doesn't include the sigmoid in z_j^L .

Soft-Max and Log-Likelihood

- When the task involves the prediction of a probability distribution, Soft-Max activation is useful.
 - It shares the benefit of using a cross-entropy.
- def **Soft-Max activation**: $z_j^L = \frac{e^{a_j^L}}{\sum_k e^{a_k^L}}$.
 - \sum_k considers the other neurons in the layer.
 - By using the exponential, the highest number get the most of the weight.
 - Usually the maximal number results in almost 1, while the others are much below.
 - A soft version of evaluating the *maximum* of a function.
 - $\sum_j z_j^L = \sum_j \frac{e^{a_j^L}}{\sum_k e^{a_k^L}} = \frac{\sum_j e^{a_j^L}}{\sum_k e^{a_k^L}} = 1$ (required in a probability distribution).
 - It guarantees that each activation $z_j^L \in (0, 1)$ (not included) and $\sum_j z_j^L = 1$.
 - The collective behavior of the output units can be seen predicting a distribution of probabilities.
 - A distribution of probabilities Where $z_j^L = P(j | x)$.
- def **Log-likelihood cost function**: $C \equiv -\ln z_y^L$.
 - y is not the one-hot encoding representing the correct class, y is the number correct class.
 - z_y^L is the output of the **neuron corresponding to the correct class**.
 - It behaves as a cost function:
 - If the network compute the correct output, then $z_y^L \approx 1$ and $C = -\ln z_y^L \approx 0$.
 - If the network compute the *wrong* output, then $z_y^L \approx 0$ and C will assume a large value.
 - Only a single neuron is considered since the others are already **taken into account in the Soft-Max**.
- Using Soft-Max + Log-Likelihood architecture **does not suffer of slow learning**.
 - $\frac{\partial C}{\partial w_{jk}^L} = z_k^{L-1}(z_j^L - y_j)$ (it depends on how far the output is from the desired result).
 - $\frac{\partial C}{\partial b_k} = z_j^L - y_j$.

1.3.5.2 Parameter initialization

- It can have a significant impact on generalization performance and convergence speed.
 - Unfortunately there is very little theory suggesting how to properly initialize the parameters.
 - Initializing them to large values is not ideal according to seen activation functions.
- Main approaches involve:
 - Symmetry breaking:
 - Keeping the variance of the parameters constant.

Symmetry breaking

- By initializing all parameters to the same value, all neurons will compute the same function.
- A common approach is to **initialize the parameters randomly**.
 - Using either a uniform distribution in the range $[-\epsilon, +\epsilon]$ or using $\mathcal{N}(0, \epsilon^2)$.
- The choice of ϵ is very important too.
 - Too big or too small weight will contribute to very large (**gradient explosion**) or very small gradients (**gradient implosion**).

He initialization

- def **He initialization**: $\epsilon = \sqrt{\frac{2}{M}}$ (where M is the number of inputs of the neuron being initialized).
- A few attempts have been made to find good initialization values.
 - For ReLU activation units the *He initialization* should make the gradient approximately equal to 1.
- The main idea is to **keep the variance constant between layers**.
 - *Assumptions*:
 - Each layer l of the network evaluates to: $a_i^l = \sum_{j=1}^M w_{ij} z_j^{l-1}$ and $z_j^l = \text{ReLU}(a_i^l)$.
 - The weights are initialized drawing from a Gaussian $\mathcal{N}(0, \epsilon^2)$.

- The outputs of units at layer $l - 1$ have zero mean and variance λ^2 .
- It can be shown that $\text{var}[z_j^l] = \frac{M}{2}e^2\lambda^2$.
- Then to keep the variance constant between layers, $\epsilon = \sqrt{\frac{2}{M}}$.

Xavier initialization

- **def Xavier initialization:** $\epsilon = \frac{1}{\sqrt{M}}$.
- Designed for activation functions **symmetric** around 0 (sigmoid and tanh).
- It samples initialization weights from $U[-\epsilon, \epsilon]$.
- **def Normalized Xavier initialization:** $\epsilon = \frac{6}{\sqrt{N+M}}$.
 - Where N is the number of units in the layer.

1.3.5.3 Convergence

- Interesting facts can be proven about the convergence of GD method.
- The components of w evolve independently.
 - After T steps, it can be shown that $\alpha^{(T)} = (1 - \eta\lambda_i)^T \alpha_i^{(0)}$.
 - Provided that $|1 - \eta\lambda_i| < 1$, the limit as $T \rightarrow \infty$ leads to $\alpha_i = 0$.
 - Which implies that the minimum of the error function has been reached.
- GD leads to a **linear convergence** in the neighborhood of a minimum.
 - The **order of convergence** is linear with rate $1 - \eta\lambda_i$.
 - Since $\lim_{T \rightarrow \infty} \frac{\alpha_i^{(T)} - 0}{(\alpha_i^{T-1} - 0)^1} = 1 - \eta\lambda_i$.
- The **rate of convergence** is governed by $1 - (\frac{2\lambda_{\min}}{\lambda_{\max}})$.
 - Where λ_{\max} and λ_{\min} are the maximal and minimal eigen-values.
 - By increasing η the speed of convergence can be improved.
 - But $|1 - \eta\lambda_i| < 1$ must be ensured, implying $\eta < 2/\lambda_{\max}$.
 - The **fastest convergence** is obtained when $\eta = 1/\lambda_{\max}$.
 - In this case the rate of convergence in the direction of λ_{\max} is 0.
 - Which implies the minimum will be reached in a single step.
 - Assuming to set $\eta = \frac{1}{\lambda_{\max}}$, the direction in which convergence is the slowest is λ_{\min} .
 - In this case the rate of convergence is $1 - \frac{\lambda_{\min}}{\lambda_{\max}}$.
 - The worst case scenario is when λ_{\min} is small compared to λ_{\max} (convergence rate ≈ 1).
 - Convergence will be then very slow.
 - The reciprocal of $\frac{\lambda_{\min}}{\lambda_{\max}}$ is the **condition number** of the Hessian matrix.
 - The larger this number, the slower will be the convergence of the GD.
 - With very dis-equal λ_i , the corresponding hyper-paraboloid will be very irregular.
 - Taking a point on this surface, its gradient won't point exactly to the minimum (**slower convergence**).
 - The gradient of a point is orthogonally directed to the tangent of the level line.
 - While on a perfect hyper-sphere, its gradient will point exactly to the minimum (**faster convergence**).

Momentum

- When convergence is slow, a momentum term can be added to the GD formula.
 - This adds inertia to the motion and smooths out oscillations.
- **def Network update with momentum:** $\Delta w^{(\tau-1)} = -\eta \nabla E(w^{(\tau-1)}) + \mu \Delta w^{(\tau-2)}$.
 - Where μ is the **momentum parameter** (typical value: 0.9).
 - In region of low curvature, the momentum is increasing the learning rate by a factor of $\frac{1}{1-\mu}$.
 - In region of high curvature, the effective learning rate will be close to η .
 - In those regions the GD is oscillatory, successive contributions of the momentum tend to cancel.

Learning rate scheduling

- It is advantageous to change the learning rate η during learning.
- In practice, best results are obtained using a larger value for η at the start of training.
 - And then reducing the learning rate over time.

- $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} - \eta^{(\tau-1)} \nabla E(\mathbf{w}^{(\tau-1)})$.
- ex Common learning schedules:
 - **Linear:** $\eta^{(\tau)} = (1 - \tau/K)\eta_0 + (\tau/K)n_K$.
 - η_0, η_K and K are three parameters and the formula is applied up to K steps.
 - After these K steps, η is kept fixed at η_K .
 - **Power law:** $\eta^{(\tau)} = \eta_0(1 + \tau/s)^c$.
 - The scaling factor s determines how fast the learning rate decreases.
 - $c < 0$ determines the shape of the decrease (usually $c = -0.5$).
 - **Exponential decay:** $\eta^{(\tau)} = \eta_0 e^{-\tau/s}$.
 - s is the scaling factor and $c < 1$ is the decay factor (usually $c \in [0.95, 1)$).
- Extra refinements by optimizing separately η for each direction in parameter space (for each w_i):
 - **AdaGrad (Adaptive Gradient):**
 - Reduce each learning rate parameter using the accumulated sum of squares of all derivatives calculated for that parameter.
 - The idea is to **reduce learning rate for parameters** that have received **large updates in the past**.
 - **RMSProp (Root Mean Square Propagation):**
 - Replace the sum of squared gradients of AdaGrad with an exponentially weighted average.
 - Fixes a problem with AdaGrad where the weight updates tend to become too small.
 - **Adam (Adaptive moments):**
 - Combines RMSProp with momentum.
 - Probably the most used optimization method in DL.

1.3.5.4 Normalization

- Two main problems in DL:
 - Coping with values that vary in very different ranges.
 - Having to deal with **vanishing** and **exploding** gradients.
- Weight normalization try to keep the values computed by the network in a reasonable range.
- Three main approaches: **data normalization**, **batch normalization**, **layer normalization**.

Data normalization

- If the dataset has input variables that span very different ranges.
 - Then a change in one dimension will produce a much larger change in the output wrt a change in another dimension.
- Mean and variance for each dimension is computed and then all data points are rescaled.
 - $\mu_i = \frac{1}{N} \sum_{n=1}^N x_{ni}$, $\sigma_i^2 = \frac{1}{N} \sum_{n=1}^N (x_{ni} - \mu_i)^2$.
 - $\hat{x}_{ni} = \frac{x_{ni} - \mu_i}{\sigma_i}$.
 - Examples distribution will behave as a Gaussian with $\mu = 0$ and $\sigma^2 = 1$ (nice for learning).
 - And applying this to normally distributed examples won't change anything.

Batch normalization

- The same reasoning can be applied to variables (weights) at each hidden layer.
 - Unfortunately, normalization for those values **cannot be done once for all**.
 - The computation need to be performed every time the variables are updated.
- Batch normalization works by normalizing, across the examples of the mini-batch.
 - Normalizing the values computed at each layer of the network by each unit i .
- One can normalize the pre or the post activations (both work well in practice).
 - Pre-activation normalization:
 - $\mu_i = \frac{1}{K} \sum_{n=1}^K a_{ni}$, $\sigma_i^2 = \frac{1}{K} \sum_{n=1}^K (a_{ni} - \mu_i)^2$.
 - $\hat{a}_{ni} = \frac{a_{ni} - \mu_i}{\sqrt{\sigma_i^2 + \delta}}$.
 - K is the size of the mini-batch.
 - δ avoids numerical issues when σ_i^2 is small (eg when using small batches).
- This kind of normalization **reduces the representational capability** of the hidden units.
 - To compensate for this, one can **rescale the pre-activation values** to have $\mu = \beta_i$ and $\sigma = \gamma_i$.
 - $\tilde{a}_{ni} = \gamma_i \hat{a}_{ni} + \beta_i$ (with γ_i and β_i learnable parameters).
 - While originally mean and variance across a minibatch were computed by a complex function of all weights and biases.
 - Now they are determined by two simple independent parameters.

- Which turn out to be much easier to learn by GD.
- By doing this, normalization is not undone.
 - Before, it was hard to find the right compromise between μ and σ^2 just by acting on hundreds of weights.
 - With γ_i and β_i , only those two weights govern this (lines act way better).
- Once training completes, mini-batches to compute the normalization factors won't be available.
 - A moving average of those factors is kept during the training and used at **inference time**.
 - $\bar{\mu}_i^{(\tau)} = \alpha \bar{\mu}_i^{(\tau-1)} + (1 - \alpha) \mu_i$ (with $0 \leq \alpha \leq 1$).
 - $\bar{\sigma}_i^{(\tau)} = \alpha \bar{\sigma}_i^{(\tau-1)} + (1 - \alpha) \sigma_i$ (with $0 \leq \alpha \leq 1$).

Layer normalization

- Normalization across the **hidden unit values** for each **data point separately**.
 - Instead of normalizing across examples within a mini-batch for each hidden unit separately.
- It updates the pre-activation values as it follows:
 - $\mu_n = \frac{1}{M} \sum_{i=1}^M a_{ni}$, $\sigma_n^2 = \frac{1}{M} \sum_{i=1}^M (a_{ni} - \mu_n)^2$.
 - $\hat{a}_{ni} = \frac{a_{ni} - \mu_n}{\sqrt{\sigma_n^2 + \delta}}$.
 - n ranges over the examples and M is the number of hidden units in the layer.
- **Additional learnable parameters** β_i and λ_i are introduced (similar to batch normalization).
 - In this case there is no need to keep moving averages to normalize data at inference time.

1.4 Hopfield networks and restricted Boltzmann Machines

1.4.1 Hopfield networks

- Introduced by the physicist John Hopfield in 1982 (2024 Nobel in Physics).
- Used to represent **content addressable memory**:
 - Given a reasonable portion of the memorized information and a partial corrupted portion.
 - It can be completed by recovering other associated information (**pattern completion**).
- The system preserves this property also when some components are broken (**fault tolerance**).
- *Attractor networks*: useful to memorize fundamental memories which attract other inputs.
 - Given a partial (and/or noisy) figure it retrieves the original figure.
- **Pattern completion** characterizes a lot of human memory.
 - These network (and the attractor network in general) are used to model several cognitive phenomena.
- Network of McCulloch-Pitts neurons interconnected, where each neuron can have two states (1 and -1).
 - Given N units:
 - Each neuron is connected with all the other neurons but itself (**fully connected**).
 - Weights are symmetrical $w_{ij} = w_{ji}$.

1. **Input.** To the network is presented x^* : configuration +1 and -1 of length N . For each neuron j , $y_j(0) = x^*(j)$

2. **Iteration until convergence.** Update the elements of the network in an asynchronous way by choosing a unit at random. The rule to update the elements is

$$y_j(n+1) = \begin{cases} 1 & \text{if } \sum_{i=0}^N w_{ji} y_i(n) > 0 \\ -1 & \text{if } \sum_{i=0}^N w_{ji} y_i(n) < 0 \\ v_j(n) & \text{if } v_j(n) = 0 \end{cases}$$

3. **Convergence.** Repeat the operation until we find a **stable state**, s.t. for each unit j $y_j(n+1) = y_j(n)$.

4. **Output.** This stable state is the output of the network.

Figure 4: Information withdrawal phase in Hopfield networks.

1.4.1.1 Information withdrawal phase

- The **pattern completion** happens during the **withdrawal/retrieval phase**.
 - *State* of a network with N neurons at a given iteration: activations of the N neurons (1, -1).
 - *Input*: vectors of 1 and -1 of dimension N .
 - Each neuron represent an element of the input.
 - Then calculate new activation of a neuron at a time, randomly chosen, until there is no change.
- **def Activation of j at iteration n :** $y_j(n) = \varphi(v_j(n))$.
 - $v_j = \text{netinput to } j = \sum_{i=0}^N w_{ji} y_i(n)$.
 - $\varphi(v_j)(n) = -1$ if $v_j(n) < 0$, 1 if $v_j(n) > 0$, and $\varphi(v_j)(n-1)$ if $v_j(n) = 0$.
 - Activations are calculated in an **asynchronous way**.
- A **stable state** (where $\forall i, y_i(n+1) = y_i(n)$) will be the output of the network.
 - **Stable states are not unique**, they depends on the initial input and/or evaluation order.
 - prop The opposite of a stable state is a stable state.
 - A initial input (eg noisy image) is perturbed to a stable state (eg actual image).
 - Stable states act as **attractors**: given any state, it will converge in one of the stored states.

1.4.1.2 Storage phase

- The information withdrawal phase assumes weights are already computed.
 - Those weights are computed in the **storage phase**.
 - Once the memories are stored, the pattern completion can start.
- **Memorized information works as attractor**.
- **def Weights computation:** $w_{ji} = \frac{1}{M} \sum_{k=1}^M f_k(i) \cdot f_k(j)$ where $j \neq i$.
 - Given M fundamental memories and f_1, \dots, f_M vectors of dimension N .
 - For each fundamental memory the product between the i -th and the j -th element.

- If the have different sign, $f_k(i) \cdot f_k(j) < 0$ contributes to weaken w_{ji} .
- If the have the same sign, $f_k(i) \cdot f_k(j) > 0$ contributes to strengthen w_{ji} .
- The learning algorithm is based on **Hebb's principle**:
 - Strengthen the connections between units with the same activation.
 - Weaken the connections between units with opposite activation.
 - Hebb's principle has a neurobiological counterpart:
 - In the synapses between units which are often active at the same time are strengthened.
 - While synapses between neurons which are not simultaneously active are weakened.
 - *Neurons that fire together wire together.*

1.4.1.3 Convergence of Hopfield networks

th Convergence: given a starting state the network will always arrive to a stable state.

- The activities cannot be changed forever.
- *Proof:*
 - Given N units, there are 2^N possible states of the network.
 - To each of these states is associated an **energy** value.
 - **def Energy function:** $E = -\frac{1}{2} \sum_i \sum_j w_{ij} y_i y_j$.
 - E measures the predisposition of the network to change state.
 - Opposite of **harmony** (bad).
 - Harmony increases with the number of positive $w_{ij} y_i y_j$.
 - Higher harmony \rightarrow less predisposition to change state.
 - Lower harmony \rightarrow more predisposition to change state.
 - Each change of state of the network leads to a lowering of the energy.
 - Each $w_{ij} y_i y_j$ is positive if y_i and y_j have the same sign and w_{ij} is positive.
 - Or they have opposite sign, and w_{ij} is negative.
 - The activations enforce each other (**low energy \rightarrow low predisposition to change**).
 - Each $w_{ij} y_i y_j$ is negative if y_i and y_j have the same sign and w_{ij} is negative.
 - Or they have opposite sign, and w_{ij} is positive.
 - Instable (**high energy \rightarrow high predisposition to change**).
 - Consider how E changes with the activation of the k -th unit.
 - E and E' denote energy before and after the change of y_k .
 - $E = \frac{1}{2} \sum_i \sum_j w_{ij} y_i y_j = -\frac{1}{2} (\sum_{i \neq k} \sum_{j \neq k} w_{ij} y_i y_j + 2 \sum_{j \neq k} w_{kj} y_k y_j)$.
 - $E' = -\frac{1}{2} (\sum_{i \neq k} \sum_{j \neq k} w_{ij} y_i y_j + 2 \sum_{j \neq k} w_{kj} y'_k y_j)$.
 - $E - E' = -\frac{1}{2} (2 \sum_{j \neq k} w_{kj} y_k y_j - 2 \sum_{j \neq k} w_{kj} y'_k y_j) = - \sum_{j \neq k} w_{kj} y_j (y_k - y'_k)$.
 - $\sum_{j \neq k} w_{kj} y_j$ is v_k .
 - There are two cases:
 - y_k passes from $+1$ to -1 :
 - Hence $y_k - y'_k > 0$ and $\sum_j w_{kj} y_j < 0$.
 - Then $-\sum_{j \neq k} w_{kj} y_j (y_k - y'_k) > 0$ and $E > E'$.
 - y_k passes from -1 to $+1$:
 - Hence $y_k - y'_k < 0$ and $\sum_j w_{kj} y_j > 0$.
 - Then $-\sum_{j \neq k} w_{kj} y_j (y_k - y'_k) > 0$ and $E > E'$.
 - At each change E lowers.
 - Therefore, there is a moment in which there is **no reachable state with lower energy**.
 - Since the states are finite, at a given point there's a state that doesn't change.
 - This state is a **stable state**.

1.4.1.4 Qualities and drawback of Hopfield networks

- Qualities:
 - They **complete partial patterns**.
 - They **generalize**.
 - Given an input similar to what has been memorized, they recover the corresponding information.
 - They are **fault tolerant**.
 - If some synapses get broken (*brain damage*) the output is still reasonable.
 - They allow the **extraction of prototypes**.

- If the NN learns several similar informations, it creates their prototype and this is the explicitly presented state.
- The learning rule is **Hebb-like**.
 - Which is biologically plausible and there is evidence that it exists in the brain.
- The networks can account for **context effects**.
 - Humans remember better what is learned if they are put in the same context.
- Drawbacks:
 - Not all the stable states are fundamental memories memorized during storage.
 - There are **spurious states**:
 - The opposite of a stable state is a stable state.
 - Combinations of stable states are stable states.
 - Not all fundamental memories are stable states.
 - Storage capacity with few errors given N units = $0.14N$ (*rule of thumb*).
 - Limited storage capacity.
 - Errors.
 - In the brain synapses are not symmetrical.
 - In the brain there are no stable states but states transitorily stable, which evolve in **successive states**.
 - There are extensions of Hopfield networks that learn sequences of states.
- Hybrid models are being used.
 - Standard convolutional model are poor at pattern completion.
 - eg Convolutional + HN for pattern completion.

1.4.2 Restricted Boltzmann machines

- **Spurious states** are the main limitation of Hopfield networks.
 - The capacity of a totally connected net with N units is only about $0.15N$ memories.
 - An inefficient use of the bits required to store the weights.
 - A phase of **unlearning** can be used after learning to get rid of spurious states.
 - Hinton suggests this is the function of dreams.
- Three main idea behind Restricted Boltzmann machine:
 - Visible vs. **Hidden units**.
 - **Stochastic** units.
 - New learning algorithm to correct errors (*fantasies*).
 - HN use *one-shot learning* during the storage phase.

1.4.2.1 Hidden states

- Main features:
 - Two levels: visible and hidden neurons.
 - Connections only between visible and hidden units.
 - All visible units are connected to all hidden units.
 - Weights are symmetrical: $w_{ij} = w_{ji}$.
 - Possible neurons' activation states: 0 or 1 (and not -1 and -1).
 - Neurons are **stochastic**.
 - Input and output are computed on the visible layer.
- Instead of using the net to store memories, use it to construct **interpretations** of sensory input.
 - The input is represented by the visible units.
 - The interpretation is represented by the states of the hidden units.
 - Construct interpretations: detect and **represent regularities** in visual activations.
 - Hidden units specialise in recognising (and generating) **portions of visual patterns**.

1.4.2.2 Stochastic units

- Noisy networks find better energy minima.
- A HN always makes decisions that reduce the energy.
 - This makes it impossible to escape from local minima.
- Stochastic binary units:
 - Replace the binary threshold units by binary stochastic units that make biased random decision.
 - def **Activation probability**: $P(S_i = 1) = \frac{1}{1+e^{-\Delta E_i/T}}$ (sigmoid).
 - def **Energy gap**: $\Delta E_i = E(S_i = 0) - E(S_i = 1)$.
 - Objective: minimize the energy.

- $E = - \sum_k \sum_j w_{kj} v_k h_j$.
- $\Delta E_i = \sum w_{ij} s_j$ for all j connected to i .
- If $\Delta E_i < 0$, $p(s_i = 1) < 0.5$.
- If $\Delta E_i = 0$, $p(s_i = 1) = 0.5$ (no difference by assigning 1 or 0).
- If $\Delta E_i > 0$, $p(s_i = 1) > 0.5$.
- The *temperature* T controls the amount of noise.
- $T = 1$ is assumed (can be ignored it).
- **Sampling:**
 - Sampling is used then to attribute a state (0 or 1) using the corresponding probability.
 - The same VL activation is presented multiple times and corresponding HL activation is collected.

1.4.2.3 Execution of a RBM

- Procedure:
 - Present a visual input.
 - Calculate activation at the hidden level.
 - Recalculate activation at the visual level.
 - Possible, start again.
 - Continue until *convergence* (no significant change).
- By evaluating the hidden layer only, the most probable activation at the visible level is provided.
 - An activation of the HL represents a **codified version** of the corresponding VL activation.
 - Experimentally it has been observed each HL neuron *specialises* in a sub-set of VL neurons.
 - It also works with different (or corrupted) patterns.
- RBM are **generative models**: they are used to generate visual states.
 - The process of generating them usually requires several steps.
 - Calculate visual states, then hidden, then visual again, etc.
 - Until equilibrium (a state in which there are no big changes in the probability distribution).
 - They learn to **reproduce** at the visual level the training set.

1.4.2.4 Weights learning in RBM

- Goal: maximize the probability of obtaining at the VL the vectors of Tr .
 - Tr : set of visible vectors to be learned by the RBM (in order to reconstruct or generate).
- A **fantasy** is the reproduction (at the VL) of an original pattern after a very high number of steps.
 - The weights are learned in order to minimize difference between pattern and fantasies.
 - Actually, Hinton discovered that just **two steps** are needed («a very surprising short-cut»).
 - Importance of weights in RBM:
 - Hidden units specialise in recognising (and generating) **portions of patterns**.
 - The probability of unit i to get state 1 depends on the weights to i .
 - Roughly, positive weight w_{ij} augments the impact of neuron j to $p(s_i = 1)$.
 - By observing weights matrices to HU, it can be understood which inputs values augment their probability to activate.
- After learning the weights starting from Tr , RBM is ready to be used:
 - Generation:** generate a visual activation pattern from a random hidden activation pattern.
 - Reconstruction:**
 - Start with a visible input.
 - See what visual activation the RBM reproduces after a few passage V-H-V, etc.

Contrastive divergence (learning algorithm)

- Execution:
 - Start with a training vector on the visible units.
 - Update all the hidden units in parallel (since they're independent).
 - Update all the visible units in parallel to get a *reconstruction*.
 - Update the hidden units again.
- def **Weights update**: $\Delta w_{ij} = \epsilon(\langle v_i h_j \rangle^0 - \langle v_i h_j \rangle^1)$.
 - v_i and h_j are the activations of visual unit i and hidden unit j .
 - ϵ is the learning rate.
 - Just **two step** ($\langle \cdot \rangle^0$ and $\langle \cdot \rangle^1$) are needed.
- Everything must be iterated for several epochs.
- Overall algorithm:

- Weights are initialised randomly.
- At each epoch n :
 - An element of Tr is chosen.
 - For each weight, Δw_{ij} is computed, and $w_{ij}(n + 1) = w_{ij}(n) + \Delta w_{ij}$.
- Until the maximum number of epochs is reached.

1.4.2.5 Deep (or stacked) RBM

- Idea:
 - First train a layer of features that receives input directly from the pixels.
 - This layer discovers patterns of activities, regularities, in the layer below.
 - Then treat the activations of the trained features as if they were pixels and learn features of features in a second HL.
 - This second HL discovers patterns of activities in the layer below correlations between features discovered at the previous level.
 - Each level models the correlation between activities in the level below.
 - Then do it again.
- It's proven that with each additional layer of features the probability for generation Tr is **improved**.
- SBM can be **fine tuned** according to the desired applications.
 - If used for *classification*, by adding a classification layer.
 - The first really successful deep autoencoders [Hinton & Salakhutdinov, 2006].
 - If used for *generation*, with **Deep belief nets**.
 - If used for finding few very descriptive features with **autoencoders**.
- As in CNN, the more deep RBM are specialized on more complex features.

1.5 Convolutional neural networks

1.5.1 Convolutional neural networks

- Designed to solve computer **vision problem** (eg image classification, object detection, etc).
 - Traditionally, computer vision was on *hand-crafted features*.
 - Today features are no longer hand-craft but **learned** by CNN architectures.
 - Usually, those features cannot be described in intelligible terms.
 - Traditional NN don't scale well for image data.
- Key advantages of CNN: local connections, shared weights, pooling, hierarchy of features.
 - Passage through all these phases is necessary since it's not possible to pass from raw pixel to abstract concepts.
 - The focus of DNN is to go through all these levels of abstraction.
 - Designed to encode invariances and equivariances **specific** to image data.
- Multilayer perceptrons don't scale well on big size images.
 - eg With a 1000px square image, the corresponding network should have 3M weights.
 - They also loose the **spatial information** of an image.

1.5.1.1 CNN basic elements

- **Convolution** layers, **pooling** layers, **fully-connected** layers, activation function.
- Two terminological variants:
 - In the complex layer terminology, the convolutional layer is composed by:
 - Convolution stage → Detector stage (eg ReLU) → Pooling stage.
 - In the simple layer terminology, the convolutional layer is the convolution stage only.
- CNN contain multiple such layers of convolution and pooling in succession.
 - In which the output channels of a particular layer form the input channel of the next layer.
 - After convolution blocks there are one or more **fully connected layers**.

1.5.1.2 CNNs applications

- CNNs have state-of-the-art accuracy for **image classification tasks**.
 - Their development was accelerated through the introduction of **ImageNet**.
 - Large-scale benchmark dataset: 14M natural images, hand labelled into 22K categories.
 - Image Classification, object detection, image segmentation, neural style transfer.
 - Used also for other tasks, as analysis of sequential data.
- Recent alternative architectures based on Transformers have become competitive with CNN.

ImageNet Large Scale Visual Recognition Challenge (ILSVRC, 2010-2017)

- ImageNet subset with 1K non-overlapping categories, 1.28M training, 50K validation and 100K test images.
- Key milestones:
 - AlexNet (2012, deep breakthrough).
 - VGG, GoogleLeNet, ResNet (advanced depth and accuracy).
- The model outputs a ranked list of 1K class probabilities for each image (output of softmax layer).
 - If the correct label is within the top-5 predictions, it counts as correct.
 - Top-5 accuracy = $\frac{\text{# images where true label} \in \text{top-5 predictions}}{\text{total images}}$.
 - Top-5 error rate = $1 - \text{Top-5 accuracy}$, 5% for human, 2.3% in 2017 for a CV system.
 - Superhuman performance maybe because humans cannot distinguish specific classes [Bishop].
- Further developments:
 - Models have been evaluated on noisy and blurred images [2017].
 - Slightly different dataset has been proposed (eg ImageNetV2) [2019].
 - Top-1 accuracy = $\frac{\text{# number of correct predictions}}{\text{total images}}$ [2023].
 - ImageNet still drives progress to date, but Top-1 accuracy is stagnating.

Notable CNNs

- LeNet (one of earliest CNN models) [LeCun et al., 1989].
 - Two convolutional layers, pooling and fully connected layers.
 - Used for digit recognition tasks (MNIST dataset).
- **AlexNet** [Krizhevsky, Sutskever & Hinton, 2012]:

- Won the 2012 competition with top-5 error rate of 15.3%.
- Use of the ReLU activation function.
- Application of GPUs to train the network.
- Use of dropout regularization.
- VGG-16 model [2014].
- ResNet family.
- Inception Family (GoogLeNet, V2, V3).

1.5.1.3 Convolution layers

- Hidden units are connected to a **portion** of the image, the unit's **receptive field**.
 - Instead of linearizing the image, and loosing spatial information.
- Weight vector: **filter** (or **kernel**) that determines the kind of features for which the unit activates the most.
 - Unit activation: $z = \text{ReLU}(w^T x + w_0)$.
 - **w** and **x** are vector's representation of pixel values and weights, while w_0 the bias.
- **Feature maps**:
 - Nearby units are connected to adjacent regions of the image with the same weight vector.
 - Detect the **same feature** in a different area of the image.
 - **Weight sharing** (all units share the same weights).
 - Units of the hidden layer form a **feature map**.
 - All units of a feature map detect the **same feature** in a different area of the image.
 - The filter goes over the image and determines the activation of the associated units in the FM.
- Convolutional layers contain a **set of feature maps**, each with a different filter.
 - eg Examples of convolution filters: identity, edge detection, sharpen, Gaussian blur.

Properties of convolution layers

- **Sparse weights**.
- **Parameter (weight) sharing**.
- **Locality-based analysis**.
- **Equivariance to translation**.
 - A feature detected in a position can be detected in another position never seen before by means of the same filter.
 - It will activate a unit in the same layer but possibly in a different location.
 - It will activate a nearby unit in the same feature map.
- **Translation equivariance \neq translation invariance**.
 - At classification level, particular object should be assigned the same classification irrespective of position.
 - **Scale invariance**: changes to object size should also leave its classification unchanged.
 - For MLP to learn invariance one would need to consider huge datasets (like by using data augmentation).

Convolution layers hyperparameters

- **Kernel size** (eg 3×3 , usually odd numbers).
 - Larger kernel size \rightarrow smaller feature map.
- **Stride**: determines how far the filter moves across the input.
 - The filter is moved in steps of size S .
 - Larger stride \rightarrow smaller feature map.
- **Padding**: add extra pixels around the edge of the input to control output size.
 - So the feature map can have the same dimensions as the original image.
 - Larger strides reduce the output size, while padding preserves more information.
- **Convolutional filters number** (of feature maps in a convolutional layer).

Filters

- In classical image processing, kernels' weights were **handcrafted** by experts.
 - Instead, CNNs **learn the best weights** (via backpropagation).
- Learned filters are much less neat and interpretable than the examples made by humans.
- **Multidimensional convolution**:
 - A multi-dimensional filter takes input from across the R, G and B channels.
 - eg A 3×3 filter has 27 weights (plus a bias) and can be visualized as a 3×3 tensor.

- This holds for the first level of the kernels.
 - For higher level filters the number of input channel is in general higher than 3.
- Filters in deeper levels cannot be visualized directly.
 - The input that maximizes activation for all neurons of a given FM can be visualized.
 - **Saliency maps:**
 - Used to identify those regions of an image most significant in determining the class label.
- Multiple independent filter lead to multiple feature maps, forming a single **convolutional layer**.
 - The filter operates on a portion of the input volume.
 - Each slice of neurons (same depth) denotes a feature map.
 - Weights are shared in a feature map.
 - Neurons in the same FM process different portions of the input volume in the **same way**.
 - Each feature map can be seen as the result of a specific filtering of input.

1.5.1.4 Rectified Linear Unit

- MLP historically used the sigmoid activation function.
 - In deep network, its use is problematic for back propagation (**vanishing gradient**).
 - Sigmoid for large and small values derivatives is close to 0.
 - In weight update: $\delta^l = ((\mathbf{W}^{l+1})^T \delta^{l+1}) \odot \sigma'(\mathbf{a}^l)$.
 - If σ' is close to 0 the product is very small → vanishing gradient.
- **def Rectified linear unit activation function (ReLU):** $f(u) = \max(0, u)$.
 - The derivative is 0 for negative values, and 1 for positive values.
 - It introduces non-linearity to the model.
 - This leads to sparse activations (part of the neurons are off).
 - Which leads to simpler activation patterns, which makes it **more generalizable**.
 - Unit acts as a **feature detector** that signals when it finds a sufficiently good match to its kernel.

1.5.1.5 Pooling

- A pooling function replaces the output of the net at a certain location with a summary statistic of the nearby output.
 - The **max pooling** operation reports the maximum output within a rectangular neighborhood.
 - The **average pooling** operation reports the average output within a rectangular neighborhood.
- It **reduces dimensionality** (downsampling).
- It selects the **most informative values**.
- It implements **invariance wrt (small) translation**.
- Pooling layers have no weights.
 - Pooling does not have learnable parameters but only hyperparameters (fixed with the architecture).
 - eg A standard choice is max-pooling over 2×2 kernels with stride 2.

1.5.1.6 Cost functions

- For classification with CNN, common CFs are: **Soft-Max** and **Cross-Entropy**.
- **def Soft-Max:** $z_k = f(v_k) = \frac{e^{v_k}}{\sum_{c=1, \dots, s} e^{v_c}}$.
 - It consist of a final layer of s neurons (one for each class) fully connected to the previous layer.
 - The neuron inputs v_k are calculated as usual, but the activation function for the k -th neuron is the one above.
 - Values z_k can be viewed as probabilities: they belong to $[0, 1]$ and their sum is 1.
- **def Cross-Entropy loss:** $-\sum_{i=1}^S t_i \cdot \log(p_i)$.
 - It measures how much the predicted distribution p **differs** from the desired output.
 - t_i is the desired output for neuron i , while p_i is the predicted value.
 - Functioning:
 - If the model predicts $p_j \approx 1$ for the correct class, the prediction is very close to the truth.
 - $\log(1) = 0 \rightarrow$ low cross-entropy.
 - If the model predicts $p_j \ll 1$ (possibly $p_j = 0$) for the correct class, the prediction is very wrong.
 - $\log(0) = -\infty \rightarrow$ high cross-entropy.
 - Backpropagation is then applied with this loss function.
 - Cross entropy between two discrete distributions: $H(p, q) = -\sum_v p(v) \cdot \log(q(v))$.
 - It measures how much q **differs** from p .
- By using Soft-Max and Cross-entropy, BP involves calculating gradients for both the output and hidden units.

1.5.1.7 Avoiding overfitting

- **Dropout regularization:**
 - During training, some neurons are randomly ignored (*dropped out*) with a given probability p .
 - It speeds up the training process.
 - More training iterations needed, but those will be faster.
 - Encourages **sparsity**, thus enhances generalization and avoids overfitting.
 - Typically implemented as a *layer* that zeroes the previous neurons' output with probability p .
 - Placed after **activation functions**.
- **Early stopping:**
 - Monitors the model's performance on a **validation set** during training.
 - Stop the training when performance starts to degrade (indicating overfitting).
 - A way to keep the model from continuing to learn noise in the Tr after the optimal point.
- **Data augmentation:**
 - Not a regularization technique in the traditional sense.
 - It artificially increases the size of the training dataset by applying random transformation to Tr images.
 - It helps the model generalize better by exposing it to variation to be found in the real world.
- **Batch normalization:**
 - It normalizes the output of each layer.
 - Which can reduce overfitting by introducing noise to the training process.
- **Weight decay:**
 - Adds a penalty to the loss function that encourages weights to be small, thus preventing overfitting.

1.5.1.8 Neuroscientific basis of CNN

- In the mammalian visual cortex, there are simple neurons that detect simple features.
 - Hubel & Wiesel paved the way in this research by analyzing cat's striate cortex.
- Several studies tried to establish commonalities between CNN and brain processing.

1.5.2 Transfer learning

- The training of complex CNN on large dataset can require a lot of machine time (even on GPUs).
 - One trained, the classification of a new pattern is generally fast (eg 10-100ms).
- **Fine-tuning:**
 - Start with a pre-trained network trained on a similar problem.
 - The output layer is replaced with a new layer of Soft-Max neurons (adapting classes number).
 - As initial values for weights those of the pre-trained network are used.
 - Except for connections between penultimate and final layers whose weights are initialized random.
 - New training iteration with SDG are carried out to optimize weights.
 - To respect to the peculiarities of the new dataset (not necessary to be of large size).
 - eg Examples of fine-tuning:
 - A network trained on object classification → lesion classification (**scarcer dataset**).
- **Feature reusing:**
 - A pre-trained NN is used without further fine-tuning.
 - Generated features by the NN during the forward step on the new dataset are extracted (at intermediate layers).
 - These features are used to train an external classifier.
 - To classify the patterns belonging to the new application domain.

1.6 Recurrent neural networks

1.6.1 Recurrent neural networks

- Specialized to work on **sequences**:
 - Speech recognition, music generation, DNA analysis, sentiment classification, translation, NER, LM, etc.
 - The input and the output can have same or different dimensions.
 - Standard NN aren't well suited for seq-to-seq tasks.

- Inputs and outputs can be different lengths in different examples.
- Doesn't share features learned across different positions of text.
- **Context** must be taken into account.
- **Elman networks** [1990]:
 - Precursor of simple RNN, built for cognitive purposes.
 - Introduced to explain how infants could learn to identify words from a continuous stream of phonemes.
 - eg «many years ago a boy and girl lived by the sea they played happily».
 - It showed how this network can induce categorical grammar (noun/verb) of input sentences with no previous knowledge.
 - Just by learning it from example.
 - Architecture:
 - Input units.
 - Hidden units.
 - The output of HU is fed to output units but also to **context units**.
 - At the next iteration, the new input (eg next word in the sentence) and previously stored context units are fed to the current hidden units (**recursion**).
 - Output units.
 - At the $i + 1$ -th time step, the i -th activation is taken into account.
 - Unrolled vs Rolled representation (the latter represent the recurrent component).
 - The rolled representation highlights that the **weights are shared**.
 - def **RNN Activation**: $a^{(i)} = g(W_{ax} * x^{(i)} + b_a + W_{aa} * a^{(i-1)})$ and $\hat{y}^{(i)} = g'(W_{ya} a^{(i)} + b_y)$
 - $a^{(0)}$: zero vector.

1.6.1.1 Training a RNN

- Training is performed using **Backpropagation Through Time** (BPTT).
 - BPTT begins by unfolding a recurrent neural network in time.
 - The unfolded network contains k inputs and outputs.
 - But every copy of the network shares the same parameters.
 - Then, BP is used to find the gradient of the loss function wrt all the network parameters.
- Problems:
 - **Vanishing or exploding gradient** (technical problem).
 - Not very good at capturing **long distance dependencies**.

1.6.2 Long Short Term Memory

- To capture long distance dependencies, a sort of **memory** is needed.
- A **LSTM cell** substitutes the hidden layer of a RNN.
 - Each cell receive an input from the example and the activation of the previous LSTM cell (a la RNN).
 - Units have **gates** that decide when to update memory cell states, what to keep, what to forget.
- Three type of gate: **forget**, **input** and **output gate**.
- Applications:
 - Speech recognition, machine translation, syntactic parsing, handwriting recognition, image captioning.
- LSTM can be **stacked** too (deep LSTM).
 - The output of the LSTM below is used for the upper one.

1.6.2.1 LSTM architecture

- **Cell state**:
 - The horizontal line running through the top of the diagram.
 - It transports informations down the entire chain, with only some minor linear interactions.
 - LSTM have the ability to remove or add information to the cell state, carefully regulated by **gates**.
- def **LSTM Forget gate**: $f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$.
 - What information will be eliminated from the cell state.
 - Given h_{t-1} and x_t , f_t is a vector of values between 1 and 0 (with same dimension as C_{t-1}).
 - Values of f_t are then multiplied element-wise with C_{t-1} .
 - Thus gating what values of C_{t-1} are kept, what are thrown away.
- def **LSTM Input gate**: $i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$.
 - What information should be added (eg information on new subject).

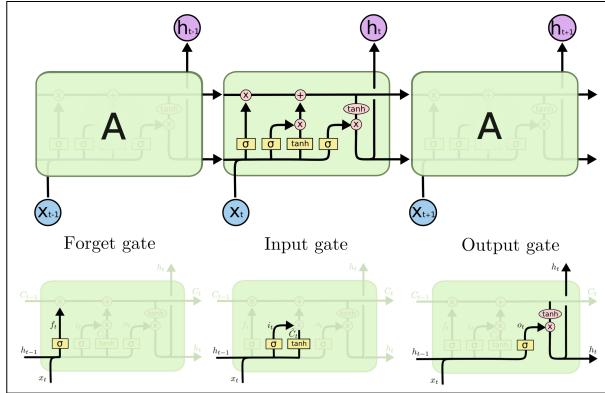


Figure 5: LSTM architecture with corresponding gates.

- Input gate is multiplied with **candidate cell state** $\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$.
 - And the result is then added to the previous cell state.
 - As a result of previous operations by forget and input gate, a **new cell state** is computed.
 - $C_t = f_t \cdot C_{t-1} + i_t \cdot \tilde{C}_t$.
 - def **LSTM Output gate**: $o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$.
 - Which part of the cell state will be kept.
 - From the new cell state C_t , the **output** h_t is computed.
 - $h_t = o_t \cdot \tanh(C_t)$.
 - o_t is applied to cell state C_t to really compute (filter) what to output.
-

1.7 Cognitive Assessment

1.7.1 NN and cognition

- Three different approaches:
 - **ANN as cognitive models**.
 - Cognitive models of how humans might learn language or analyze visual information.
 - First NN for language aimed to be cognitively plausible models of language learning in infants.
 - eg Elman's first RNN, Rumelhart and McClelland's PDP model.
 - Tasks: learn word-meaning associations, learn past-tense, etc.
 - Perceptron aimed to model simple visual cells in the mammalian cortex.
 - CNN are considered as good models of visual information is processed in the brain.
 - **Cognitive assessment of NN models**.
 - Use tools of CogSci and linguistics to probe models' performance and compare to human performance.
 - Contemporary NN are focuses to application with looser relation to cognitive plausibility.
 - However, cognitive tools ave been used to assess NN performance.
 - **ANN as models to understand the brain**.
 - Several works to predict and describe activity in areas of the brain using CNNs (and family).
 - Di Carlo's claim: CNN are a good model of how visual information is processed in the brain ($V1 \rightarrow V2 \rightarrow PIT \dots$).
 - Learning how the brain processes visual information can also improve CNNs.
 - eg CNN models with a neural hidden layer that better match primary visual cortex ($V1$) are also more robust to adversarial attack.

1.7.2 ANN as cognitive models

1.7.2.1 Interactive Activation model [Rumelhart & MCCelland, 1981-82]

- A possible mechanistic explanation of **context effect**.
 - People are faster and more accurate at identifying the O in real word SPOT vs. TKOR.
 - This suggest that humans use orthographic rules (not just familiarity).
- Question: do human represent these rules in their minds and use them in the service of perceiving letters.
 - No, humans are not using orthographic rules when reading.

- People simply process visual input using the **interactive activation model**, which explains the context effect.
 - Pipeline: Visual input → feature level ↔ letter level ↔ word level.
 - Levels have self-loops and feed information back to the level underneath.

1.7.2.2 Debate between nativism and empiricism

- **Nativism** (Chomsky, Fodor):
 - The mind contains **innate, domain-specific structures** (eg Universal grammar).
 - Learning is impossible without rich built-in constraints.
 - *Poverty of stimulus*: input is insufficient, innate knowledge must fill the gaps.
- **Empiricism** (English tradition as in Lock and Hume, now the *USA East coast*):
 - The mind begins with **few domain-specific structures**.
 - Rich learning mechanisms extract structure from experience.
 - Statistical regularities in input are sufficient for acquisition.
- **Connectionism** (Rumelhart, McClelland, Elman):
 - **No built-in rules** or domain-specific constraints.
 - Structures **emerge** from exposure to data (eg past tense learning, word segmentation, phonotactics).
 - Connectionist models were historically developed *as a challenge* to nativist claims.
 - eg **Elman's (recurrent) neural network**:
 - Task: predict the next letter in a sequence of letters.
 - The network was presented with one letter at a time.
 - It takes into account the context in order to make a prediction via the **context layer**.
 - The network learns alone word boundaries.
 - Prediction error accuracy is somehow a **proxy of word boundaries**.
 - Inside a single word the error becomes lower, and it increases when a new word begin.
 - Task: predict the next word in a sequence of words.
 - By analysing inner representation, words can be clustered according to their similarities into **syntactic categories**.
 - The network learns to form a inner representation for verbs and one for nouns.
 - This is a contribution to the empiricist view.
 - No previous knowledge of syntactic structures needed.

1.7.3 Cognitive assessment of NN models

1.7.3.1 Assessing the Ability of LSTM to Learn Syntax-Sensitive Dependencies [Linzen, Dupoux, Goldberd, 2016]

- **Cognitive science as a tool to evaluate NN models.**
- Objectives of the paper:
 - Stress test for LSTM (performance side):
 - Success in language processing is typically attributed to theirs ability to capture **long-distance statistical regularities**.
 - On the cognitive side:
 - Limits of RNN as model of (child) language learning (as in Elman, 1990-91).
 - **There are rules that infants learn but LSTM don't learn.**
 - Therefore either LSTM is not the right model, or there is innate knowledge (authors more on the nativist side).

Structured representations

- Arguments for **structured representations**:
 - Many word co-occurring statistics can be captured by treating sentences as an unstructured list of words.
 - Most naturally occurring agreement cases in the Wikipedia (training corpus) are easy.
 - Solvable without syntactic information, based only on the sequence of nouns preceding the verb.
 - Other dependencies however are **sensitive to the syntactic structure** of a sentence.
 - Given the difficulty in identifying the subject from the linear sequence of the sentence.
 - Dependencies such as subject-verb agreement serve as an argument for structured syntactic representations in human.
 - But LSTM don't use any type of structured representations.

- Sadly (compared to Elman's results), LSTMs don't behave well on complex sentence with several and far apart attractors.

Task considered

- Task considered: subject-verb agreement.
- Training set: derived from Wikipedia.
- Three different training conditions:
 - Most favorable condition:
 - Training with explicit supervision directly on the task of guessing the number of a verb based on the words that preceded it.
 - Error rates of LSTM evolve as a **function of distance** (with no intervening nouns).
 - Grammatically judgment training objective.
 - Full sentences annotated as to whether or not they violate subject-verb number agreement.
 - Training of a model without any grammatical supervision.
 - Instead using a *language modeling* objective (next word prediction).

Results and improvements

- Results (number prediction task):
 - Most naturally occurring agreement cases in Wikipedia are easy.
 - And can be solved without syntactic information.
 - This lead to high overall accuracy in all models.
 - The accuracy of supervised model is lower on **harder cases**, even if it managed to recover.
 - More mistakes when no overt cues to syntactic structure (function words) are available.
- Related tasks:
 - Verb inflection (eg write/writes, access to semantics too).
 - Grammaticality judgments (weaker supervision, no syntactic clues boundaries).
 - Language model (worst performer of all).
- Main theoretical results:
 - LSTM make errors on **harder sentences** (mostly if trained in the language model setting).
 - It highlights the limitation of LSTMs that wouldn't be apparent when **analysing average performance**.
 - Therefore, **explicit supervision** is necessary for learning the agreement dependency using LSTM.
 - **Limiting its plausibility** as a model of child language acquisition (similarly to Elman, 1990, for RNNs).
- Some (not general) single hidden units were found to be specialized in detecting singular or plural.
- Possible improvements:
 - LSTM can be encouraged to develop more sophisticated generalizations by **oversampling** grammatically challenging training sentences.
 - Language modeling objective is not by itself sufficient for learning structure-sensitive dependencies.
 - **Joint training objective** (with supervisions as in SVA task) can be used.
 - To supplements LM on tasks for which syntax-sensitive dependencies are important.

Other neural architectures

- In following papers by different authors, newer neural architectures were tested.
- NN based on Transformers perform well on syntactic tasks (eg subject-verb agreement).
 - However, they may be *right for the wrong reasons*.
 - Learning wrong heuristics, learning statistical approximations, etc.
- *Embers of autoregression* [McCoy et al., 2024]:
 - LLM learn to predict the next words, and this has its advantages but also disadvantages.
 - They displays sensitivity to **output probability**, **input probability** and **task frequency**.
 - Even state-of-the-art models tends to fail on **lower probability sentences**.
 - If the rule is well defined, they don't infer the underlying rule, they learn a **statistical approximation**.
- **Argument of the Poverty of the stimulus** re-adapted:
 - Today's best LLMs are trained on vastly more data than a child is exposed to.
 - Estimation: less than 100M words for a 12yo human.
 - Hence, the models will necessarily be un-human-like (a la Chomsky).
 - *BabyLM challenge*: to train language models in low-resource data settings.
 - Where amount of linguistic input resembles the amount received by human language learners.

- Find the best architecture who perform best given the same initial conditions.
- Another objective of the challenge is to democratize research.

1.7.3.2 Humans learn proactively in ways language don't [Han & McClelland, 2025]

- LLMs now match or exceed average human performance.
 - Maybe LLMs may soon become genuine examples of general intelligence or unified theories of human cognition.
 - This paper tries to answer these grandiose claims that appear in the literature.
- Task considered (for humans and LLMs): **chain rank task**.
 - **Train phase:**
 - Participants learn about this chain from individual sentences describing the relations of neighboring pairs of entities.
 - The participant's goal is to learn whether the correct completion for each sentence is *true* or *false*.
 - This allows the CRT to be presented to humans and LLMs with maximally similar inputs and outputs.
 - eg Five entities (persons: Kennedy, Roosevelt, Lincoln, Einstein and Presley), considered as pairs.
 - The relation is *better at chess, worse at chess*.
 - **Test phase:**
 - Participants are required to make judgments of entity pairs from the underlying chain, including:
 - The original learning items.
 - New items that requires participants to make either:
 - **Reverse generalizations:** $A > B \rightarrow B < A$.
 - **Transitive generalization:** $A > B \wedge B > C \rightarrow A > C$.
- Performed on LoRA, ICL and ICL+CoT (aligned or unaligned).

Results

- LLM learns to generalize much more slowly than humans.
- While on learned relationships LLMs are optimal, they are worse in reverse and n -hop generalization.
- Via a post-experimentation form provided to all participants, it emerges:
 - Humans rely on one of two strategies for representing and reasoning about the learning items:
 - *Ranker strategy:*
 - Actively recording the learning items into a single representation of entity ranking.
 - Ranking that can be recalled for any test item.
 - *Non-ranker strategy:*
 - Memorizing each learning item independently.
 - Recalling only relevant learning items for each item.
 - During the testing stage, rankers achieved higher accuracy but also make faster judgments.
- → Stronger performance is characterized by the **proactive formation of efficient representations that aid learning and generalization**.
 - Thesis: human intelligence is explained by **proactivity** and **curiosity** (and not only memorization).

1.7.3.3 Atoms of recognition in human and computer vision [Ullman, Assif, Fataya, Harari, 2016]

- CNN have been proposed several times as a very close model to how human brain process visual information.
 - Goal: see how far this analogy can go.
- Central problem in the study of human vision:
 - Discovering **visual features** and **representation** used by the brain to recognize object.
 - NN models rival human performance on image classification.
 - Investigates commonalities in representation and learning process between current NN and human visual system.
- **Minimal recognizable configurations** (MIRC).
 - Small portion of images from which a human can recognize the subject of the picture.
 - By rendering smaller (sub-MIRCs), no recognition is possible.
 - Since they are small, they **reveal features** that are important in recognition.
 - The role of features is revealed uniquely at the minimal level (essential contribution).
- Human and networks models are tested to check if they **perform in the same way**.
 - If the human vision uses the same features to understand the content of an image.

Results

- In humans, better recognition of MIRC than in network models.
- In humans, **sharp decrease between MIRC and sub-MIRC recognition.**
 - Not in neural networks models.
 - All MIRC features are necessary for human recognition.
- Networks models:
 - They are not good at recognizing minimal images.
 - No decrease between MIRC and sub-MIRC.
- Therefore **different features are used (in different ways)** by humans and NN.

Internal interpretation

- Additional limitation is the ability to perform a detail internal interpretation of MIRC.
- Human can consistently recognize multiple components internal to the MIRC.
- Such internal interpretation is beyond the capacities of current NN.
 - It can contribute to accurate recognition.
 - Because a false detection could be rejected if it does not have **expected internal interpretation**.

Feed forward and top-down processes

- Features are used in the visual system as a part of the cortical feed-forward process or by a top-down process.
 - Top-down process is currently missing from the purely feed-forward computational model.
 - While in human, there is a **combination of bottom-up and top-down processes**.
- Top-down processes are likely involved:
 - Detailed interpretation appears to require features and interrelations that are class-specific.
- Two main stages:
 - Initial activation of class candidates, which is incomplete and with limited accuracy.
 - The activated representations then trigger the application of class-specific interpretation and validation processes.
 - Which recover richer and more accurate interpretation of the visible scene.

1.8 Representation learning

1.8.1 Representations

- **Representation learning** underpins all current developments in DL.
 - It enables the sharing of statistical power across tasks.
 - Particularly beneficial when dealing with multiple modalities or domains.
 - Also facilitate **knowledge transfer** to tasks with few/zero examples are available, but a task representation exists.
 - Changing representation can make a problem very difficult or very easy.
- A **good representation** is one that makes a subsequent learning task easy.
 - FFN training by supervised learning can be seen as performing a sort of representation learning.
 - Most part of the network is used to **learn better representations** (more and more abstract).
 - And then a simple classifier is used at the top of the network.
- Learning objective can be tailored to force the representation to have some nice properties.
 - Most representation learning problems face a trade-off between:
 - Preserving as much information about the input as possible.
 - Attaining nice properties (eg *independence* of the features).

1.8.1.1 Properties of good representations

- An ideal representation is one in which the **features** within the representation correspond to the **underlying causes** of the observed data.
- Two key properties: they are often **distributed** and **disentangled**.
- Local vs. distributed representations:
 - A **local representation** uses a single, discrete value for each concept (eg one-hot encoding).
 - Problem: all concepts are equally distant (no similarity is learnt).
 - This representation is **sparse** and **inefficient**.
 - A **distributed presentation** describes each concept using multiple and continuous features.
 - Each feature is involved in representing many concept.
 - It creates a **rich similarity space** (semantically close concepts are close in distance).
 - -ex DR help in **representing an exponential number of regions using** a linear number of parameters.
 - This doesn't automatically translate in an exponential advantage for the classification algorithm.
 - The VC dimension of a NN of linear threshold units is only $O(w \log w)$ (w number of weights).
 - The **Vapnik-Chervonenkis** dimension is a measure of the capacity of a model (hypothesis) space.
 - It is used to deduce bounds on the generalization capabilities of ML.
 - Keeping the VC dimension of the hypothesis space small is usually a sensible goal.
 - Despite the representation being able to distinguish between many zones, not all zones can be represented.
 - DR can be **interpretable**, but this is **not the norm**.
- In a **disentangled representation** single features (or single features directions in the feature space) correspond to **single distinct factors** of variation in the data.
 - It allows to reason about and manipulate high-level concepts independently.

1.8.2 Transfer Learning and Domain Adaptation

- TL: Where what has been learned in one setting (distribution P_1) is **reused** to improve generalization in another setting (P_2).
 - In domain adaptation, the task is the same but, eg, the data distribution may differ.
- Generalization of the idea of greedy PT.
 - Where the transferred representations were between unsupervised and supervised tasks.

1.8.2.1 Transfer Learning

- The learner must perform two or more different tasks.
 - But it is assumed that many of the factors that explain variations in P_1 are relevant in P_2 .
- **Depth** seems to be crucial in this process.

- As an architecture uses deeper representations, the learning curve on new categories of transfer setting P_2 becomes much better.
- eg Image recognition on cancerous cells.
 - Strategy: fine-tune a general model (like ImageNet) on smaller medical dataset.
- TL and DA assume there is a common (**shared**) **representation** that explain variations in the input data.
 - These variation are then adapted to the various tasks by different output layers.
 - eg Image recognition.
 - The shared part is the initial part of the network.
- But there are situations, where what is shared between different tasks is not the semantics of the input.
 - But the semantics of the output and inputs need to be adapted to be compatible with that.
 - The shared part is the last part of the network.
 - eg Speech recognition.

1.8.2.2 Domain adaptation

- The task remains the same between each setting, but the input distribution is slightly different.
 - eg Sentiment analysis (from movie reviews to product reviews).
- Create a representation that captures the underlying patterns in the data while **ignoring domain-specific nuances**.
- *Unsupervised Domain Adaptation by Backpropagation* [Ganin & Lempitsky, 2015]:
 - The source is MNIST while the target is MNIST-M (MNIST with different background).
 - A part of the network is used to learn a shared representation.
 - To this shared representation two classification heads are attached:
 - The first one classify digits.
 - The second one classify domains.
 - During training, each time the second head succeeds, the shared network is penalized.
 - The gradient is inverted (**gradient ascent**).

1.8.2.3 Concept drift

- When the concept underlying the data distribution shifts in time.
 - A model learnt at a given point in time need, then, to be updated to take into account the drift.
- In this case too, data from a given setting (previous than the shift) is exploited to get an advantage in the new setting.
- The representation part is not changed, the classification part is.

1.8.2.4 One-shot and zero-shot learning

- **One-shot learning:**
 - An extreme case of transfer learning, where only a single labeled example is given for the new setting.
 - An already-learnt representation (eg a manifold of known class) is used.
- **Zero-shot learning:**
 - One needs to adapt to a new setting without seeing any labelled example.
 - It can be seen as including three random variables:
 - The traditional input x .
 - The traditional outputs y .
 - An additional random variable describing the task T .
 - The model is trained to model $p(y | x, T)$.

1.8.3 Greedy layer-wise Unsupervised Pre-training

- Layers are trained in an **unsupervised** way (supervised ways also available).
- Supposed to be **only a first step before joint training** starts.
- Instrumental to the revival of deep NN.
 - Classical deep nets had trouble in propagating information.
 - Specifically, deep networks had problems with **vanishing** and **exploding** gradients.
 - *Solution:* break up the training into the training of smaller networks.
- Main idea:
 - Each layer is trained using **unsupervised** learning, taking the output of the previous layers.
 - And producing as output a new representation of the data, whose distribution is hopefully simpler/better.
 - Most often the procedure relies on a single-layer representation learning.

```

def greedy_layer_wise_unsupervised_pretraining(X,y)
    f = lambda z: z          # identity function
    data = X                  # repr 0

    for k in range(m):
        f_k = UnsupervisedLearn(data)
        f = lambda x: f_k(f(x))
        data = f_k(data)      # repr k+1

    if fine_tuning:
        f = FineTune(f, X, y)

    return f

```

Figure 6: Greedy layer-wise Unsupervised Pre-training pseudo-code.

- eg RBM, single layer autoencoders, other models that learn latent representations.
- y is used only in the **fine-tuning** phase.
- In 2006 it was shown that it could be used to find **good initialization** to train deep architectures.
 - Nowadays, other techniques are components are used (ReLU, He, Xavier initialization, Adam dropouts, etc).
 - Nowadays, unsupervised is largely abandoned, except in the field of NLP.
 - DL techniques based on supervised learning, regularized with dropout or batch normalization usually outperform UPT.
 - In many fields (NLP and CV), **self-supervised learning** (kind of UPT) is widely used.
 - But this approach was the initial demonstration that such training was indeed achievable.
- Variants:
 - Can also be used as initialization for deep unsupervised models.
 - Greedy layer-wise Supervised Pre-training.
 - Simultaneous supervised and unsupervised learning.
 - Allows incorporating the constraints imposed by the output layer from the outset.

1.8.3.1 Main ideas beneath Unsupervised Pre-training

- UPT acted as a **powerful regularizer** and provided much better **initialization** for network weights.
- Two main ideas are important to understand why and when unsupervised pre-training work:
 - The **choice of initial parameters** can have a significant **regularizing effect** on the model.
 - Initially it was assumed that this would result in approaching a local minimum.
 - Training is longer without pre-training.
 - PT and non-PT models start and stay in different regions of function spaces.
 - All trajectories of a given type initially move together and then diverge.
 - This suggests that each trajectory moves into a different apparent local minimum.
 - Nowadays **local minima** are no longer considered a serious problem.
 - Critical points are more likely **saddle points** rather than spurious local minima.
 - Local minima **concentrate near the global optimum**.
 - **Learning about input distribution** can help to learn the mapping from inputs to outputs.
 - In many tasks, features learned in unsupervised stage can be useful also in supervised stage.
 - Unsupervised pre-training often helps when the **original representation is poor**.
 - eg Learning word embeddings instead of using one-hot vector.

1.8.3.2 Advantages and disadvantages of Unsupervised Pre-training

- Advantages:
 - It is useful to view unsupervised PT as *regularizer*:
 - When the number of examples is small.
 - When the number of unlabeled examples is large.
 - When the **function to be learned is extremely complicated**.
 - In contrast to other regularizers, UPT **doesn't force the learnt function to be simple**.
 - Rather it helps in discovering feature functions that are helpful in the unsupervised task.
 - UPT can be a more appropriate regularizer if the true underlying functions are:
 - **Complicated and shaped by regularities** of the input distribution.
- Disadvantages:
 - As a regularization technique, UPT has the problem that it is **difficult to calibrate**.

- In most regularization techniques, there is a single parameter to set regularization strength.
- With unsupervised training, either the network is initialized using PT or it is not.
- The hyper-parameters of the PT phase needs to be adjusted and this can be extremely slow.

1.8.4 Self-supervised learning (SSL)

- The modern successor to UPT, it learns from unlabeled data by creating a **pretext** task.
 - Pretext: a reason given in justification of a course of action that is not the real reason.
- Core idea: hide/change **some part of the input** and train the network to predict the original input.
 - The supervision signal comes from the data itself, not from human labels.
 - This forces the model to learn a rich semantic representation of the data.
- eg Pretext tasks in Vision:
 - Image inpainting: predict a missing patch of an image.
 - Colorization: predict the color version of a greyscale image.
 - Jigsaw puzzle: learn to assemble shuffled image patches correctly.
- eg Pretext tasks in Language:
 - **Masked Language modeling (MLM)**: predict a randomly masked word in a sentence.
 - The core idea behind models like BERT.

1.8.4.1 SSL via contrastive learning

- A dominant family of SSL methods where the goal is to learn an embedding space where:
 - Semantically **similar** examples are pulled **together**.
 - Semantically **dissimilar** examples are pushed **apart**.
- Procedure:
 - Take an example (the *anchor*).
 - Create two different random augmentation of it, resulting in a **positive pair**.
 - Take examples from the rest of the batch, the **negative examples**.
 - Train the model to:
 - Maximize the similarity of the positive pair's representations.
 - Minimize the similarity to all negative examples.
- def **Information Noise Contrastive Loss**: $\mathcal{L}_{\text{InfoNCE}} = - \log \frac{\exp(\text{sim}(z_i, z_i^+)/\tau)}{\sum_j \exp(\text{sim}(z_i, z_j)/\tau)}$.
 - Key contrastive objective where:
 - sim is cosine similarity, with temperature τ which sharpens similarity scores.
 - (z_i, z_i^+) is the positive pair, all the other z_j in batch are negatives.
 - Goal: maximize similarity of positives, minimize similarity to negatives.
- CL requires **negative samples** to avoid the trivial collapse where all representation become identical.
 - Positive pairs suffice to introduce a notion of similarity.
 - But **many negatives** are crucial to **define dissimilarity and to prevent collapse**.
 - Negatives encourage the model to **spread representations apart**.
 - So that different inputs occupy distinct regions of the embedding space.

CL systems

- **SimCLR**:
 - It popularized contrastive learning in vision by demonstrating that with the right choices, simple CL can rival supervised learning.
 - Key ideas:
 - **Strong data augmentation** (random crop + resize, color jitter, gaussian blur, solarization).
 - **Projection head**: a small MLP head improves representation quality.
 - **Large batch size** (4096-8192): more negatives → better **contrastive signal**.
 - It showed that **very large batch sizes** (i.e. negative samples) allows for the **best results**.
- **Momentum Contrast** (MoCo):
 - It removes SimCLR's dependency on massive batch sizes.
 - Key ideas:
 - Maintain a **queue** of thousands of negative embeddings.
 - Use a **momentum encoder** to produce consistent keys.
 - Queue provides many negatives without huge GPU memory.
 - Effect: Moco allows contrastive learning with batch size ≈ 256 .
 - **Momentum encoder (ME)**:

- MoCo requires the ME to be updated slowly so that the produced **feature representations** (keys) **change gradually** over time.
- This stability is crucial because MoCo maintains a queue of **cached key embeddings** generated from past batches.
- If the ME changed too quickly, the embeddings stored in the queue would **become stale**.
 - And no longer consistent with the current model, which would weaken or even break the CL objective.
- The slowly updated ME ensures:
 - Cached key remain compatible with current queries.
 - The queue provides a large, consistent set of negatives.
 - Training remains stable and effective.

1.8.5 Semi Supervised Learning (SemiSL)

- A possibly large amount of examples is given, but only few of them are labelled.
 - The task is to exploit the unlabelled examples to better perform on the supervised task.
- **Causality Hypothesis:** an **ideal representation** is one in which the **features** within the representation correspond to the **underlying causes** of the observed data.
 - It underlies a large deal of research motivated by the idea that disentangling the causal factors in $p(\mathbf{x})$ could be a good step for learning $p(\mathbf{y} | \mathbf{x})$.
 - This motivated SemiSL approach.
 - But sometimes this fails since unsupervised learning of $p(\mathbf{x})$ is of no help to learn $p(\mathbf{y} | \mathbf{x})$.
 - eg When there is no regularities in the data.
 - Other times \mathbf{y} is among the salient causes of $p(\mathbf{y})$.
 - In these cases learning $p(\mathbf{x})$ can be very useful.
 - eg When there is some regularities underlying the data.

1.8.5.1 Causal factors

- If \mathbf{h} represents all **factors causing** \mathbf{x} , and \mathbf{y} is assumed to be related to one of them:
 - The generative process can be conceived as $p(\mathbf{h}, \mathbf{x}) = p(\mathbf{x} | \mathbf{h})p(\mathbf{h})$.
 - The data has marginal probability: $p(\mathbf{x}) = \sum_{\mathbf{h}} p(\mathbf{h}, \mathbf{x}) = \sum_{\mathbf{h}} p(\mathbf{x} | \mathbf{h})p(\mathbf{h}) = \mathbb{E}_{\mathbf{h}}[p(\mathbf{x} | \mathbf{h})]$.
 - The best possible model of \mathbf{x} is the one that uncovers the above *true* structure.
 - With \mathbf{h} as a latent variable that explains the observed variations in \mathbf{x} .
- **Most observations are formed by an extremely large number of underlying causes.**
 - The brute force approach of encoding *all* possible factors of variations does not work.
- It is then necessary to decide what to encode into \mathbf{h} .
 - To find a strategy to guide the network to keep only the **relevant** part of \mathbf{h} .
 - Two main strategy:
 - Use a supervised signal to guide the unsupervised process.
 - Use a much larger \mathbf{h} when using only unsupervised learning.
- Another emerging strategy is to **change the definition of what is salient**.
 - Historically, one would optimize against a fixed criterion (often similar to MSE), but this is problematic.
 - eg MSE applied to pixels of images implicitly specifies that a cause is only relevant if affect these brightness of a large number of pixels.
 - Other definitions of salience are possible.
 - In GAN, it's defined implicitly by a game played by an encoder trying to fool a discriminator.
- **Causal factors are robust.**
 - With different domains or task natures, or temporal non-stationarity, the **causal mechanisms remain invariant**.
 - While the marginal distribution over the underlying causes can change.

1.9 Autoencoders

1.9.1 Autoencoders

- Unsupervised learning tools used to **improve supervised networks**.
 - eg Image colorization, increase resolution, image inpainting, machine translation, etc.
- A NN trained to attempt to copy its input to its output via a representation \mathbf{h} built by its hidden layers.
 - Encoder: $\mathbf{h} = f(\mathbf{x})$.
 - Decoder: $\mathbf{r} = g(\mathbf{h})$.
 - Not expected to faithfully copy every input to its output.
 - Instead, they are forced to prioritize which aspects of the input should be preserved.
- AE are a classic example of **self-supervised learning**, where the **supervised signal** is the input itself.
- Traditionally used for dimensionality reduction or feature learning.
 - Today they are used for **generative models** due to connections established with **latent variable models**.
 - **Latent variable model:**
 - The model has both observed \mathbf{x} and latent variables \mathbf{h} .
 - The goal is to learn the distribution $p(\mathbf{x}, \mathbf{h})$.
 - Similar to what happens in representation learning.
- Goal:
 - Only copying the input to the output is not useful.
 - Instead, training the AE should result in \mathbf{h} taking on useful properties.
 - Approach: force \mathbf{h} to have a smaller dimension than \mathbf{x} .

1.9.1.1 Dimensionality reduction

- Dimensionality reduction has been the first motivation to study autoencoders.
- Learning AE that map input to lower dimensional output.
 - Many tasks can be more easily solved, and models of smaller spaces consume less memory and runtime.
 - It often places semantically related examples nearby, helping **generalization**.
- **Semantic hashing:**
 - Store all database entries in a hash table mapping binary code vectors to entry.
 - The hash table allows retrieval of similar elements (they share the same binary code).
 - Swap single bits in the code to search for slightly less similar elements.
 - To produce these binary codes, one typically use sigmoid units in the final layer and train them to saturation.
 - By injecting additive noise just before the sigmoid.

1.9.1.2 Undercomplete Autoencoders

- An autoencoder whose code dimension is less than the input dimension ($\text{len}(\mathbf{x}) \gg \text{len}(\mathbf{h})$).
- The learning process for UAE is usually the minimization of a loss function, $L(\mathbf{x}, g(f(\mathbf{x}))$.
 - Where L is a loss function penalizing $g(f(\mathbf{x}))$ for being dissimilar from \mathbf{x} (eg MSE).
 - When the decoder is linear and L is MSE, a UAE learns to span the same subspace as PCA.
 - PCA finds the mapping $f(\mathbf{x}) = \arg \min_{\mathbf{h}} \|\mathbf{x} - g(\mathbf{h})\|_2$, imposing g as a linear model.
 - When $L = \|\cdot\|_2$ and the reconstruction layer is modeled by linear units, the two approaches can solve the same problem.
- Data to be reconstructed is assumed to live on a linear manifold.
 - The best approach must be determined and why.
 - PCA is easier so when its requirements are granted.
 - An AE is better with a highly non-linear.

1.9.2 Regularized Autoencoders

- The difference between $\text{len}(\mathbf{x})$ and $\text{len}(\mathbf{h})$ determines how much the AE is forced to learn only the *most relevant* part to the input.
 - If the size of \mathbf{h} is too big, the AE could simply learn the identity function.
- RAE uses a loss function that encourages the model to have other properties besides I/O copying ability.
- Properties exploited for regularization:
 - Sparsity of the representation → **Sparse autoencoders**.
 - Robustness to noise or missing inputs → **Denoising autoencoders**.

- Smallness of the derivative → **Contractive autoencoders**.

1.9.2.1 Sparse Autoencoders

- The training criterion involves a sparsity penalty $\Omega(\mathbf{h})$ on the code layer \mathbf{h} , plus the reconstruction error.
 - $\Omega(\mathbf{h}) = \lambda \|\mathbf{h}\|_1 = \lambda \sum_i |h_i|$.
- def **Sparse Autoencoders Loss**: $L(\mathbf{x}, g(f(\mathbf{x}))) + \Omega(\mathbf{h})$.

Generative models

- Training a SAE, by minimizing reconstruction error plus a sparsity penalty, is approximately equivalent to maximizing the model's likelihood.
 - $\min L(\mathbf{x}, g(f(\mathbf{x}))) + \Omega(\mathbf{h}) \approx \min -\log p_{model}(\mathbf{x})$.
 - Then **sparsity penalty is not just a regularizer**.
 - It naturally arises from **modeling the joint distribution** over data and latent variables when a **Laplace prior** is assumed on the latent factors.
 - Adopting a probabilistic view reveals that **SAE are effectively learning a generative model**.
- **Generative models**: a model that learns the joint distribution $p(\mathbf{x}, \mathbf{y})$ of the data.
 - After the model is trained, it can generate new data points of a given class by sampling from it.
 - Sometimes \mathbf{y} is not an observable variable.
 - \mathbf{y} is then called a **latent variable** and usually denoted by \mathbf{h} .
 - GM differ from a **discriminative model**.
 - Since DM learn to approximate the conditional probability $p(\mathbf{y} | \mathbf{x})$.
 - A DM cannot therefore generate new data points.
- Given a GM with visible variable \mathbf{x} and latent variables \mathbf{h} .
 - $p_{model}(\mathbf{h}, \mathbf{x}) = p_{model}(\mathbf{h}) p_{model}(\mathbf{x} | \mathbf{h})$.
 - The **likelihood** of \mathbf{x} given the model can be computed as the marginalization of the latent variables \mathbf{h} .
 - **Likelihood**: $p_{model}(\mathbf{x}) = \sum_{\mathbf{h}} p_{model}(\mathbf{h}, \mathbf{x})$.
 - **Log-Likelihood**: $\log p_{model}(\mathbf{x}) = \log \sum_{\mathbf{h}} p_{model}(\mathbf{h}, \mathbf{x})$.
 - AE can be seen as approximating this sum with a **point estimate for just one highly likely value** for \mathbf{h} .
 - By taking the logarithm of $p_{model}(\mathbf{x})$ the log-likelihood of \mathbf{x} under the model is obtained.
 - $\log p_{model}(\mathbf{x}) = \log \sum_{\mathbf{h}} p_{model}(\mathbf{h}, \mathbf{x})$.
 - Given an $\tilde{\mathbf{h}}$ generated by the decoder:
 - $\log p_{mod}(\mathbf{x}) = \log \sum_{\mathbf{h}} p_{mod}(\mathbf{h}, \mathbf{x}) \approx \log p_{mod}(\tilde{\mathbf{h}}, \mathbf{x}) = \log p_{mod}(\tilde{\mathbf{h}}) + \log p_{mod}(\mathbf{x} | \tilde{\mathbf{h}})$.

Loss decomposition

- The Loss can be divided in a **reconstruction loss** ($L(\mathbf{x}, g(f(\mathbf{x})))$) and a **sparsity loss** ($\Omega(\mathbf{h})$).
- **Reconstruction loss**:
 - By minimizing the reconstruction loss, the AE is learning to approximate the log-likelihood of the data.
 - By minimizing it, the AE is learning to approximate the conditional distribution $p_{model}(\mathbf{x} | \mathbf{h})$.
 - It is trying to find the \mathbf{x} that is the best reconstruction given \mathbf{h} .
 - That is, solving $\arg \max_{\mathbf{x}} p_{model}(\mathbf{x} | \mathbf{h}) = \arg \min_{\mathbf{x}} -\log p_{model}(\mathbf{x} | \mathbf{h})$.
 - A minimization problem is preferred for optimization.
 - While logarithms are preferred to avoid numerical issues thanks to their additivity.
- **Sparsity loss**:
 - By setting $\Omega(\mathbf{h}) = \lambda \sum_i |h_i|$ minimizing the sparsity term correspond to maximizing the log-likelihood of the $p(\mathbf{h})$ term assuming a Laplace prior over each component of \mathbf{h} independently.
 - Using the L_1 norm of \mathbf{h} is known to **encourage sparsity**.
 - Since h_i are independent, $p(\mathbf{h}) = \prod P_{model}(h_i)$.
 - By taking the logarithms of it, \prod turns into a \sum .
 - If $p_{mod}(h_i) = \frac{\lambda}{2} e^{-\lambda|h_i|}$ then $-\log p_{mod}(h_i) = \sum_i (\lambda|h_i| - \log \frac{\lambda}{2}) = \Omega(\mathbf{h}) + \text{const}$.
 - λ can be treated as an hyper-parameter.
 - const , depending only on λ , is not optimized during learning and can be ignored.

Training a SAE

- Training the network correspond to minimize the **negative log-likelihood of the data under the model**.
 - Obtained by the results related to the loss decomposition.

- def **Sparse Autoencoder Training**: $\arg \min_{\theta} [-\log p_{model}(\mathbf{x})] \approx \arg \min [-\log p(\mathbf{x} | \tilde{\mathbf{h}}) - \log p(\tilde{\mathbf{h}})] \approx \arg \min_{\theta} [L(\mathbf{x}, g(f(\mathbf{x}))) + \Omega(\mathbf{h})]$.
- The sparsity penalty then is not a regularization term.
 - It's just the consequence of modeling the joint distribution taking into account the latent variables and assuming them to have a Laplace prior.
- This view provides different motivations for training a SAE:
 - It is a way of approximately training a generative model (eg to generate new examples).
 - Learnt features are useful since they describe the latent variables that explain the input.

Generating new data with SAE

- In principle, a learned model can be used to generate new data.
 - By sampling from the prior $p_{model}(\mathbf{h})$ and then reconstructing \mathbf{x} using the decoder $g(\mathbf{h})$.
- During the training process, the decoder has only been trained to reconstruct data points close to Tr examples.
 - In practice, by naively sampling from prior, it is likely to get \mathbf{h} values very far from the ones seen during training.
 - Yielding very poor samples.
- SAE can be used to **generate new data points**, but one should **calibrate the sampling process**.
 - To ensure that the sampled \mathbf{h} values are close to the ones seen during training.

1.9.2.2 Denoising Autoencoders

- More in general, one can view both the encoder and the decoder as modeling some distribution.
 - $p_{encoder}(\mathbf{h} | \mathbf{x}) = p_{model}(\mathbf{h} | \mathbf{x})$.
 - $p_{decoder}(\mathbf{x} | \mathbf{h}) = p_{model}(\mathbf{x} | \mathbf{h})$.
 - In general, the E and D distribution are not necessarily conditional distributions compatible with a unique joint distribution $p_{model}(\mathbf{x}, \mathbf{h})$.
 - Training E and D as a **denoising autoencoder** will tend to make them compatible asymptotically.
 - With enough capacity and examples.
 - DAE guarantees that those p will be compatible, when previously it wasn't guaranteed.
- Rather than adding a penalty Ω to the cost function, they changes the reconstruction error of the CF.
- def **Denoising Autoencoders Loss**: $L(\mathbf{x}, g(f(\tilde{\mathbf{x}})))$.
 - Where $\tilde{\mathbf{x}}$ is a copy of \mathbf{x} that has been corrupted by some form of noise.
 - DAE must therefore undo this corruption rather than simply copying their input.
- It has been shown that denoising training forces f and g to implicitly learn the structure of $p_{data}(\mathbf{x})$.

DAE Stochastic view

- A corruption process $C(\tilde{\mathbf{x}} | \mathbf{x})$ that produces corrupted samples is introduced.
- The AE then learn a **reconstruction distribution**.
 - def **Reconstruction distribution**: $p_{reconst}(\mathbf{x} | \tilde{\mathbf{x}}) = p_{decoder}(\mathbf{x} | \mathbf{h}) = g(\mathbf{h})$ and $\mathbf{h} = f(\tilde{\mathbf{x}})$.
- Typically, GD minimization on the negative log likelihood $-\log p_{decoder}(\mathbf{x} | \mathbf{h})$ can be performed.
 - As long as the encoder is deterministic, the whole AE can be **trained end-to-end using SGD**.
 - The DAE can be seen as performing SGD on $-\mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} \mathbb{E}_{\tilde{\mathbf{x}} \sim C(\tilde{\mathbf{x}} | \mathbf{x})} \log p_{decoder}(\mathbf{x} | \mathbf{h} = f(\tilde{\mathbf{x}}))$.

Manifolds

- Many ML algorithms exploit the idea that data concentrates around **low-dimensional manifolds**.
 - Or a small set of such manifolds.
 - AE take this idea further and aim to learn the structure of the manifold.
- DAE can be seen as **approximating a vector field**.
 - The corruption process moves the examples away from the lower dimensional manifold.
 - The AE is **learning to project back** these examples to the manifold.
 - The corrupted example is projected to the manifold, obtaining \mathbf{h} .
 - \mathbf{h} can then be used to obtain the original example.
 - To assume that the lower dimensional manifold is *smooth* has been proved empirically right.
 - The DAE learns a **vector field that provides the fastest path to the manifold**.
- **Manifold tangent planes**:
 - An important characterization of a manifold is the set of its tangent planes.
 - At a point \mathbf{x} on a d -dimensional manifold:

- The tangent plane is given by d basis vectors that span the local directions of variation allowed on the manifold.
- These local directions specify how one can **change x infinitesimally while staying on the manifold**.
- AE and manifolds:
 - If the data generating distribution concentrates near a low-dimensional manifold, this yields representation that implicitly capture a local coordinate system for this manifold.
 - Only the variations tangent to the manifold around x need to correspond to changes in $\mathbf{h} = f(\mathbf{x})$.
 - Hence the encoder learns a mapping from the input space \mathbf{x} to a representation space.
 - A mapping that is **only sensitive to changes along the manifold directions**.
 - But that is **insensitive to changes orthogonal to the manifold**.

1.9.2.3 Contractive Autoencoders

- def **Contractive Autoencoders Loss**: $L(\mathbf{x}, g(f(\mathbf{x}))) + \lambda \sum_i \|\nabla_{\mathbf{x}} h_i\|^2$.
 - As in SAE, optimize a regularized objective $L(\mathbf{x}, g(f(\mathbf{x}))) + \Omega(\mathbf{h}, \mathbf{x})$, but with a different form of Ω .
 - This force the model to learn a function that does not change much when \mathbf{x} changes only slightly.
 - With a very high λ , a constant plane is obtained by minimizing the derivative.
- As a side effect of the contractive penalty, the encoder learn to map a **neighborhood of inputs** onto a **smaller neighborhood of outputs**.
 - Without any other competing force (or with a very high λ) the penalty would drive f to be learnt as a constant function.
 - Which maps the whole input space onto a single point.
- CAE is **contractive only locally**.
 - If \mathbf{x} and \mathbf{x}' are different enough, they can be mapped to point very far apart than the original points.
- In the limit of small Gaussian input noise, the **denoising reconstruction** error is equivalent to a **contractive penalty on the reconstruction function**.
 - DAE make the **reconstruction function** to resist to small errors in the input (decoder part).
 - CAE make the **feature extraction function** to resist small perturbations of the input (encoder part).

Learning manifolds

- Regularized AE learns manifolds by balancing two opposing forces.
 - In case of CAE, these two forces are reconstruction error and the contractive penalty $\Omega(\mathbf{h})$.
- The compromise between these two forces yields an AE whose derivatives $\frac{\partial f(\mathbf{x})}{\partial \mathbf{x}}$ are mostly tiny.
 - Only a **small number of hidden units** (corresponding to a small number of directions in the input) may have **significant derivatives**.
- Directions \mathbf{x} with large $J_{\mathbf{x}}$ rapidly change \mathbf{h} .
 - **Such directions are penalized during learning** (due to contractive penalty).
- So the ones *surviving* are those really necessary to model the data.
 - And hence they are likely to be directions which **approximate the tangent planes of the manifold**.
 - By loosing those directions, the input cannot be reconstructed, so those are kept.

1.9.3 Variational Autoencoders

- An AE whose **encodings distribution is regularised** during the training.
 - In order to ensure that its latent space has **good properties allowing to generate** some new data.
 - The regularization method is based on *variational inference*, an inference method popular in statistics.
- AE are not good generative models.
 - To do generation:
 - After training the encoder is discarded.
 - A random \mathbf{h} is picked and served to the decoder.
 - The decoder will then generate an output.
 - While this simple schema is intuitive, it does not work well in practice.
- In seen AE, encoding networks are not regularized.
 - The semantic *areas* are concentrated and well-spaced.
 - When an example (eg the randomized one) that is not nearby any of those area, will return a meaningless output.
 - Encoding networks need to be **regularized**.
 - Larger semantic *areas* all nearby between each other.
- To impose this kind of regularization, images are **encoded into a distribution**.
 - During training, the decoder is trained to decode points sampled from these distributions.

1.9.3.1 Probabilistic model

- A probabilistic model where a latent variable \mathbf{h} is sampled from a distribution $p(\mathbf{h})$.
 - And then \mathbf{x} is sampled from $p(\mathbf{x} | \mathbf{h})$.
- It is also assumed:
 - $p(\mathbf{h}) \sim N(\mathbf{0}, \mathbf{I})$.
 - $p(\mathbf{x} | \mathbf{h}) \sim N(g(\mathbf{h}), c\mathbf{I})$ (centered around the reconstruction, isotropic distribution).
- Even under these assumptions, the inference problem is **intractable** due to the normalization factor in the denominator.
 - $p(\mathbf{h} | \mathbf{x}) = \frac{p(\mathbf{x}|\mathbf{h})p(\mathbf{h})}{p(\mathbf{x})} = \frac{p(\mathbf{x}|\mathbf{h})p(\mathbf{x})}{\int p(\mathbf{x}|\mathbf{u})p(\mathbf{u}) d\mathbf{u}}$.
 - $p(\mathbf{x})$ is unknown, it's therefore approximated with an integral (that is also unknown, intractable).
 - Variational inference is a technique allowing to solve this problem.
 - By **approximating the hard to compute probability with a simpler one**.
 - Specifically, a distribution $q_{\mathbf{x}}(\mathbf{h}) \sim N(f_1(\mathbf{x}), f_2(\mathbf{x})) \approx p(\mathbf{h} | \mathbf{x})$ is used.
 - $p(\mathbf{h} | \mathbf{x})$ is approximated using a Gaussian.
- *Goal:* to find f_1 and f_2 allowing the best possible approximation of $p(\mathbf{h} | \mathbf{x})$.
 - $\arg \min_{f_1, f_2} KL(q_{\mathbf{x}}(\mathbf{h}) \| p(\mathbf{h} | \mathbf{x})) = \dots = \arg \max_{f_1, f_2} E_{\mathbf{h} \sim q_{\mathbf{x}}} [-\frac{\|\mathbf{x} - g(\mathbf{h})\|^2}{2c}] - KL(q_{\mathbf{x}}(\mathbf{h}) \| p(\mathbf{h}))$.
- *Goal:* to find the decoder function g that maximises the likelihood of $p(\mathbf{x} | \mathbf{h})$ under assumption $\mathbf{h} \sim q_{\mathbf{x}}(\mathbf{h})$.
 - def **VAE optimization problem:** $\arg \min_{f_1, f_2, g} E_{\mathbf{h} \sim q_{\mathbf{x}}} [\frac{\|\mathbf{x} - g(\mathbf{h})\|^2}{2c}] + KL(q_{\mathbf{x}}(\mathbf{h}) \| p(\mathbf{h}))$.
 - With $q_{\mathbf{x}}(\mathbf{h}) \sim \mathcal{N}(f_1(\mathbf{x}), f_2(\mathbf{x})) \equiv \mathcal{N}(\mu_{\mathbf{x}}, \sigma_{\mathbf{x}})$ and $p(\mathbf{h}) \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.
 - It has the form a **regularized network** (with a reconstruction term and a regularization term).
 - Easily solvable via GD.

Reparametrization trick

- While optimization is easily done via GD, the extraction of a point in the embedding space is **not differentiable**.
 - Randomly picking a point from a distribution is not a differentiable operation.
- Sampling from $N(\mu_{\mathbf{x}}, \sigma_{\mathbf{x}})$ is equivalent to sampling $\zeta \sim N(\mathbf{0}, \mathbf{I})$ and then $\mathbf{h} = \sigma_{\mathbf{x}} + \zeta \mu_{\mathbf{x}}$.
 - Points are rescaled and shifted (standard translation from standard Gaussian to other normal distributions).
 - A point is not extracted from $N(\mu_{\mathbf{x}}, \sigma_{\mathbf{x}})$, which depends on $\mu_{\mathbf{x}}$ and $\sigma_{\mathbf{x}}$, which are part of the network.
 - But is extracted from $N(\mathbf{0}, \mathbf{I})$, which contains only constant terms.
 - Then \mathbf{h} is computed as described.
- During BP, $\mu_{\mathbf{x}}$ and $\sigma_{\mathbf{x}}$ are then easily reachable from \mathbf{h} .
 - Derivatives to ζ (where sampling is performed) are not needed.
 - Optimization is only needed for part of the NN which construct $\mu_{\mathbf{x}}$ and $\sigma_{\mathbf{x}}$.

1.9.3.2 VAE Loss

- def **VAE Loss:** $C\|\mathbf{x} - g(\mathbf{h})\|^2 + KL(N(\mu_{\mathbf{x}}, \sigma_{\mathbf{x}}) \| N(\mathbf{0}, \mathbf{I}))$.
- It can be shown that $KL(N(\mu_{\mathbf{x}}, \sigma_{\mathbf{x}}) \| N(\mathbf{0}, \mathbf{I})) = \frac{1}{2}[-1 - \log(\sigma_{\mathbf{x}}^2) + \sigma_{\mathbf{x}}^2 + \mu_{\mathbf{x}}^2]$.
 - Where $\mu_{\mathbf{x}} = f_1(\mathbf{x})$ and $\sigma_{\mathbf{x}} = f_2(\mathbf{x})$.

1.10 Generative Adversarial Networks

1.10.1 Generative models

- A systems that:
 - Take a training set of samples drawn from a distribution p_{data} .
 - Learn to represent and estimate of that distribution.
- To learn $p(x)$ is way more difficult than to learn $p(y | x)$ (discriminative models).
- The distribution p_{model} can be estimated explicitly.
 - Or the model can only give the possibility to draw samples from it.
 - GANs are usually used to draw examples even if they can be designed to do both.
- Advantages of studying generative models:
 - The ability to **represent high-dimensionality** is important in many domains.
 - GM can be incorporated in (model based) **reinforcement learning**.
 - The GM model can be *queried* by the RL system to validate assumptions.
 - They can be used as the **basis of Semi Supervised Learning** systems, etc.

1.10.1.1 Maximum likelihood models

- GANs are compared to **maximum likelihood models** (a type of generative models).
- def **Maximum Likelihood Optimization**: $\theta^* = \arg \max_{\theta} p_{model}(\{x^{(i)}\}_{i=1}^m; \theta)$.
 - Parameters θ that maximize the likelihood of the training data given the model.
 - Almost always that optimization is made in **log-space**.
 - $\theta^* = \arg \max_{\theta} p_{model}(\{x^{(i)}\}_{i=1}^m; \theta) \approx \arg \max_{\theta} \mathbb{E}_{x \sim p_{data}} [\log p_{model}(x; \theta)]$.
 - **Maximizing** the ML w.r.t. θ is the same as **minimizing** the KL divergence of p_{data} and p_{model} .
 - $KL(p_{data} \| p_{model}) = \dots = \mathbb{E}_{x \sim p_{data}} [\log p_{data}(x)] - \mathbb{E}_{x \sim p_{data}} [\log p_{model}(x; \theta)]$.
 - $\mathbb{E}_{x \sim p_{data}} [\log p_{data}(x)]$ is the log-likelihood of data.
 - $\mathbb{E}_{x \sim p_{data}} [\log p_{model}(x; \theta)]$ is the log-likelihood of the model.
 - Since the minimization is w.r.t. to θ , only the second element must be maximized.
- Maximum likelihood models can be subdivided into ones with *explicit density* and ones with *implicit density*.
 - Explicit density can be subdivided into tractable density and approximate.

Explicit Density models with tractable density

- The distribution is fully modelled.
- **Fully Visible Belief Nets**:
 - Strong assumption: the distribution is tractable.
 - Chain rule of probability to factor the probability of model x into a product of 1D probabilities.
 - $p_{model}(x) = \prod_{i=1}^n p_{model}(x_i | x_1, \dots, x_{i-1})$.
 - They are the basis of sophisticated generative models.
 - *Problem*: sample must be generated one entry at a time → generating a new sample is $O(n)$.
- **Nonlinear Independent Component Analysis**:
 - A **simple distribution** over z coupled with a **non-linear transformation** g that warps space in complicated ways can yield a **complicated distribution** over x .
 - If there is a vector of latent variables z and a continuous differentiable, invertible transformation g such that $g(z)$ yields a sample from the model in x space, then:
 - $p_x(x) = p_z(g^{-1}(x)) |\det(\frac{\partial g^{-1}(x)}{\partial x})| = p_z(g^{-1}(x)) |\det(J_x(g^{-1}(x)))|$.
 - The density p_x is tractable if the density of p_z is tractable and the Jacobian of g^{-1} is tractable.
 - *Problem*: these models impose constraints on the choice of g .
 - The invertibility restriction impose that the number dimensions of x equals to the ones of z .

Explicit Density models with approximate density

- The distribution is approximated.
- **Variational approximations**:
 - Use a deterministic approximation to overcome the problems of having to deal with an intractable distribution.
 - The main idea is to define and maximize a lower bound on the intractable distribution.
 - $\mathcal{L}(x; \theta) \leq \log p_{model}(x; \theta)$.
 - This is the kind of approach taken by **variational autoencoder**.
 - Very often the approximation is based on multivariate gaussians.

- VAE learns make several approximation to the true likelihood using Gaussian distribution.
- *Drawback:*
 - When the approximation is too crude even with a perfect optimization and infinite data the gap between \mathcal{L} and the true likelihood can make the results poor.
 - VAE often obtain very good likelihood, but produce lower quality samples (w.r.t. GANs).
 - If compared to FVBMs, VAE are hard to train (but GANs are even harder to train).
- **Markov Chain Approximations:**
 - Use some form of stochastic approximation.
 - Goal: to find a way to efficiently draw samples from $p_{model}(\mathbf{x})$ when the distribution is intractable.
 - **Markov Chain** method (*MC Monte Carlo*) draw examples by repeatedly sampling from simpler distributions.
 - According to a transition operator: $\mathbf{x}' \sim q(\mathbf{x}' | \mathbf{x})$.
 - Convergence:
 - By repeating updating \mathbf{x} according to the transition operator q , MC methods can sometimes guarantee that \mathbf{x} will eventually converge to sampling from $p_{model}(\mathbf{x})$.
 - *Challenge: slow convergence.*
 - *Learning phase:* drawing of examples via MCMC can be too costly.
 - *Inference phase:* inefficient if compared to GANs.

Implicit Density models

- Some models model the distribution only implicitly.
- **Generative stochastic models:**
 - Learns a Markov transition operator that allows to draw from the implicit model.
 - *Problem:*
 - By approximating examples using a MC, these models have the same problems as MCMC.

1.10.2 Generative Adversarial Networks

- GANs model the distribution only implicitly.
 - The distribution is implicit but can be used to generates sample.
 - Offer a single step sample generation method.
- Features:
 - GANs can generate samples in parallel.
 - The generator function has very few restrictions.
 - No costly Markov chain approximations.
 - No variational bound is needed.
 - Subjectively GANS are regarded to produce better samples.

1.10.2.1 GANs functionality

- The idea is to set up a game between two players.
 - **Generator (G):** create samples intended to come from the same distribution as the training data.
 - **Discriminator (D):** examines samples to determine whether they are real or false.
- Definitions:
 - G and D are two differentiable functions (i.e. two NNs).
 - G takes an input \mathbf{z} , is defined in terms of parameters θ_G and outputs a value \tilde{x} from same space of \mathbf{x} .
 - D takes an input \mathbf{x} , is defined in terms of parameters θ_D and outputs a value $y \in \{\text{False}, \text{Real}\}$.
- Assumption: the discriminator tries to predict the probability that the input is real (regression).
 - $D(\mathbf{x}) \approx 1$ if the discriminator believes that the $P(y = \text{Real} | \mathbf{x})$ is high.
- **Cost functions:** $L_G(\theta_G, \theta_D)$ and $L_D(\theta_G, \theta_D)$.
 - ! They are both defined with respect to both θ_G and θ_D .
- Each player want to minimize its cost function but can only do so by **acting on its own parameters**.
 - The optimization can be seen as a zero-sum game between two players.
 - Not always true, depending on the definition of L_G and L_D .
 - The solution of a zero-sum is a **Nash equilibrium**.
 - A point in the (θ_G, θ_D) space such that the point is local minimum of L_D w.r.t. θ_D and a local minimum of L_G w.r.t. θ_G .
 - Nowadays, constraints for having a zero-sum game are abandoned for easier training.

1.10.2.2 GANs training

- The training process consists of simultaneous SGD.
 - Sample two minibatches (one from real data, one from generated data).
 - Evaluate the two losses and update θ_D using the gradients from L_D and θ_G using the ones from L_G .
- Sometimes more steps are performed on the discriminator before going back to the generator.

Discriminator's Cost function

- def **Discriminator's Cost function**: $L_D(\theta_D, \theta_G) = -\frac{1}{2}\mathbb{E}_{x \sim p_{data}}[\log D(x)] - \frac{1}{2}\mathbb{E}_{z}[\log(1 - D(G(z)))]$.
- The usual cross-entropy used when minimizing a binary classifier with sigmoid output units.
- Initially, all variants of GANs used this exact cost for the discriminator.
 - While, costs for the generator changes from model to model.
- $\frac{1}{2}$ is used since usually 50% of examples are real and 50% examples are not.
 - But since it is an optimization problem, constants are not impactful.
- The distribution of z is usually a Gaussian distribution.

Generator's Cost function

- def **Generator's Cost function (zero-sum game)**: $L_G = -L_D$.
 - Each time the discriminator lower its loss, the generator is penalized by the same amount.
 - The players are just trying to minimize/maximize (**minmax**) the same function.
 - $\theta_{G^*}, \theta_{D^*} = \arg \min_{\theta_G} \max_{\theta_D} V(\theta_D, \theta_G)$, with $V(\theta_D, \theta_G) = -L_D(\theta_D, \theta_G)$.
 - A **saddle (or minmax) point** is searched in the loss surface.
 - The same loss surface is evaluated, but from two different perspectives.
 - At a relative minimum along one axial direction and at a relative maximum along the crossing axis.
 - Theoretical advantages of this setup:
 - This game corresponds to minimizing the **Jensen-Shannon divergence** between the data and the model distribution.
 - def **Jensen-Shannon divergence**: $JSD(P||Q) = \frac{1}{2}KL(P||M) + \frac{1}{2}KL(Q||M)$.
 - Where M is the median distribution $M = \frac{1}{2}(P + Q)$.
 - The JSD is a symmetrical type of KL .
 - The game **converge to equilibrium** if both players' policies can be updated in **function space**.
 - This is not possible in practice though.
 - When the **discriminator minimizes** the cross-entropy, the generator **maximizes** it.
 - The gradient of the loss **vanishes** when the discriminator rejects the generator examples with high confidence.
- def **Generator's Cost function (non-saturating game)**: $L_G = \frac{1}{2}\mathbb{E}_z \log D(G(z))$.
 - The generator minimize a cross-entropy term tailored to its view of the problem.
 - The generator maximizes the probability that the discriminator is mistaken.
 - It lacks the theoretical properties of the minmax loss, but it doesn't suffer from gradient vanishing.
 - The generator loss has been chosen **heuristically** (not theoretically sound, it *just works*).
 - The game is no longer a zero-sum one.

1.10.2.3 Why do GANs work

- GANs performances were attributed to the minimization of JSD instead of KL .
 - KL is not symmetric and MLE minimizes $KL(P_{data}||p_{model})$.
 - Minimizing the JSD is more akin to minimizing $KL(p_{model}||p_{data})$.
- More recent results suggest that JSD doesn't explain why GANs make sharper samples.
 - Using ML to optimize GANs doesn't show problems in selecting a small number of modes and generating sharp images.
 - By selecting a single mode (corresponding to an image) the result won't be the average of several modes.
 - Averaging several modes results in a *blurrier* image.
 - GANs often choose to generate from very few modes.
 - Fewer than the number allowed by the model capacity.
 - **Model collapse problem**: generating always the same image.
 - Reverse KL would select as many modes as allowed by the model.
- This suggests GANs choose to generate a small number of modes due to a **defect in the training procedure**.
 - Rather than due to the divergence they aim to minimize.

- The reason why this happens is still not clear.
 - Maybe it **makes different approximations** than other models.
 - Maybe it **optimizes a different family of functions**.

1.10.2.4 Deep Convolutional GANs

- Nowadays, most GANs are loosely based on DCGANs.
- Main insights for this architecture are:
 - **Batch normalization:**
 - Most layers for both D and G , batches for G and D are normalized separately.
 - The last layer of G and the first layer of D are not batch normalized.
 - **No pooling nor unpooling layers:**
 - When G needs to increase the spatial dimension of the representation it used *transposed convolution (deconvolution)* with a stride greater than 1.
 - **Adam instead of SGD with momentum.**

1.10.3 Wasserstein GANs

- GANs are very difficult to train.
 - The game between G and D doesn't converge easily.
- def **Wasserstein distance** (*earth-mover distance*): $W(p, q) = \inf_{\gamma \in \Pi(p, q)} \mathbb{E}_{(x, y) \sim \gamma} \|x - y\|$.
 - A distance between two probability distributions p and q .
 - $\Pi(P, Q)$: the set of all joint distribution $\gamma(x, y)$ whose marginals are respectively p and q .
 - $(x, y) \sim \gamma$: (x, y) distributes according to γ .
 - Interpretable as the **minimum amount of work** required to transform the probability mass p to q .
- W has better properties than other distances and divergences used for GANs.
 - It allows **learning to converge** in many situations when other measures fail.
 - But the infimum in the definition of W is **intractable**.

1.10.3.1 Lipschitz continuity

- A *solution* to the intractability of W .
- **WGANS**: the family of functions considered $\{f_{\theta_D}(x)\}$ is **assumed to be Lipschitz continuous**.
 - def **Lipschitz continuous function**: f if $\exists c: \|f(x) - f(y)\| \leq c\|x - y\|$ for all x and y in $\text{dom}(f)$.
 - The smaller the constant, the smoother the function.
 - The function cannot change too much in a small region of the domain.
- Under this assumption:
 - **Discriminator training**: maximize W by maximizing $\mathbb{E}_{x \sim p_{\text{data}}} [D_{\theta_D}(x)] - \mathbb{E}_{z} [D_{\theta_D}(G_{\theta_G}(z))]$.
 - $\mathbb{E}_{x \sim p_{\text{data}}} [D_{\theta_D}(x)]$: the average score for real examples.
 - $\mathbb{E}_z [D_{\theta_D}(G_{\theta_G}(z))]$: the average score for fake examples.
 - **Generator training**: minimize $L_G(\theta_G, \theta_D) = -\mathbb{E}_z [D_{\theta_D}(G_{\theta_G}(z))]$.
 - Those quantities are easy to compute (therefore W **intractability is avoided**).
- Everything can be trained by BP with a small trick to ensure that the learned function is Lipschitz continuous.
 - Every time θ_D parameters are updated, they are **clipped** in $[-c, c]$ (where c is a user-defined constant).
 - eg With $[-5, 5]$, both 14 and 16 are clipped to 5, while -4.2 is left as it is (no normalization).
- Another difference is the activation function of the last layer of D is **linear** (instead of sigmoid).
 - D is no longer meant to model a probability distribution (there D is often referred as the **critic**).

1.10.3.2 WGANS advantages

- The loss is **meaningful** and **correlates with G convergence** and sample quality.
 - Other losses do not correlate with sample quality.
 - In these cases, with a smaller loss the sample quality is still low or even worse than before.
- The discriminator f is usually trained near optimality.
 - Way more epochs are assigned to the discriminator.
 - When D is properly trained, then G is trained further (for a small amount of epochs).
 - It can be shown that the loss is an **estimate of W given a factor** (determined by the constant c).
 - Empirical evidence shows that this correlates with the **quality of generated samples**.
 - The discriminator **should** be trained till optimality.

- When the critic is trained to completion, it simply provides a loss to the generator that can be trained as any other NN.
- Empirical evidence shows that WGANs are **much more robust** than GANs when one varies the generator.

1.11 Pruning

[Notes from a.a. 2024/2025, Pruning is not part of the exam curriculum in a.a. 2025/2026]

1.11.1 Deep networks

- Features of deep networks:
 - High number of hidden networks.
 - Parameters number is a proxy of complexity, but it's not its only source.
- Methods to both boost perform and make ANN robust to noise are needed.
- Recent trends in ANNs:
 - **Handcrafted** (eg VGG, ResNet).
 - **Efficient hand-crafted** (eg Inception).
 - Variation of handcrafted models with variations to make them more efficient.
 - **Neural-architecture search**.
 - A technique for automating the design of ANN.
 - Learning the architecture is adding one layer of complexity.
 - Other optimization metrics (eg best accuracy while minimizing number of parameters).
 - **Hardware-aware Neural-architecture search**.
 - In the optimization phase, the hardware is considered.
 - Information from the HW after a deploy of the model is collected.
 - Information like memory footprint, latency, energy consumption, etc.
 - And then this information is incorporated in the model.
 - A new optimization phase is launched after the information has been integrated.

1.11.1.1 Deep models outside the datacenters

- Communication has a cost, to have computation *on the edge* is desirable.
- Mobiles phones, DL accelerators, FPGAs have little memory.
- Methods: quantization (avoid FP-unit), knowledge distillation, pruning.
 - eg 8-bit quantization bit is good enough for most image classification.
 - **Knowledge distillation**:
 - A large *teacher* model and a shallow *student* model.
 - The *teacher* has learnt input-output mapping.
 - This knowledge is tried to be distilled in the *student*.
 - *Issue*: how *shallow* the model should be is not known to store this mapping.

1.11.2 Pruning

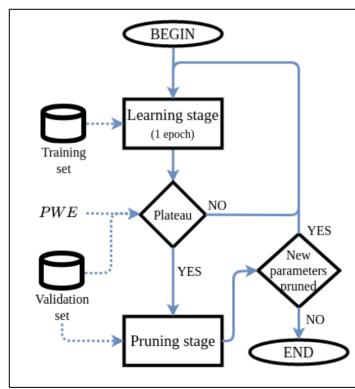


Figure 7: General Pruning scheme.

- **Sparsification**: to take some model parameters and set them to 0.
- **Pruning**: removing parameters or entire units from the DL model.
 - Pruning relates with sparsification.
 - The weight matrix (representing a layer) becomes sparse.
 - Removed units are set to 0, introducing nonetheless some representation overhead.
- Advantages:
 - The number of parameters is reduced (for special design like ASICs the gain is real).

- Modern GPUs implement functions to completely avoid operations involving zeros.
- If they are removed in a *structured* way (entire blocks), there is no representation overhead.
 - The gain is real even in general frameworks.
- Pruning is also biological plausible.
 - The peak synaptical connectivity is at 7 years of age, then the neuronal density decrease.

1.11.2.1 Traditional pruning methods

- Limits of traditional pruning methods:
 - Computationally extremely expensive.
 - They do not work so well in multi-layer architectures.
 - Slow (a lot of iterations to converge).

Skeletonization [Mozer & Smolensky, 1988]

- Motivation: computing was extremely expensive at the time.
- Principal idea:
 - Iteratively train the network to a certain performance criterion (until convergence).
 - Compute a measure of relevance: $\rho_i = E_{\text{without unit } i} - E_{\text{with unit } i}$.
 - A single unit i is removed and the error is computed without it.
 - If ρ_i is very high, the unit i is fundamental to mode performance.
 - Trim the least relevant units (once they are ranked).
- This scheme is the most common in several pruning approaches.

Optimal brain damage [Le Cun et al., 1989]

- Use of **second-order derivative information** to find a trade-off between **complexity** and **training error**.
 - Going therefore beyond gradient.
- The focus here switches from units to **single parameters**.
- MSE is used (CrossEntropy was not popular at the time).
- Principal idea: $\delta E = \sum_i g_i \delta u_i + \frac{1}{2} h_{ii} \delta u_i^2 + \frac{1}{2} \sum_{i \neq j} h_{ij} \delta u_i \delta u_j + O(\|\delta U\|^3)$.
 - With $g_i = \frac{\partial E}{\partial u_i}$ the gradient and $h_{ij} = \frac{\partial^2 E}{\partial u_i \partial u_j}$ the Hessian.
 - Let the error function being approximated through Taylor expansion.
 - A perturbation over the parameter vector (δU) will produce a perturbation on the error.
 - Goal: to find the largest subset of parameters whose pruning will cause the least increase of E .
 - **Intractable** in this form (eg for 2.6K parameters network, the Hessian would be 6.5×10^6).
- $\delta E = \frac{1}{2} \sum_i h_{ii} \delta u_i^2$ ($h_{ii} \delta u_i^2$ is the **neuron saliency**).
 - To approximate a result, the **diagonal** is used.
 - Deleting one parameter will not cause impact on the others (in general not true).
 - Assumption: the model has been already trained and a local minimum is found.
 - Therefore the gradient term of the error should be closed to zero (ignorable).
 - All the h_{ii} are non-negative (hence every perturbation will cause the error to go up or stay the same).
 - Estimating the diagonal of deviation nowadays can be done with the two back propagation passes.
 - It increases the complexity but acceptably.
- Algorithm:
 - Choose a reasonable network architecture.
 - Train the network until a reasonable solution is obtained (train until convergence).
 - Compute the second derivatives h_{ii} for each parameters.
 - Compute the saliences for each parameter.
 - Sort the parameters by saliency and delete some low-saliency parameters (*thresholding*).
 - Iterate to training step.
 - First work to propose **retraining**.
 - The number of parameters to be pruned is not known in advance.
- With retraining the MSE remains small even with fewer parameters (bigger pruning).
 - Without retraining, the MSE grows with fewer parameters.

1.11.2.2 Newer pruning methods

- **Newer pruning methods:**
 - A revival of pruning methods developed after 2015.
 - The number of parameters of modern architectures became huge.

- A slow process of parameters pruning is still used as traditionally was.

Learning both weights and connections [Han et al., 2015]

- **Magnitude pruning** (previously proposed).
 - Simpler idea than from the past.
 - If a parameter has a very small value (its **magnitude**), it is pruned from the network.
 - No concept of saliency or other estimator.
 - But the model is **regularized**.
- Principal idea:
 - Parameters are randomly initialized (nowadays, pre-trained models are used).
 - **Training stage**: parameters are updated then trained with standard GD until performance is achieved.
 - **Parameter sparsification**: parameters below T (hyperparameter) are removed, pruning connections.
 - Drop-out stochastically disable entire units, but they don't disappear from the net.
 - The underline parameter is still in the model.
 - In parameter sparsification, the parameter is **permanently removed** after training.
 - **Neuron sparsification**: neurons without input arcs input are pruned from the network.
 - This can lead to a degradation in network performance.
 - **Fine-tuning**: fine-tune the model, recovering the performance and iteratively prune again.
 - Called also re-training (by LeCun) or regularization.
- Assumption: parameters having low magnitude are also **less important**.
 - A regularization function that is trying to push as much as possible value of the parameters of the NN as close as possible to 0 is desirable.
- **Regularization**:
 - From optimization theory, LASSO (L1) is the most common choice.
 - L0 is non-differentiable (there are proxies).
 - Empirically, **L2 regularization** leads to **sparser models** (under same performance constraints).
 - Fine-tuning is the key.
 - **L2 regularization with iterative prune and retrain** works better than other variants.
 - Parameters used to store *noise* are pruned away.
 - Pruning help in the **signal-to-noise** ratio.
 - It maximizes the storage capacity of the network.
 - L2 is very cheap to implement inside an optimizer.
- Unstructured pruning can also unveil some structure of input data (even in deeper architectures).
- Network pruning can save $9\times$ to $13\times$ parameters with no drop in predictive performance.
- In order to prune the network, **various re-training** are needed (**extremely expensive**).
 - **Goal**: reach the highest sparsity with no task-related performance degradation.
 - The most effective approaches are also the ones **more computationally intensive**.
 - One-shot (or few-shot) pruning approaches are in general worse.
 - One-shot pruning: trained → prune → fine-tune and then stop.
 - Not as effective as re-training approach.
 - New challenge: achieve sparsity with less computation at training time.

Lottery ticket hypothesis [Frankle & Carbin, 2019]

A randomly-initialized, dense neural network contains a **sub-network** that is initialized such that, **when trained in isolation**, it can **match the test accuracy** of the original network after training for at most the same number of iterations.

- **Pruning at initialization** (even before training).
- The sub-network exists already at initialization.
 - If it can be found, a lot of computation can be saved.
 - Computation for training the model (less parameters to train).
 - Computation for pruning (no iterative pruning anymore, just pruning at initialization, **zero-shot pruning**).
- **Identification of lottery winners**:
 - A dense architecture is taken and training and following pruning is performed.
 - **Rewind process**: just the parameters which have not been pruned will be rolled-back to the value they had at initialization (before training).
 - A smaller model is obtained, onto which the training is performed.
 - The **same performance** is obtained training the model with **less parameters**.

- But the transition from **dense to small model** shouldn't need **full training information**.
- So there are parameters *winning at the lottery of initialization* in deep models.
 - It is possible to successfully train a model from the initialization phase.
 - Removing the largest part of the parameters and **training just the remaining fraction**.
- In This work, just the **existence of lottery winners** has been shown.
 - How to identify them at initialization (or in the first learning stages) was not proposed.
- Several methods for pruning models at initialization were then proposed.

Rigging the lottery [Evcı et al., 2019]

- Principal idea:
 - Start from a random, sparse configuration of the model.
 - Then, some connections can be either chopped (pruning) or re-instated (**growing**).
- **Pruning during training.**
 - Instead of waiting to reach convergence.
 - Pruning is performed during a certain amount of iterations (eg after every epoch).
 - The network therefore became sparser and sparser.
- But maybe a parameter has low magnitude at training just because convergence is not reached.
 - No certainty that only *useless* parameters has been removed.
- Instead of only removing some parameters (**drop**), some parameters are **reintroduced** (**growing**).
 - **Growing:** connections are created in a sparse model.
 - *Proposal:* grow connection based on the **gradient** of the units.
 - A high gradient for a missing parameter is a signal to reintroduce the parameter.
 - An iterative scheme that is locally reducing the parameters number and that might converge.
- **Massive disoptimality:**
 - Tremendously sensitive to hyper-parameters choice (eg pruning or growing rate at each iteration, etc).
 - Tremendously sensitive to dataset, architecture, etc.
- First try at solving the lottery ticket problem, but not an effective one.

SNIP [Lee et al., 2019]

- *SNIP: Single-shot network pruning based on connection sensitivity.*
- **Prune just-once**, instead of running an iterative algorithm involving regularization+pruning.
- A randomly initialized network (not even trained) is considered:
 - The **gradient** is computed for each parameters (like *gradient accumulation*).
 - *Assumption:* if the gradient is low at initialization, the parameter will not evolve much during training.
 - These parameters can be removed from the network.
 - The ranking of parameters to remove is not magnitude-based, but **sensitivity-based**.
 - This is performed right after initialization, hence the main direction of the gradient is preserved.
- Similar to *Optimal Brain Damage* (but with gradient instead of second derivatives).
 - In OBD: the network is converged \Rightarrow the gradient is $\approx 0 \rightarrow$ second order information is needed.
 - In SNIP: since not at convergence, gradient information is enough.
 - Loss variation is assumed constant during training (*bold assumption*).
- Pruning at initialization using a **gradient informing the strategies** could have been a good choice to reduce the complexity.
 - This is true for MNIST, but for bigger datasets the reality is different.
- Iterative approach are generally better than zero or few-shot approaches.

Lottery winners identification

- To identify winners at initialization time is not feasible [Frankle et al, 2020 2021]:
 - Few **iterations of warm-up** is needed.
 - eg 2K for ResNet-20 on CIFAR-10, 8 out of 90 epochs for ImageNet).
 - Before there are huge (but decreasing) gradients and rapid motions in weight space.
 - Without warm-up, all lottery winners found are *unstable*.
 - This *resizes* the original lottery winner hypothesis.
- NN are initialized to maximize FP and BP signal, to find the local minimum as fast as possible.
 - Adam optimizer is typically used to converge quickly.
 - SDG is typically used for robustness.

1.11.2.3 Structured and unstructured sparsity

- Unstructured sparsification focuses on the connection.
 - Pruning connections.
 - Each neuron is still part of the network (possible bottleneck).
- Structured sparsification focuses on the neurons.
 - Pruning neurons.
 - Pruned neurons can be ignored.
- Testing on embedded devices.
 - Compression with NN (eg MPEG-7):
 - Pruning → simplification → entropy coding → bit stream → decompression.
- Unstructured strategies prune way more parameters than the structured counterpart.
 - eg Even with 99% reduction in parameters, the memory footprint can goes from 47MB to 39MB.
 - Even with entropy coding, the memory footprint saving is not much.
 - While structured approach, the parameter reduction is less but the memory footprint saving is bigger.
- There are hybrid approaches that uses both structured and unstructured sparsity.

1.11.2.4 Pruning and GPUs

- With parallel computation, the bottleneck is in either caching or the critical path (maximum model's path).
 - If caching is sufficiently fast, the bottleneck is the critical path.
- The possibility of pruning entire layers must be established.

Layer collapse

- When a specific layer in a NN fails to effectively differentiate the input features.
 - Causing all outputs from the layer to converge to similar or identical values regardless of input.
 - Seen as a downside of poor initialization or hyper-parameter fine-tuning (i.e. the model does not learn).
 - But if the model learns while some layers are collapsed, the computation of this layers can be **skipped**.
- Layer collapse **without skip connections**:
 - One way to reduce model's depth is to **remove non-linearities**.
 - Layer fold is one approach parametrizing the negative slope of a PReLU.
 - α is a parameter: when $\alpha = 0 \rightarrow$ ReLU, when $\alpha = 1 \rightarrow$ identity activation.
 - With an identity activation, the layer is successfully **linearized**.
- With fully-connected layers:
 - The solution is straight-forward using simple linear algebra.
- With convolutional layers (stride = 1, no padding):
 - The solution is straight-forward using simple linear algebra.
- With convolutional layers (stride = 1, with padding = 1):
 - Up to 2024, no closed-form solutions has been found.
 - Most architectures using skip connections make use of padding to maintain same dimensionality of the output.

1.12 Transformers

1.12.1 Introduction to transformers

- The main idea behind LLM is:
 - To have a **large model**.
 - Trained on a **large dataset**.
 - To **predict the next token in a sequence**.
- They unlocked **efficient parallel processing** and the ability to **model long-range dependencies**.
- Before Transformers, the main tools were RNNs and LSTMs.
 - These models have some limitations.
 - They are **sequential** and cannot be easily parallelized.
 - They have difficulties in capturing **long-range dependencies**.

1.12.1.1 Foundation models

Scaling hypothesis: performance improves smoothly as we increase the **model size**, **dataset size**, and amount of **compute** used for training. For optimal performance **all three factors must be scaled up in tandem**.

- Large models that are trained on a wide range of tasks.
 - And can be **fine-tuned** for specific tasks.
- Transformers make them possible because of the following reasons:
 - They are **scalable** and can be trained on very large datasets.
 - Exploiting **parallelism** and distributed computing.
 - They can be trained in a **self-supervised** way (no need for labeled data).
 - The **scaling hypothesis** asserts that:
 - Simply by increasing the scale of the model (as measured by the number of parameters)
 - And training on a commensurately large dataset.
 - Significant improvements can be achieved, even with no architectural changes.
 - Performance has a power-law relationship with each of the three scale factors.

1.12.2 Attention

- A mechanism that allows a model to focus on different parts of the input when making predictions.
- Originally introduced as an enhancement to RNNs for machine translation.
 - Later showed that significantly better results can be achieved using attention mechanism alone.
 - **Eliminating the recurrence mechanism** completely.
- A transformer can be viewed as a way to build a **richer form of embedding**.
 - In which a given vector is mapped to a location that depends on the other vectors in the sequence.
 - While previous NN architectures (once trained) are fixed on word orders.
- Why attention work better than recurrence:
 - Attention layers let **every token interact with every other token in one step**.
 - Whereas recurrent models process tokens **one at a time**.
 - This remove the sequential bottleneck, enabling **massive parallelization**.
 - And helping capture **long-range dependencies** more effectively.

1.12.2.1 Transformer processing

- Input data to a transformer is a sequence of vector $[x_n^T]_{n \in [N]}$ of dimensionality D .
 - Each vector is called a **token** (eg a word in a sentence, a patch within an image, etc).
 - Tokens are collected in a matrix $\mathbf{X} \in \mathbb{R}^{N \times D}$, which is the input to the transformer.
- The fundamental building block of a transformer is the **transformer layer**.
 - A function that takes $\mathbf{X} \in \mathbb{R}^{N \times D}$ as input and produces a matrix $\tilde{\mathbf{X}} \in \mathbb{R}^{N \times D}$ as output.
 - $\tilde{\mathbf{X}} = \text{TransformerLayer}[\mathbf{X}]$.
 - The transformer layer is composed by two blocks: the **attention** block and the **transform** block.

1.12.2.2 Attention weights

- Assume to want to compute new embeddings $\mathbf{y}_1, \dots, \mathbf{y}_N$ for tokens x_1, \dots, x_N .
 - In such a way that the embedding for y_n depends on the embeddings of all other token.
 - Instead of moving information *left-to-right* (as in RNNs), everything is connected.
 - $\mathbf{y}_n = \sum_{m=1}^N a_{nm} \mathbf{x}_m$ (where a_{nm} are the **attention weights**).

- Requirements for attention weights:
 - Capture the similarity between the tokens \mathbf{x}_n and \mathbf{x}_m .
 - $a_{nm} \geq 0$.
 - $\sum_{m=1}^N a_{nm} = 1$ (to form a distribution).
- def **Dot-product Self-attention**: $a_{nm} = \frac{\exp(\mathbf{x}_n^T \mathbf{x}_m)}{\sum_{m'=1}^N \exp(\mathbf{x}_n^T \mathbf{x}_{m'})}$.
 - **Similarity** is computed via **dot-product** (common method).
 - It computes the a between the tokens in the same sequence using dot-product.
 - It uses **softmax** to form a distribution.
 - It takes the exponential and then normalize the result to sum to one.
 - With unrelated \mathbf{x}_n and \mathbf{x}_m (i.e. orthogonal), the attention weights will be ≈ 0 .
- def **Dot-product Self-attention** (matrix version): $\mathbf{Y} = \text{Softmax}[\mathbf{X}\mathbf{X}^T]\mathbf{X}$.
 - Fast way to compute all the attention weights at once.

1.12.2.3 Self-attention

- **Query, key and value**:
 - **Query**: the *user* request to get the value.
 - **Key**: the information that the system uses to accompany the value and that should be match with query.
 - **Value**: the information should be returned when the query match the key.
- **Soft-attention**:
 - Continuous variables are used to measure the degree of match between the query and the keys.
 - These variables are used to weight the influence of the value vectors on the output.
- Question: given an embedding of a token \mathbf{x}_n , to which information should this token attend to compute its new embedding and how much.
 - \mathbf{x}_n can be used to produce a **query**, and each \mathbf{x}_m in the sequence can be seen as giving rise to:
 - A **key** than can be matched against the **query** \mathbf{x}_n , to get a sense of how similar they are.
 - A **value** that can be used to compute how much influence token \mathbf{x}_m should have on the new embedding of \mathbf{x}_n .
- def **Self-attention** (matrix version): $\mathbf{Y} = \text{Softmax}[\mathbf{X}\mathbf{X}^T]\mathbf{X} = \text{Softmax}[\mathbf{Q}\mathbf{K}^T]\mathbf{V}$.
 - **Value**: the token \mathbf{x}_n that will be used to compute output tokens.
 - **Key**: the token \mathbf{x}_n for value \mathbf{x}_n .
 - **Query**: the token \mathbf{x}_n that will be used to compute the attention weights for output \mathbf{y}_n .
 - Self-attention since the **same sequence** is used to determine **all three components**.
 - **Q, K** and **V** are actually three different linear projections of **X**.
 - These projections allow the model to learn **separate similarity spaces** for queries, keys and values.

Trainable parameters

- The given formula is **deterministic** and doesn't depend on the parameters model (not trainable).
- Each feature with token \mathbf{x}_n contributes equally to the attention weights.
 - Whereas the flexibility to focus more on some features than on others is desirable.
- **Trainable parameters** are introduced to compute the attention weights, $\tilde{\mathbf{X}} = \mathbf{X}\mathbf{U}$.
 - Where $\mathbf{U} \in \mathbb{R}^{D \times D}$ is a matrix of trainable parameters.
 - Analogous to a layer in a standard NN.
 - Therefore the formula for new embeddings is $\mathbf{Y} = \text{Softmax}[\mathbf{X}\mathbf{U}\mathbf{U}^T\mathbf{X}^T]\mathbf{X}\mathbf{U}$.
 - But the matrix $\mathbf{X}\mathbf{U}\mathbf{U}^T\mathbf{X}^T$ is **symmetric** (while significant asymmetries should be supported).
 - eg *Chisel* should be strongly associated with *tool* (since every chisel is a tool).
 - But *tool* should only be weakly associated with *chisel* (not every tool is a chisel).
 - Also, the **same matrix U** is used to define both value vectors and attention coefficients (not ideal).
 - Therefore different separate matrices are defined for queries, keys and value.
- **Trainable parameters**:
 - $\mathbf{Q} = \mathbf{X}\mathbf{W}^{(q)}$.
 - With dimensionality $D \times D_k$ (where D_k is the dimensionality of the **key** vectors).
 - $\mathbf{K} = \mathbf{X}\mathbf{W}^{(k)}$.
 - With dimensionality $D \times D_k$, so that the dot product with **query** vectors $\mathbf{Q}\mathbf{K}^T$ is well defined.
 - $\mathbf{V} = \mathbf{X}\mathbf{W}^{(v)}$.
 - With dimensionality $D \times D_v$ (where D_v is the dimensionality of the output **value** vectors).

- def **Self-attention**: $\mathbf{Y} = \text{Softmax}[\mathbf{Q}\mathbf{K}^T]\mathbf{V}$.

Scaled self-attention

- The gradient of the softmax function can become exponentially small for inputs of high magnitude.
 - Which can lead to **vanishing gradients** during training.
- The values of the dot product can be **scaled** by a factor \sqrt{D} .
 - Where D is the dimensionality of the key vectors.
- def **Scaled self-attention**: $\mathbf{Y} = \text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Softmax}\left[\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{D_k}}\right]\mathbf{V}$.
- **Choice of the scaling factor $\sqrt{D_k}$** :
 - \mathbf{q} and \mathbf{k} vectors are assumed to be independent and have zero mean and unit variance.
 - D_k is then exactly the **standard deviation** of each one of the dot products in $\mathbf{Q}\mathbf{K}^T$.
 - The standard deviation of $\mathbf{q}^T\mathbf{k}$ matters since:
 - **Softmax saturates for large-magnitude inputs**.
 - Very large or very small logits entering the softmax leads to **vanishing gradients**.
 - The **typical magnitude** of $\mathbf{q}^T\mathbf{k}$ grows as $O(\sqrt{D_k})$.
 - Since $\text{std}(\mathbf{q}^T\mathbf{k}) = O(\sqrt{D_k})$, increasing the dimensionality of key/query vectors makes the dot-product logits naturally grow in magnitude.
 - Larger logits make the softmax more likely to saturate, causing the above gradient issue.

1.12.2.4 Multi-head self-attention

- Used to handle when there might be **multiple patterns of attention** that are relevant at the same time.
 - eg IN NL, some patterns might be relevant to tense whereas other might be relevant to subject-verb agreement.
- Given H heads indexed by $h \in [H]$, for each h :
 - $\mathbf{H}_h = \text{Attention}(\mathbf{Q}_h, \mathbf{K}_h, \mathbf{V}_h)$.
 - Where $\mathbf{Q}_h = \mathbf{X}\mathbf{W}_h^{(q)}$, $\mathbf{k}_h = \mathbf{X}\mathbf{W}_h^{(k)}$ and $\mathbf{V}_h = \mathbf{X}\mathbf{W}_h^{(v)}$.
 - The heads are first concatenated and then multiplied by a matrix $\mathbf{W}^{(o)}$.
 - def **Multi-head self-attention**: $\mathbf{Y}(\mathbf{X}) = \text{Concat}[\mathbf{H}_1, \dots, \mathbf{H}_H]\mathbf{W}^{(o)}$.
 - Dimensionality:
 - $\mathbf{Y}(\mathbf{X})$: $N \times D$.
 - $\text{Concat}[\mathbf{H}_1, \dots, \mathbf{H}_H]$: $N \times HD_v$.
 - $\mathbf{W}^{(o)}$: $HD_v \times D$.
 - $\mathbf{W}^{(o)}$ is both used:
 - To **reduce dimensionality** (to have $N \times D$).
 - To **learn to best combine attention heads to produce a useful output**.
 - Typically $D_v = \frac{D}{H}$ so that the concatenated matrix has the same dimensionality as the input matrix.
 - So that $N \times D$ for the concat, and $D \times D$ for $\mathbf{W}^{(o)}$.
 - In this case $\mathbf{W}^{(o)}$ is not needed for dimensionality reduction, but for the other reason.

1.12.3 Transform block

- A few additional improvements can be made to self-attention to make it more expressive and easier to train.
- To improve training efficiency:
 - A **residual connection** around self-attention is added.
 - Followed or proceeded by **layer normalization**.
 - $\mathbf{Z} = \text{LayerNorm}[\mathbf{Y}(\mathbf{X}) + \mathbf{X}]$ or $\mathbf{Z} = \mathbf{Y}(\text{LayerNorm}[\mathbf{X}]) + \mathbf{X}$.
- The output is then passed through a non-linear NN with D input units and D output units (MLP).
 - This can be a two-layer feedforward NN with ReLU activations.
 - $\tilde{\mathbf{X}} = \text{LayerNorm}[\text{MLP}[\mathbf{Z}] + \mathbf{Z}]$ or $\tilde{\mathbf{X}} = \text{MLP}[\text{LayerNorm}[\mathbf{Z}]] + \mathbf{Z}$.

1.12.4 Positional encoding

- **Permutation equivariant of self-attention**:
 - Consider $\mathbf{X} \in \mathbb{R}^{N \times d}$ a sequence of token embeddings and \mathbf{P} a permutation matrix of the rows of \mathbf{X} .
 - \mathbf{P} is a square binary matrix obtained by permuting the rows of an identity matrix.
 - $\mathbf{P}\mathbf{X}$ (left multiplication) \rightarrow permuting the rows i and j of \mathbf{X} iff $P_{ij} = 1$.

- $\mathbf{P}^T = \mathbf{P}^{-1} \implies \mathbf{P}^T \mathbf{P} = \mathbf{P} \mathbf{P}^T = \mathbf{I}$.
- The main matrix operation in the transformer are $\mathbf{X} \mathbf{X}^T \mathbf{X}$.
 - Note that $\mathbf{P} \mathbf{X} (\mathbf{P} \mathbf{X})^T \mathbf{P}^T \mathbf{P} = \mathbf{P} \mathbf{X} \mathbf{X}^T \mathbf{P}^T \mathbf{P} = \mathbf{P} (\mathbf{X} \mathbf{X}^T) \mathbf{P}$.
 - \implies **Self-attention is equivariant** to the permutations of the input tokens.
- This is a problem, since **order** of tokens is usually important (eg in text sentences).
- The goal is not to change the transformer architecture.
 - Instead, to **encode the position of the tokens in the input sequence**.
 - Idea: to encode the position n of the n -th tokens as an additional input vector \mathbf{r}_n .
 - And combine it with the token vector \mathbf{x}_n before passing it to the transformer.

1.12.4.1 Appending or adding positional information

- **Appending** the positional encoding to the token vector is not desirable.
 - Since the dimensionality of the input space and subsequent layers increase.
- The positional encoding can be simply **added** to the token vector.
 - This can seems dangerous since it induces a corruption or loss of information.
 - But two randomly chosen uncorrelated vectors tend to be nearly **orthogonal in high-dimensional spaces**.
 - Allowing the network to process them separately.
 - The residual connection allows the positional information to not get lost too.
- Due to linear processing in the transformer, a **concatenated representation exhibits properties to an additive one**.

1.12.4.2 Ideal and bad positional encodings

- An ideal encoding should:
 - Be **unique** for each position.
 - Be **bounded** (each element of the encoding representation should have a **finite range**).
 - Generalize to **sequences of arbitrary length**.
 - Have a consistent way to express **relative positions**.
- eg Bad positional encodings:
 - **One-hot encoding:**
 - It's unique and bounded.
 - But it doesn't generalize to sequences of arbitrary length.
 - And not make it easy to reason about relative positions.
 - **Assigning an integer to each position:**
 - It's unique, but not bounded.
 - It may start to corrupt the vector significantly as the sequence length increases.
 - **Assigning a real number in $[0, 1]$ to each position:**
 - It's bounded, but not unique since it depends on the length of the sequence.

1.12.4.3 Sinusoidal positional encoding

- There are many approach to define positional encodings.
 - One of the most popular is the **sinusoidal positional encoding**.
- def **Sinusoidal positional encoding**: $\mathbf{r}_n = [\sin(w_1 \cdot n), \cos(w_1 \cdot n), \dots, \sin(w_{D/2} \cdot n), \cos(w_{D/2} \cdot n)]$.
 - Where n is the token position, $w_i = \frac{1}{10000^{2i/D}}$ (frequency) and D the size of the representation.
- prop This encoding makes it easy to reason about **relative positions**.
 - Two reasoning sustains this assertion:
 - The dot product between two positional encoding \mathbf{r}_n and \mathbf{r}_m depends only on $n - m$.
 - And not on the absolute positions n and m .
 - Encoding of $n + m$ can be expressed as a **linear combination of the encodings** of n and m .
 - It is always possible to find \mathbf{M} that depends only on k , such that $\mathbf{r}_{n+k} = \mathbf{M} \mathbf{r}_n$.

1.12.5 Applications of Transformers

- Three modes:
 - **Prediction (encoders)**
 - **Generation (decoders)**
 - **Translation or seq2seq (encoder-decoders)**.

1.12.5.1 Tokenization

- Tokens are generally small groups of characters.
 - They might include words in their entirety.
 - Created in a pre-processing step that convert a string of words and punctuation into a string of tokens.
- Tokenization can be applied to various types of data (images, etc).
- **Byte pair encoding:**
 - One of the most common tokenization method.
 - It initially treats the set of unique characters as 1-character-long n -grams (the initial tokens).
 - Then, successively, the **most frequent pair of adjacent tokens** is merged into a new, longer n -gram.
 - All instances of the pair are replaced by this new token.
 - This is repeated until a vocabulary of prescribed size is obtained.
- Rule of thumb: one token generally corresponds to ~ 4 character of text for common English.
 - This translates to roughly $\frac{3}{4}$ of a word (eg 100 tokens ≈ 75 words).

Detokenization

- Modern tokenizer must take *special care* to ensure that text can be reconstructed **exactly**.
 - With a naive concatenation, the original spacing may be lost.
 - A proper **detokenization** step is therefore required to recover the original text.
- Modern LLM tokenizers **encode** whether a **space precedes each token directly inside the token itself**.
 - During tokenization a special prefix character is added whenever a token begins a new word.
 - Detokenization simply consists of joining the tokens **exactly as the tokenizer produced them**.
 - After which the special space-prefix markers are converted back into normal spaces.

1.12.5.2 Decoder transformers (*generation*)

- Goal: to use a transformer architecture to construct a **autoregressive model**.
 - **def Autoregressive model:** $p(\mathbf{x}_1, \dots, \mathbf{x}_N) = \prod_{n=1}^N p(\mathbf{x}_n | \mathbf{x}_1, \dots, \mathbf{x}_{n-1})$.
 - But in the attention weights, everything is connected, is not sequential.
 - $p(\mathbf{x}_n | \mathbf{x}_1, \dots, \mathbf{x}_{n-1})$ are expressed by a transformer network learned from data.
- The architecture consists of a **stack of transformers**:
 - That a sequence of tokens.
 - And produce a sequence $\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_n$ of dimensionality D as output.
 - Then a linear transformation followed by a softmax to compute the **distribution** over the K output tokens.
 - $\mathbf{Y} = \text{Softmax}[\tilde{\mathbf{X}}\mathbf{W}^{(p)}]$.

Masking

- Decoder models generate text by sampling from $p(\mathbf{x}_n | \mathbf{x}_1, \dots, \mathbf{x}_{n-1})$ at each step n .
 - Older predictions are **used as input** to the model to generate the next token.
- With other NN architecture for processing sequences it would be nice to **process entire sequences at once**.
 - But the model would be able to see the future tokens (making the task trivial).
- The attention mechanism gives a way to **mask out the future tokens** when computing the attention weights.
 - Thus allowing parallel processing of whole sequence.
 - The idea consists of two steps:
 - Add a special token to the input sequence that represents the start of the sequence.
 - Mask out (set to zero) the attention weights for the future tokens.
 - The attention weights are computed as $\mathbf{Q}\mathbf{K}^T$ but with attention set to 0 for **future tokens**.
 - A **mask matrix** \mathbf{M} has $-\infty$ in the upper triangular part.
 - $\mathbf{Y} = \text{Softmax}[\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{D_k}} \circ \mathbf{M}\mathbf{V}]$.
 - Therefore \mathbf{y}_i will depends only on \mathbf{x}_1 to \mathbf{x}_{i-1} .

Padding

- To allow processing multiple sequences at once, it is desirable to collect them in a sequence.
 - But this would require that all sequences have the same length.
- To solve this issue, sequences are **padded** to the same length.
 - Using a special token to represent the padding.

- A **mask matrix M** that has $-\infty$ in the positions of the **padding tokens** is used.
 - So that the attention weights for the padding tokens are zero.

Sampling strategies

- Once the distribution $p(\mathbf{x}_n | \mathbf{x}_1, \dots, \mathbf{x}_{n-1})$ is computed, the next token can be sampled from it.
- The **greedy strategy** simply chooses the token with the highest probability.
- To find the most probable sequence, the joint distribution over all tokens must be maximized.
 - $p(\mathbf{y}_1, \dots, \mathbf{y}_N) = \prod_{n=1}^N p(\mathbf{y}_n | \mathbf{y}_1, \dots, \mathbf{y}_{n-1})$
 - The higher probability sequences can be generated using a **beam search** strategy.
 - Which keeps track of the k most probable sequences at each step.
- However, the most likely sequences are not necessarily the most human-like sequences.
 - By sampling on the full distribution, sequences that are nonsensical or grammatically incorrect can be generated.
 - This arises from the typically very large size of the token dictionary.
 - Which has **long tail of tokens** with **very low probability**.
- **Top-K sampling** mitigates this issue by sampling from a **truncated distribution**.
 - A distribution that includes the top-K most probable tokens.
- **Top- p sampling (nucleus sampling)** samples from a truncated distribution that:
 - Includes the smallest set of tokens whose cumulative probability exceeds a threshold p .
- def **Temperature scaling**: $y_i = \frac{\exp(a_i/T)}{\sum_j \exp(a_j/T)}$.
 - A soft version of Top-K sampling.
 - Scales the pre-activation values by a temperature parameter T before applying the softmax function.
 - When:
 - $T \rightarrow 0$ the distribution becomes more peaked around the most probable tokens.
 - $T = 1$ the distribution is the same as the original softmax distribution.
 - $T \rightarrow \infty$ the distribution becomes uniform across all states.

1.12.5.3 Encoder transformers (*prediction*)

- Used to process sequences of tokens and produce contextual embeddings that can be used to produce a **fixed-size representation of the sequence**.
- **Masked language model**:
 - The `<class>` is put at the beginning of the input sequence.
 - During pretraining, **class token is not itself predicted**, since only masked tokens contribute directly to the loss.
 - But it still participates fully in the transformer layers.
 - It attends to all other tokens (and vice versa), so its representation becomes a **contextual summary of the whole sequence**.
 - It is never masked during pretraining.
 - A randomly chosen subset (eg 15%) of the tokens are replaced by the **mask token**.
 - The model is then trained to predict the original tokens from the masked tokens.
 - **Self-supervised learning task**.
- For sequence classification task, the **class token** is used to produce a fixed-size presentation.
 - Which is then passed through a linear layer to produce the output.
- For token classification tasks, the output of the transformer is passed through a linear layer (or more complex classifier) for each token to produce the output.
 - The output of the class token is ignored.

1.12.5.4 Encoder-Decoder transformers (*seq2seq*)

- To perform translation, the transformer architecture uses a **cross-attention mechanism**.
 - It allows the **decoder to mix information from the encoder** with the information generated so far.
 - The output of the E Z is fed with entire input sequence in the **multi-head cross attention** in D .
 - Z will be the representation of the **class token**.
- At training time, the input of D will be the *translated sequence*.
 - It is therefore a **supervised learning task**.

1.13 Graph neural networks

1.13.1 Graph neural networks

- NNs tailored to process **graphs** (their architecture are not *graph-like*).
- Graphs are a general language for **describing entities with relations/interactions**.
 - eg Social networks, biological networks, chemical compounds, etc.

1.13.1.1 Graphs

- **def Graph:** $G = (V, E)$, a tuple of a set of **nodes** and a set of **edges**.
 - Each edge is a pair of two vertices.
- **Adjacency matrix:**
 - An undirected graph G with n nodes is assumed.
 - A (binary) **adjacency matrix** A^G :
 - A square matrix of size $n \times n$.
 - Each entry A_{ij}^G is 1 if there is an edge between the node i and j .
- **Setup:**
 - A graph $G = (V, E)$ is given.
 - V is the set of nodes and E is the set of edges.
 - A is the adjacency matrix.
 - $N(v)$ is the set of neighbors of node $v \in V$.
 - Each node v has an associated **feature vector** $\mathbf{x}_v \in \mathbb{R}^d$:
 - If not, $\mathbf{x}_v = 1$ can be set.
 - Or an indicator vector to each node (one-hot encoding) can be assigned.

Tasks involving graphs

- **Node-level prediction** (eg drug-drug interaction).
 - The training can be performed on a single graph.
 - The test is performed on *unlabelled* nodes.
- **Link-level prediction** (eg recommendation).
 - The training can be performed on a single graph.
 - The model can learn patterns and relationships that indicate preferences.
 - The model can use interactions and node features to predict the likelihood of preferences.
- **Graph-level prediction** (eg time of arrival).

1.13.2 Graphs and Machine learning

- The traditional method to handle graphs using ML:
 - Extract features from the graph (eg node degree, centrality, subgraph patterns, etc).
 - Use these features as input to a ML model (eg SVM, NN, RF, etc).
- **Graph representation learning:**
 - DNN have changed the way data is learnt and represented.
 - Feature engineering is replaced by **learning representation**:
 - Learn a function that maps a graph to a vector representation.
 - Use the learned representation for downstream tasks (eg node classification, etc).

1.13.2.1 Node embeddings

- Early approach in GNN.
- Goal: define $ENC()$ so that $sim(u, v) \approx \mathbf{z}_u^T \mathbf{z}_v$.
 - Where sim is a similarity measure between nodes u and v .
 - A simple way to get graph embeddings is to aggregate (eg average) the node embeddings.

Encoder-decoder framework

- **Encoder:** $\mathbf{z}_v = ENC(v)$.
 - A function that maps nodes to **vector representations**.
- **Similarity:** $sim(v, u)$.
 - A function that measures the similarity between two nodes in the **graph space (domain knowledge)**.
 - **Random walk:** a sequence of steps where at each step, a node is chosen randomly from the neighboring nodes of the current node.

- **Random Walk similarity (RW)**: a node similarity based on visiting node u on a random walk starting from node v .
- **RW** is a flexible stochastic definition that incorporates both **local and high-order neighborhood information**.
- eg Expected commute time between two nodes, etc.
- **Decoder** $y_{vu} = DEC(\mathbf{z}_v, \mathbf{z}_u)$.
 - A function that compute the similarity between two nodes in the **embedding space**.
 - Usually, a function of the dot product between embeddings of two nodes.
 - $y_{vu} = DEC(\mathbf{z}_v, \mathbf{u}) = f(\mathbf{z}_v^T \mathbf{z}_u)$.
- Once fixed DEC , the goal is to learn the function ENC such that the similarity in the embedding space y_{vu} **approximates the similarity in the graph space**.

Learning

- *Goal:* to learn the parametrized function ENC_θ such that $\mathbf{z}_v^T \mathbf{z}_y = ENC_\theta(v)^T ENC_\theta(u)$.
 - \approx probability that u and v co-occur on a random walk over the graph.
- **def Node embeddings loss (for RW):** $\arg \min_\theta \sum_{v \in V} \sum_{u \in N_R(v)} -\log\left(\frac{\exp(\mathbf{z}_v^T \mathbf{z}_u)}{\sum_{w \in V} \exp(\mathbf{z}_v^T \mathbf{z}_w)}\right)$.
 - A heuristic that tries to approximate the above probability.
 - Where $N_R(v)$ is the set of neighbors of v according to the random-walk strategy R .
 - This expression is **easily optimized using SGD**.
 - Derivation:
 - $-\log\left(\frac{\exp(\mathbf{z}_v^T \mathbf{z}_u)}{\sum_{w \in V} \exp(\mathbf{z}_v^T \mathbf{z}_w)}\right) = \dots = C_v \mathbf{z}_v^T \mathbf{z}_w$.
 - v is assumed fixed, the expression is minimized when $\mathbf{z}_v^T \mathbf{z}_w$ is maximized.
 - Which is, when the similarity between v and u is maximized.
 - Since u is drawn from v neighbors, the **similarity between v and its neighbors is effectively maximized**.
 - f is assumed to be a monotonic function (eg a normalization of the dot product).
 - This loss is **quadratic**, in practice **negative sampling optimization** is used.
 - The denominator becomes $\sum_{w \notin N_R(v)} \exp(\mathbf{z}_v^T \mathbf{z}_w)$.

1.13.2.2 Graph embeddings

- The node embeddings can be aggregated into a graph embedding \mathbf{z}_G .
 - Using different strategies to obtain a graph embedding.
 - The most common strategies are:
 - Sum/average of the node embeddings.
 - Create a super-node that represents the entire graph.
- Early attempts tried to directly learn the embeddings of the nodes.
 - Instead of learning a function that maps the graph to a vector representation.
 - Idea:
 - Learn the embedding \mathbf{z}_v for all nodes $v \in V$.
 - Use the embeddings for downstream tasks.
- Downstream tasks:
 - **Node clustering**: cluster points $\{\mathbf{z}_u\}$ in the embedding space.
 - **Node prediction**: predict the label of a node u based on \mathbf{z}_u .
 - **Link prediction**: predict the existence of edge (u, v) based on $\{\mathbf{z}_u, \mathbf{z}_v\}$.
 - Combining the embeddings: concatenate, Hadamard, sum/avg, distance, etc.
 - **Graph classification**: predict the label of the entire graph based on \mathbf{z}_G (hardest task).

1.13.2.3 Limitations of earlier approaches

- Why not use the predefined similarity (which node embeddings try to reproduce) on graph nodes.
 - Those similarity are usually:
 - Implicit or expensive to compute.
 - Not suitable as ML inputs.
 - Tied to a fixed graph.
 - Hard to generalize or scale.
 - Node embeddings act as a **low-dimensional, learnable surrogate** for graph similarity.
- Limitations of earlier approaches:
 - **Transductive method**.

- Cannot obtain embeddings for nodes not in the training set.
 - Therefore cannot be applied to new graphs.
- Inductive methods: a model is built and used to predict on new data.
- Cannot capture structural similarity.
- Cannot utilize node, edge and graph features.

1.13.3 Modern approaches to Graph neural network

- Assumptions:
 - Arbitrary size and complex topological structure (no spatial locality like grids).
 - No fixed node ordering or reference point.
 - Often dynamic (graph changes over time) and have multimodal features.
- Using an adjacency matrix with extra information for features is problematic.
 - By swapping the labelling of two nodes, the input matrix changes (not desirable).
 - Main issues:
 - **Sensitive to node ordering.**
 - Not applicable to graphs of different sizes (the input layer of the DNN wouldn't match).

1.13.3.1 Invariance and equivariance

- **Permutation invariance:**
 - Graph does not have a canonical order of the nodes.
 - Graph and node **representations should be the same** for different ordering.
- **Permutation matrix:**
 - A square binary matrix that has:
 - Exactly one entry of 1 in each row and each column.
 - 0 elsewhere.
 - Multiplication:
 - If multiplied on the left of a matrix, it permutes the rows of the matrix.
 - If multiplied on the right of a matrix, it permutes the columns of the matrix.
 - The structure of the corresponding graph is the same, but the nodes are permuted.
 - The roles played by the nodes is no more the same, since their connectivity is different.
 - To restore the correct connectivity, the adjacency matrix must be permuted as well.
 - The permutation matrix is applied to swap the rows and columns of the adjacency matrix.
 - \mathbf{PAP}^T is computed.
 - Then the graph is the same as the original one, but with a different labelling.
- def **Permutation invariance:** the graph function f if $f(\mathbf{A}, \mathbf{X}) = f(\mathbf{PAP}^T, \mathbf{PX})$.
 - For any graph function $f : \mathbb{R}^{n \times n} \times \mathbb{R}^{n \times m} \rightarrow \mathbb{R}^d$ and for any permutation matrix \mathbf{P} .
 - $\mathbb{R}^{n \times n} \times \mathbb{R}^{n \times m}$, $\mathbb{R}^{n \times n}$ is for the adjacency matrix, $\mathbb{R}^{n \times m}$ for the feature matrix.
 - n is the number of nodes, m the number of features of a node.
- def **Permutation equivariance:** the node function f if $Pf(\mathbf{A}, \mathbf{X}) = f(\mathbf{PAP}^T, \mathbf{PX})$.
 - For any node function $f : \mathbb{R}^{n \times n} \times \mathbb{R}^{n \times m} \rightarrow \mathbb{R}^{n \times d}$ and for any permutation matrix \mathbf{P} .
 - n is the number of nodes, m the number of features of a node.
 - The output is permuted as is the input (it changes but it changes in the same manner).
- eg Permutation invariant and equivariant functions:
 - $f(\mathbf{A}, \mathbf{X}) = \mathbf{1}^T \mathbf{X} = \sum_i x_i$ is permutation-invariant.
 - Proof: $f(\mathbf{PAP}^T, \mathbf{PX}) = \mathbf{1}^T \mathbf{PX} = \sum_i x_i = f(\mathbf{A}, \mathbf{X})$.
 - $f(\mathbf{A}, \mathbf{X}) = \mathbf{X}$ is permutation-equivariant.
 - Proof: $f(\mathbf{PAP}^T, \mathbf{PX}) = \mathbf{PX} = Pf(\mathbf{A}, \mathbf{X})$.
 - $f(\mathbf{A}, \mathbf{X}) = \mathbf{AX}$ is permutation-equivariant.
 - Proof: $f(\mathbf{PAP}^T, \mathbf{PX}) = \mathbf{PAP}^T \mathbf{PX} = \mathbf{PAX} = Pf(\mathbf{A}, \mathbf{X})$.
 - **A MLP is neither permutation-invariant nor permutation-equivariant.**

1.13.3.2 GNN design principles

- Desiderata:
 - Ensure **permutation invariance** and **equivariance**.
 - By requiring PI and PE, the type of computation that the NN can performed is hugely restricted.
 - Use **layers** as reusable, permutation-equivariant units.
 - Stacking layers will not break permutation invariance.

- **Node-level predictions** will be automatically permutation-invariant.
- Design layers as **flexible differentiable functions**.
- Enable scalability with **parameter sharing** (to enable scalability to big data).
- A GNN is a **optimizable transformation** on all attributes of the graph (nodes, edges, global-context) **that preserves graph symmetries** (permutation invariances).
- *Goal:* design graph neural networks using permutation invariant/equivariant transformations and passing and aggregating information from neighbors.

1.13.4 Message Passing Framework

- Neighbour aggregation: generate node embeddings based on local network neighborhoods.
 - The neighborhood can be enlarged by apply the same produces for several steps.
 - NA can be described as a computational graph.
 - Computations can be done iteratively in steps (with parameter k for the neighbor depth).
 - Key distinctions are in how different approaches aggregate information across the layers.
- The MPF is a general way to define the update rule for each node.
 - A single GNN layer is considered.
- MPF consists of two main steps:
 - **Message computation:**
 - Each node v sends a message to its neighbors.
 - Based on its own features and the features of its neighbors.
 - **Message aggregation:**
 - Each node v aggregates messages received from its neighbors.
 - And updates its own features.
- For node v , $\mathbf{x}_v = \mathbf{h}_v^{(0)} \rightarrow \mathbf{h}_v^{(1)} \rightarrow \dots \rightarrow \mathbf{h}_v^{(L)}$.

1.13.4.1 Message computation

- **def Message computation:** $m_u^{(l)} = MSG^{(l)}(\mathbf{h}_u^{(l)})$.
- The message computed at level v by the node u .
- $MSG^{(l)}$ is usually a (simple, eg one-layer) NN that:
 - Takes as input the features of the node u at level $l - 1$.
 - Returns a message $\mathbf{m}_u^{(l)}$.
 - eg $\mathbf{m}_u^{(l)} = f_{\mathbf{W}^{(l)}}(\mathbf{h}_u^{(l-1)}) = \text{ReLU}(\mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)})$.
- **prop** Properties of message computation:
 - The **parameters** of the message computation function $MSG^{(l)}$ are **shared across the nodes**.
 - The function is **permutation-equivariant**.
 - Because the same function is applied to all nodes in the graph independently.
 - Both properties are crucial to ensure that the GNN can be applied to graphs of different sizes.

1.13.4.2 Message aggregation

- **def Message aggregation:** $\mathbf{h}_v^{(l)} = \text{Agg}^{(l)}(\{\mathbf{m}_u^{(l)}, \forall u \in N(v)\})$.
- Node v will aggregate the messages from its neighbors.
 - The aggregation function must be **permutation-invariant** (eg sum, mean, max).
 - Potential parameters associated with $\text{Agg}^{(l)}$ are shared across the nodes.
- Issue: information from node v itself could **get lost**.
 - Right now, the computation of $\mathbf{h}_v^{(l)}$ does not **directly** depend on $\mathbf{h}_v^{(l-1)}$.
 - Possible fix (not all GNN fix this):
 - Message: $\mathbf{m}_v^{(l)} = \mathbf{B}^{(l)} \mathbf{h}_v^{(l-1)}$.
 - Compute message from node v itself.
 - Aggregation: $\mathbf{h}_v^{(l)} = \text{Agg}^{(l)}(\{\mathbf{m}_u^{(l)}, \forall u \in N(v)\}, \mathbf{m}_v^{(l)})$.
 - Aggregates the message from node v itself after aggregating from neighbors.
 - eg Sum aggregation: $\mathbf{h}_v^{(l)} = \mathbf{m}_v^{(l)} + \sum_{u \in N(v)} \mathbf{m}_u^{(l)}$.

1.13.5 Graph Convolutional Networks

- A type of NN designed to operate on graph-structured data.
 - It generalizes the concept of **convolution** from traditional grid data to graph data.

- It aggregates feature information from a node's local neighborhood.
 - Each layer updates the feature representation of each node.
 - By combining its own features with the features of its neighbors.
- It does so by multiplying the adjacency matrix \mathbf{A} ($n \times n$) with the feature matrix $\mathbf{H}^{(l)}$ ($n \times m$).
 - For each node u and feature f the value $(\mathbf{AH})_{uf} = \sum_{k \in N(u)} A_{uk} H_{kf}$ contains:
 - The sum of f of all the nodes in the graph, weighted by edges connecting u to those nodes.
 - Not connected nodes won't contribute to the sum via multiplication via \mathbf{A} (zero values).
 - With weights 0 or 1, $(\mathbf{AH})_{uf} = \sum_{k \in N(u)} A_{uk} H_{kf} \rightarrow \sum_{k \in N(u)} \mathbf{H}_k = \sum_{k \in N(u)} \mathbf{h}_k$.

1.13.5.1 Normalization

- To avoid numerical instability, the adjacency matrix \mathbf{A} can be normalized.
 - By dividing each row by the sum of its elements.
- def Normalized adjacency matrix: $\tilde{\mathbf{A}} = \mathbf{D}^{-1} \mathbf{A}$.
 - $\tilde{A}_{uv} = \frac{A_{uv}}{\deg(u)} = \frac{A_{uv}}{|N(u)|}$.
 - \mathbf{D} : the diagonal degree matrix, $D_{uu} = \sum_v A_{uv}$.
- def Symmetric normalization: $\tilde{\mathbf{A}} = \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$.
 - $\tilde{A}_{uv} = \frac{A_{uv}}{\sqrt{D_{uu} D_{vv}}} = \frac{A_{uv}}{\sqrt{\deg(u) \deg(v)}} = \frac{A_{uv}}{\sqrt{|N(u)||N(v)|}}$.
 - $\mathbf{D}^{-1/2}$: the diagonal matrix with the inverse of the square root of the diagonal elements of \mathbf{D} .
 - Enhancement over the previous solution, standard nowadays for GCN.

1.13.5.2 GCN layers

- def GCN layer: $\mathbf{H}^{(l+1)} = \sigma(\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \mathbf{H}^{(l)} \mathbf{W}^{(l)})$.
 - $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$: adjacency matrix with added self-loops.
 - Self-loops are used to store information from the node itself (desiderata in message aggregation).
 - $\tilde{\mathbf{D}}$: diagonal degree matrix of $\tilde{\mathbf{A}}$, $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$.
 - $\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2}$: normalized $\tilde{\mathbf{A}}$.
 - $\mathbf{W}^{(l)}$: learnable weight matrix.
 - σ : activation function (eg ReLU).
- Benefits:
 - Locality: leveraging local graph structure makes them effective for tasks like node classification and link prediction.
 - Scalability: by stacking multiple GCN layers, information from larger neighborhoods can be captured.
- Everything can be described with the MPF:
 - $\mathbf{h}_u^{(l)} = \sigma(\sum_{v \in N(u) \cup \{u\}} \mathbf{W}^{(l)T} \frac{\mathbf{h}_v^{(l-1)}}{\sqrt{|N(v)||N(u)|}})$.
 - $\sum_{v \in N(u) \cup \{u\}}$ performs aggregation, while $\mathbf{W}^{(l)T} \frac{\mathbf{h}_v^{(l-1)}}{\sqrt{|N(v)||N(u)|}}$ performs message computation.

1.13.6 Training a GNN

- def Last layer node embeddings: $\mathbf{z}_v = \mathbf{h}_v^L = ENC_\theta(\mathbf{A}, \mathbf{x}_v)$.
 - Given a GNN with L layers and a GNN model ENC_θ (eg GCN, GraphSAGE, GAT).
- On top of the GNN, further layers (eg fully connected layers) are defined to perform downstream tasks.
 - $\mathbf{y}_v = DEC_{\theta'}(\mathbf{z}_v)$.
- To train the GNN, a loss function is defined and optimized via gradient-based methods.

1.13.6.1 Supervised learning with GNNs

- A node-level task is assumed.
- The goal is to minimize the following generic loss function: $\min_{\theta, \theta'} \sum_v \mathcal{L}(\mathbf{y}_v, DEC_{\theta'}(ENC_\theta(\mathbf{A}, \mathbf{x}_v)))$.
 - This loss can be optimized using gradient-based methods.
- The examples are the nodes of the graph.
 - To estimate the generalization error, a separate test set of nodes is kept.
 - The model is trained on the remaining nodes.

1.13.6.2 Model parameters

- In a GNN model the parameters are:
 - The parameters of the message computation and aggregation function.
 - The parameters for the downstream task (eg fully connected layers).
- In **Graph convolutional networks** the parameters are:
 - The weight matrices $\mathbf{W}^{(l)}$.
 - The parameters for solving the downstream task.

1.13.6.3 Stacking GNN layers

- In practice, multiple GNN layers can be stacked to **enhance model expressivity**.
 - But GNNs suffer from over-smoothing problem.
- **Over-smoothing problem:** all the node embeddings converge to the **same value**.
 - It's not possible to capture the structural properties of the graph.
 - The model is not able to distinguish between different nodes.
- Stacking multiple layers corresponds **enlarging the receptive field of a node**.
 - The set of nodes that influence the node's embedding.
 - Shared neighbors quickly grows when the number of hops.
 - Since the node embeddings is determined by its receptive field.
 - If two nodes have highly-overlapped receptive fields, then their **embeddings are highly similar**.
 - Stack **many** GNN layers \Rightarrow nodes will have **highly overlapped** receptive fields.
 - \Rightarrow node embeddings will be **highly similar** \Rightarrow suffer from the **over-smoothing problem**.

Enhancing GNN expressivity

- **Enhancing GNN expressivity** (main strategies):
 - **Increase the expressive power within each GNN layer.**
 - Message computation and aggregation become more complex.
 - **Add layers that do not pass messages.**
 - Add more layers before and after the GNN layers.
 - **Skip connections.**
 - Node embeddings in earlier GNN layers can sometimes better differentiate nodes.
 - Shortcuts are therefore added in the GNN.
 - Usually skip connection is implemented by summing the previous embeddings to the current ones.

2 Exercises

2.1 Mathematical foundations

2.1.1 Example 1 - Linear transformation

- Linear transformation:

- $\det(A) = a_1 \cdot a_4 - a_2 \cdot a_3 = 2 * 2 - 4 * 1 = 0.$
- How does it affects the vector spaced composed by $V_1 = [1, 0]$ and $V_2 = [0, 1]$?
 - $(A \cdot V_1)_{1,1} = 2$
 - $(A \cdot V_1)_{2,1} = 1$
 - $(A \cdot V_2)_{1,1} = 4.$
 - $(A \cdot V_2)_{2,1} = 2$
 - Both vectors lies on the same line (a dimension is lost).

2.1.2 Example 2 - Properties of derivatives

- Properties of derivatives:

- Power rule: $(x^4)' = 4x^3.$
- Linearity: $(3\sin(x) + x^2)' = 3(\sin(x))' + (x^2)' = 3\cos(x) + 2x.$
- Chain rule: $(\sin(x^2))' = \cos(x^2)(x^2)'2\cos(x^2)x.$
- Product rule: $(x^2x^3)' = 2x(x^3)' + x^2(3x^2)' = 5x^4 = (x^5)'.$
- Quotient rule: $(\frac{x^5}{x^2})' = \frac{5x^4(x^2) - x^5(2x)}{x^4} = \frac{3x^6}{x^4} = 3x^2 = (x^3)'.$

2.1.3 Example 3 - Partial derivatives

- Evaluate the derivative of z w.r.t. t :

- Where $(x, y) = (t^2, t)$ and $z = f(x, y) = x^2y^2.$
- Then $\frac{dz}{dt} = \frac{\partial z}{\partial x} \frac{dx}{dt} + \frac{\partial z}{\partial y} \frac{dy}{dt} = 2xy^2 \cdot 2t + 2x^2t \cdot 1 = 4t^5 + 2t^5 = 6t^5.$
- The same results could have been obtained simply evaluating $\frac{d}{dt}t^6 = 6t^5.$
 - By noticing that $f(x, y) = (x(t))^2 \cdot (y(t))^2 = (t^2)^2 \cdot (t)^2 = t^4t^2 = t^6.$

2.1.4 Example 4 - Discrete probability distributions

- Discrete probability distributions:

- A random variable x taking k possible values x_1, \dots, x_k is considered.
- The PMF $P(x) = \frac{1}{k}$ (**uniform probability**) is a valid PMF if:
 - It is defined over all possible states of x .
 - $\forall x \in x : 0 \leq P(x) = \frac{1}{k} \leq 1.$
 - $P(\Omega) = \sum_{x \in \{x_1, \dots, x_k\}} P(x) = \sum_{x \in \{x_1, \dots, x_k\}} \frac{1}{k} = 1.$
- Thanks to other properties it can be evaluated:
 - $P(\{x_1, x_2, x_3\} \cup \{x_4, x_5\}) = P(\{x_1, x_2, x_3\}) + P(\{x_4, x_5\}) = \frac{3}{k} + \frac{2}{k} = \frac{5}{k}.$
 - $P(\Omega - \{x_1, x_2\}) = 1 - \frac{2}{k} = \frac{k-2}{k} = P(\{x_3, \dots, x_k\}).$

2.1.5 Example 5 - Continuous probability distributions

- Continuous probability distributions:

- Consider a uniform distribution on an interval of the real numbers: $x \sim U(a, b).$
 - Where a and b are the endpoints of the interval, with $b > a.$
- The PDF $u(x; a, b)$ of the uniform distribution is:
 - $\frac{1}{b-a}$ if $x \in [a, b].$
 - 0, otherwise.
- The PDF $u(x; a, b)$ is a valid PDF if:
 - It is defined over all possible states of $x.$
 - $\forall x \in x : p(x) \geq 0.$
 - $\int_{-\infty}^{\infty} u(x; a, b) dx = \int_a^b \frac{1}{b-a} dx = \frac{1}{b-a} \int_a^b 1 dx = \frac{1}{b-a} \cdot (x|_a^b) = 1.$
 - $(x|_a^b) = (b - a).$

2.2 Introduction to neural networks

2.2.1 Example 1 - Perceptron

- **Perceptron:**
 - Data: $[1, 1, 1], [1, 1, 0], [1, 0, 1], [1, 0, 0]$ (each x_0 is the bias).
 - Given the following weights, the OR function is computed: $w_0 = 0, w_1 = 0.5, w_2 = 0.5$.
 - Expected output: -1 with $(0, 0)$, 1 for $(0, 1), (1, 0), (1, 1)$.
 - $x_0 \cdot w_0 + x_1 \cdot w_1 + x_2 \cdot w_2$:
 - $[1, 1, 1]: 1 \cdot 0 + 1 \cdot 0.5 + 1 \cdot 0.5 = 1 > 0 \rightarrow y = 1$.
 - $[1, 1, 0]: 1 \cdot 0 + 1 \cdot 0.5 + 0 \cdot 0.5 = 0.5 > 0 \rightarrow y = 1$.
 - $[1, 0, 1]: 1 \cdot 0 + 0 \cdot 0.5 + 1 \cdot 0.5 = 0.5 > 0 \rightarrow y = 1$.
 - $[1, 0, 0]: 1 \cdot 0 + 0 \cdot 0.5 + 0 \cdot 0.5 = 0 \leq 0 \rightarrow y = -1$.
 - Given the following weights, the AND function is computed: $w_0 = -1, w_1 = 1, w_2 = 1$.
 - Expected output: 1 with $(1, 1)$, -1 for $(0, 0), (0, 1), (1, 0)$.
 - $x_0 \cdot w_0 + x_1 \cdot w_1 + x_2 \cdot w_2$:
 - $[1, 1, 1]: 1 \cdot (-1) + 1 \cdot 1 + 1 \cdot 1 = 1 > 0 \rightarrow y = 1$.
 - $[1, 1, 0]: 1 \cdot (-1) + 1 \cdot 1 + 0 \cdot 1 = 0 \leq 0 \rightarrow y = -1$.
 - $[1, 0, 1]: 1 \cdot (-1) + 0 \cdot 1 + 1 \cdot 1 = 0 \leq 0 \rightarrow y = -1$.
 - $[1, 0, 0]: 1 \cdot (-1) + 0 \cdot 1 + 0 \cdot 1 = -1 \leq 0 \rightarrow y = -1$.
- Another set of weights which compute the AND function is: $w_0 = -0.8, w_1 = 0.5, w_2 = 0.5$.
- Just by modifying the weights, the computed function is completely different.

2.2.2 Example 2 - Perceptron Learning Algorithm

- Perceptron Learning Algorithm:
 - Input data: $([1, 1, 1], 1), ([1, 1, 0], -1), ([1, 0, 1], -1), ([1, 0, 0], -1)$.
 - Execution:
 - $w(0) = [0, 0, 0]$ (given this time, but usually randomized).
 - $\eta = 0.5$.
 - Examples can be evaluated in any orders.
 - $n = 0$, with $([1, 1, 1], 1)$, net input 0, net output -1 , but desired output 1 (misclassification).
 - $w(n+1) = w(n) + \eta d(n)x(n)$.
 - $w(1) = [0, 0, 0] + 0.5 \times [1, 1, 1] = [0.5, 0.5, 0.5]$.
 - Consider each pattern and consider if it is correctly predicted.
 - $([1, 1, 1], 1)$ is okay, $([1, 1, 0], -1)$ is not.
 - $n = 1$, with $([1, 1, 0], -1)$, net input 1, net output 1, but desired output -1 (misclassification).
 - $w(2) = [0.5, 0.5, 0.5] - 0.5 \times [1, 1, 0] = [0, 0, 0.5]$.
 - Consider each pattern and consider if it is correctly predicted.
 - $([1, 1, 1], 1)$ and $([1, 1, 0], -1)$ are okay, $([1, 0, 1], -1)$ is not.
 - $w(2) = [0.5, 0.5, 0.5] - 0.5[1, 0, 0] = [0, 0.5, 0.5]$.
 - $w(3) = [0, 0.5, 0.5] - 0.5[1, 1, 0] = [-0.5, 0, 0.5]$.
 - $w(4) = [-0.5, 0, 0.5] + 0.5[1, 1, 1] = [0, 0.5, 1]$.
 - $w(5) = [0, 0.5, 1] - 0.5[1, 0, 1] = [-0.5, 0.5, 0.5]$ (solution).
 - The AND problem is **linearly separable**.
 - The convergence theorem provides that a solution will be found.
 - In this example and given the patterns evaluation order, in $n = 5$ iterations.

2.2.3 Exercise 1 - Sigmoid neurons

- **Sigmoid neurons simulating perceptrons:**
 - Suppose to multiply all the weights and biases in a network of perceptrons by a positive constant $c > 0$.
 - Show that the behavior of the network doesn't change.
 - Suppose that the overall input to the same network has been chosen (fixed input).
 - Suppose that $\mathbf{w} \cdot \mathbf{x} + b \neq 0$ to any particular perceptron in the network.
 - Then replace all the perceptrons in the network by **sigmoid neurons**.
 - And multiply the weights and biases by a positive constant $c > 0$.
 - Show that in the limit as $c \rightarrow \infty$, this new network behaves the same as the original one.
 - How can this fail when $\mathbf{w} \cdot \mathbf{x} + b = 0$ for one of the perceptrons?

2.2.3.1 Solution

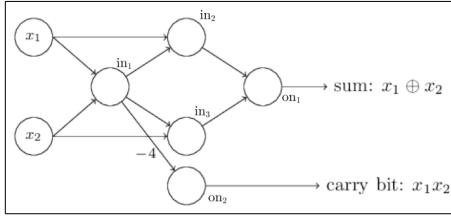


Figure 8: NN-based binary adder with carry (default weight: -2).

- Perceptron network:

- After the multiplication the equation is $\mathbf{w}_c \cdot \mathbf{x} + bc$.
- The *binary adder with carry* perceptron network is taken as reference.
- The chosen positive constant is $c = 4$.
 - With $x_1 = 0$ and $x_2 = 1$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [0, 1]$, $0 \cdot (-2 \cdot 4) + 1 \cdot (-2 \cdot 4) + (3 \cdot 4) = 4 > 0 \rightarrow y_{in_1} = 1$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [0, 1]$, $0 \cdot (-2 \cdot 4) + 1 \cdot (-2 \cdot 4) + (3 \cdot 4) = 4 > 0 \rightarrow y_{in_2} = 1$.
 - in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [1, 1]$, $1 \cdot (-2 \cdot 4) + 1 \cdot (-2 \cdot 4) + (3 \cdot 4) = -4 \leq 0 \rightarrow y_{in_3} = 0$.
 - on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [1, 0]$, $1 \cdot (-2 \cdot 4) + 0 \cdot (-2 \cdot 4) + (3 \cdot 4) = 4 > 0 \rightarrow y_{on_1} = 1$.
 - on_2 : with $\mathbf{x} = [y_{in_1}] = [1]$, $1 \cdot (-4 \cdot 4) + (3 \cdot 4) = -4 \leq 0 \rightarrow y_{on_2} = 0$.
 - Output: $\mathbf{y} = [1, 0]$, where $0 \oplus 1 = 1$ with carry = 0 (correct).
 - With $x_1 = 1$ and $x_2 = 1$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [1, 1]$, $y_{in_1} = 0$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [1, 0]$, $y_{in_2} = 1$.
 - in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [1, 0]$, $y_{in_3} = 1$.
 - on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [1, 1]$, $y_{on_1} = 0$.
 - on_2 : with $\mathbf{x} = [y_{in_1}] = [0]$, $y_{on_2} = 1$
 - Output: $\mathbf{y} = [0, 1]$, where $1 \oplus 1 = 0$ with carry = 1 (correct).
 - With $x_1 = 1$ and $x_2 = 0$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [1, 0]$, $y_{in_1} = 1$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [1, 1]$, $y_{in_2} = 0$.
 - in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [0, 1]$, $y_{in_3} = 1$.
 - on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [0, 1]$, $y_{on_1} = 1$.
 - on_2 : with $\mathbf{x} = [y_{in_1}] = [1]$, $y_{on_2} = 0$
 - Output: $\mathbf{y} = [1, 0]$, where $0 \oplus 1 = 1$ with carry = 0 (correct).
 - With $x_1 = 0$ and $x_2 = 0$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [0, 0]$, $y_{in_1} = 1$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [0, 1]$, $y_{in_2} = 1$.
 - in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [0, 1]$, $y_{in_3} = 1$.
 - on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [1, 1]$, $y_{on_1} = 0$.
 - on_2 : with $\mathbf{x} = [y_{in_1}] = [1]$, $y_{on_2} = 0$
 - Output: $\mathbf{y} = [1, 0]$, where $0 \oplus 0 = 0$ with carry = 0 (correct).
- All four cases match the original network and the corresponding *binary adder with carry* function.

- When multiplied by a positive constant, the perceptron NN behaves the same as the original one.
- Sigmoid network:
 - The *binary adder with carry* perceptron network is taken as reference.
 - The chosen positive constant is $c = 1$ (same as perceptron NN, but with sigmoid neurons):
 - With $x_1 = 0$ and $x_2 = 1$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [0, 1]$, $0 \cdot -2 + 1 \cdot -2 + 3 = 1 \rightarrow y_{in_1} = \sigma(1) = 0.731$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [0, 0.731]$, $0 \cdot -2 + 0.731 \cdot -2 + 3 = 1.538 \rightarrow y_{in_2} = \sigma(1.538) = 0.823$.
 - in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [1, 0.731]$, $1 \cdot -2 + 0.731 \cdot -2 + 3 = -0.462 \rightarrow y_{in_3} = \sigma(-0.462) = 0.386$.
 - on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [0.823, 0.386]$, $0.823 \cdot -2 + 0.386 \cdot -2 + 3 = 0.582 \rightarrow y_{on_1} = \sigma(0.582) = 0.641$.
 - on_2 : with $\mathbf{x} = [y_{in_1}] = [0.731]$, $0.731 \cdot -4 + 3 = 0.076 \rightarrow y_{on_2} = \sigma(0.076) = 0.518$.
 - Output: $\mathbf{y} = [0.641, 0.518]$, very distant from the actual solution $[1, 0]$.
 - The chosen positive constant is $c = 4$.
 - With $x_1 = 0$ and $x_2 = 1$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [0, 1]$, $0 \cdot (-2 \cdot 4) + 1 \cdot (-2 \cdot 4) + (3 \cdot 4) = 4 \rightarrow y_{in_1} = \sigma(4) = 0.982$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [0, 0.982]$, $0 \cdot (-2 \cdot 4) + 0.982 \cdot (-2 \cdot 4) + (3 \cdot 4) = 4.143 \rightarrow y_{in_2} =$

$$\sigma(4.134) = 0.984.$$

- in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [1, 0.982]$, $1 \cdot (-2 \cdot 4) + 0.982 \cdot (-2 \cdot 4) + (3 \cdot 4) = -3.872 \rightarrow y_{in_3} = \sigma(-3.872) = 0.020$.
- on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [0.984, 0.020]$, $0.984 \cdot (-2 \cdot 4) + 0.020 \cdot (-2 \cdot 4) + (3 \cdot 4) = 3.968 \rightarrow y_{on_1} = \sigma(3.968) = 0.981$.
- on_2 : with $\mathbf{x} = [y_{in_1}] = [0.982]$, $0.982 \cdot (-4 \cdot 4) + (3 \cdot 4) = -3.712 \rightarrow y_{on_2} = \sigma(-3.712) = 0.024$.
- *Output*: $\mathbf{y} = [0.981, 0.024]$, closer but still not exact from the actual solution $[1, 0]$.

- The larger chosen positive constant is $c = 123456$.

- With $x_1 = 0$ and $x_2 = 1$:
 - in_1 : with $\mathbf{x} = [x_1, x_2] = [0, 1]$, $y_{in_1} = \sigma(123456) = 1.0$.
 - in_2 : with $\mathbf{x} = [x_1, y_{in_1}] = [0, 1]$, $y_{in_2} = \sigma(123456) = 1.0$.
 - in_3 : with $\mathbf{x} = [x_2, y_{in_1}] = [1, 1]$, $y_{in_3} = \sigma(-123456) \approx 0$.
 - on_1 : with $\mathbf{x} = [y_{in_2}, y_{in_3}] = [1, 0]$, $y_{on_1} = \sigma(123456) = 1.0$.
 - on_2 : with $\mathbf{x} = [y_{in_1}] = [1]$, $y_{on_2} = \sigma(-123456) \approx 0$.
 - *Output*: $\mathbf{y} = [1, 0]$, where $0 \oplus 1 = 1$ with carry = 0 (correct).

- Even though this is a partial solution, it's clear that with $c \rightarrow \infty$ the **original function is approximated**.

2.2.4 Example 3 - Weights representation

- **Weights representation:**

- The third (and output) L of a NN has two neurons.

- Its weights vector is then (with a previous layer with three neurons): $\mathbf{W}^3 = \begin{pmatrix} w_{11}^3 & w_{12}^3 & w_{13}^3 \\ w_{21}^3 & w_{22}^3 & w_{23}^3 \end{pmatrix}$.
- $\mathbf{W}^3 \mathbf{z}^{L-1} = \begin{pmatrix} w_{11}^3 & w_{12}^3 & w_{13}^3 \\ w_{21}^3 & w_{22}^3 & w_{23}^3 \end{pmatrix} \cdot \begin{bmatrix} z_1^2 \\ z_2^2 \\ z_3^2 \end{bmatrix} = \begin{bmatrix} w_{11}^3 \cdot \mathbf{z}^2 \\ w_{21}^3 \cdot \mathbf{z}^2 \end{bmatrix}$.

2.2.5 Example 4 - First equation of BP

- Computing BP1 using the **quadratic cost**:

- $C = \frac{1}{2} \sum_j (y_j - z_j^L) \implies \frac{\partial C}{\partial z_j^L} = (z_j^L - y_j)$.
 - C is actually C_x .
 - $\frac{1}{2} \|y(x) - z^L(x)\|_2^2 = [y(x) - z^L(x)]^T [y(x) - z^L(x)] = \sum_j (y_j(x) - z_j^L(x))(y_j(x) - z_j^L(x))$.
 - The goal is to compute the derivative of C w.r.t. z_j^L (consider everything else as a constant).
 - In $\frac{1}{2}((y_1 - z_1^L)^2 + \dots + (y_j - z_j^L)^2 + \dots)$ the only term dependent on z_j^L is $(y_j - z_j^L)^2$.
 - $\frac{\partial C}{\partial z_j^L} (y_j - z_j^L)^2 = 2(y_j - z_j^L) \cdot -1 = 2(z_j^L - y_j)$.
 - -1 is the derivative of $y_j - z_j^L$ w.r.t. z_j^L .
 - The $\frac{1}{2}$ in C cancels out with the 2 in $2(z_j^L - y_j)$.
 - $\frac{\partial C}{\partial z_j^L} = (z_j^L - y_j)$.
 - To compute z_j^L , the information is conveyed through the NN up to layer L .
 - While y_j is a given constant (a value, 0 or 1).
 - Therefore $\frac{\partial C}{\partial z_j^L} = (z_j^L - y_j)$ is a simple value.
- As it is $\delta_j^L = \frac{\partial C}{\partial z_j^L} \sigma'(a_j^L)$, therefore BP1 is easily computable.

2.3 Hopfield networks

2.3.1 Exercise 1 - Hopfield networks

2.3.1.1 Solution

Information withdrawal phase

- Compute the stable state of the above network with the initial state $[-1, -1, 1]$:
 - Initial state:

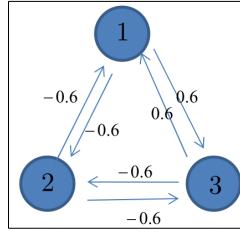


Figure 9: Hopfield networks for the example.

- $y_1(0) = -1, y_2(0) = -1, y_3(0) = 1$.
- The neurons are picked randomly:
 - $v_1(1) = -0.6 \cdot -1 + 0.6 \cdot 1 = 1.2 \rightarrow \varphi(v_1(1)) = 1 \rightarrow [1, -1, 1]$.
 - $v_2(2) = -0.6 \cdot 1 + -0.6 \cdot 1 = -1.2 \rightarrow \varphi(v_2(2)) = -1 \rightarrow [1, -1, 1]$.
 - $v_3(3) = 0.6 \cdot 1 + -0.6 \cdot -1 = 0 \rightarrow \varphi(v_3(3)) = -1 \rightarrow [1, -1, 1]$.
 - $v_1(4) = -0.6 \cdot -1 + 0.6 \cdot 1 = 1.2 \rightarrow \varphi(v_1(4)) = 1 \rightarrow [1, -1, 1]$.
- Stable state: $[1, -1, 1]$ (output of the network).
 - With a different evaluation order, the stable state can be different.
 - The other stable state is $[-1, 1, -1]$ (opposite of the one found, property of HN).

Storage phase

- Compute the weights with the initial memory $f_1 = [1, -1, 1]$:
 - $w_{12} = 1 \cdot (-1) = -1 = w_{21}$.
 - $w_{23} = (-1) \cdot 1 = -1 = w_{32}$.
 - $w_{31} = 1 \cdot 1 = 1 = w_{13}$.
- Compute the weights with the initial memories $f_1 = [1, -1, 1]$ and $f_2 = [1, 1, 1]$:
 - $w_{12} = \frac{1 \cdot -1 + 1 \cdot 1}{2} = 0 = w_{21}$.
 - $w_{23} = \frac{-1 \cdot 1 + 1 \cdot 1}{2} = 0 = w_{32}$.
 - $w_{31} = \frac{1 \cdot 1 + 1 \cdot -1}{2} = 1 = w_{13}$.
- Apply the memorization rule so that $[1, -1, 1]$ and $[-1, 1, -1]$ are stable states.
 - $w_{12} = \frac{1 \cdot -1 + -1 \cdot 1}{2} = -1 = w_{13}$.
 - $w_{23} = \frac{-1 \cdot 1 + -1 \cdot -1}{2} = -1 = w_{32}$.
 - $w_{31} = \frac{1 \cdot 1 + -1 \cdot -1}{2} = 1 = w_{13}$.

2.3.2 Exercise 2 - Restricted Boltzmann machines

- Calculate $p(h_i = 1)$:
 - Visible layer composed by v_1 and v_2 of values: 1 and 0.
 - Hidden layer composed by a single node h_1 with $w_{11} = 0.01$ and $w_{21} = 0.02$.
 - $\Delta E_i = 1 \cdot 0.01 + 0 \cdot 0.02 = 0.01$.

2.3.2.1 Solution

- Two ways of computing ΔE_i :
 - Computing $E(h_1 = 1)$ and $E(h_1 = 0)$ and then $\Delta E_i = E(s_i = 0) - E(s_i = 1)$.
 - Or Computing $\Delta E_i = \sum w_{ij} S_j = 1 \cdot 0.01 + 0 \cdot 0.02 = 0.01$.
 - Then $p(h_1 = 1)$ can be computed: $p(h_1 = 1) = 0.50251$.
- Sampling is used to attribute a state (0 and 1) using the probability.

2.3.3 Exercise 3 - Restricted Boltzmann machines

- Apply contrastive divergence by learning $[1, 0]$:
 - Given:
 - Visible layer composed by v_1 and v_2 .
 - Hidden layer composed by a single node h_1 with $w_{11} = 1$ and $w_{21} = 1$.
 - $\varepsilon = 0.1$ (learning rate).
 - Assume that sampling will output 1 in case $p(s = 1) \geq 0.5$, 0 otherwise.
 - In general, sampling will return 1 with probability p , hence not always, unless $p(s = 1) = 1$.

2.3.3.1 Solution

- Execution of contrastive divergence:
 - Epoch #1:
 - $t = 0$:
 - v_1 and v_2 are the ones in input.
 - $h_1: \Delta E_{h_1} = \sum_j w_{h_1 v_j} v_j = 1 * 1 + 1 * 0 = 1 \rightarrow p(h_1 = 1) = \frac{1}{1+e^{-\Delta E_{h_1}/T}} = 0.73.$
 - By sampling, with probability $0.73 \geq 0.5$, $h_1 = 1$.
 - $\langle v_1 h_1 \rangle^0 = 1 \cdot 1 = 1$ and $\langle v_2 h_1 \rangle^0 = 0 \cdot 1 = 0$.
 - $t = 1$:
 - To reconstruct v_1 and v_2 , h_1 is the one computed at $t = 0$.
 - $v_1: \Delta E_{v_1} = \sum_j w_{v_1 h_j} h_j = 1 * 1 = 1 \rightarrow p(v_1 = 1) = \frac{1}{1+e^{-\Delta E_{v_1}/T}} = 0.73.$
 - By sampling, with probability $0.73 \geq 0.5$, $v_1 = 1$.
 - $v_2: \Delta E_{v_2} = \sum_j w_{v_2 h_j} h_j = 1 * 1 = 1 \rightarrow p(v_2 = 1) = \frac{1}{1+e^{-\Delta E_{v_2}/T}} = 0.73.$
 - By sampling, with probability $0.73 \geq 0.5$, $v_2 = 1$.
 - $h_1: \Delta E_{h_1} = \sum_j w_{h_1 v_j} v_j = 1 * 1 + 1 * 1 = 2 \rightarrow p(h_1 = 1) = \frac{1}{1+e^{-\Delta E_{h_1}/T}} = 0.88.$
 - By sampling, with probability $0.88 \geq 0.5$, $h_1 = 1$.
 - $\langle v_1 h_1 \rangle^1 = 1 \cdot 1 = 1$ and $\langle v_2 h_1 \rangle^1 = 1 \cdot 1 = 1$.
 - Weights update:
 - $\Delta w_{h_1 v_1} = \epsilon(\langle v_1 h_1 \rangle^0 - \langle v_1 h_1 \rangle^1) = 0.1((1 \cdot 1) - (1 \cdot 1)) = 0.$
 - $w_{h_1 v_1}(2) = w_{h_1 v_1}(1) + \Delta w_{h_1 v_1} = 1 + 0 = 1$ (unchanged).
 - $\Delta w_{h_1 v_2} = \epsilon(\langle v_2 h_1 \rangle^0 - \langle v_2 h_1 \rangle^1) = 0.1((0 \cdot 1) - (1 \cdot 1)) = -0.1.$
 - $w_{h_1 v_2}(2) = w_{h_1 v_2}(1) + \Delta w_{h_1 v_2} = 1 + (-0.1) = 0.9.$
-

2.4 Recurrent neural networks

2.4.1 Example 1 - Long Short Term Memory

- LSTM forget memory:
 - Let $\sigma = 1$ for arguments > 0 , 0 otherwise.
 - Let's $W_f = [1, 1, 1, -10; 1, 1, 2, -10]$ and $b_f = 0$.
 - $[h_{t-1}, x_t] = [1, 1, 0, 1]$ and suppose $x_t = [0, 1]$ codifies “.”, the dot character.
 - Then $f_t = [0, 0]$.
 - Intuitively (and informally) erases everything from C_{t-1} .
 - Since weights are learned via BP, gates are not as interpretable.